

Does Corporate AI Talk Predict Stock Returns? Evidence from LLM-Based Disclosure Analysis and the AI Washing Phenomenon



Author: Daniel Kosten (2704227)

Supervisor: Natalie Kessler

Date: 26/06/2025

Word Count: 20,149 (Excluding References & Appendix)

Table of Contents:

ABSTRACT	6
1. INTRODUCTION.....	7
2. LITERATURE REVIEW AND HYPOTHESES DEVELOPMENT.....	15
3. DATA	24
3.1 Sample Selection and Filtering Process.....	24
3.2 Data Sources and Collection.....	25
3.3 Sample Characteristics.....	26
3.4 Descriptive Statistics	27
3.4.1 Dependent Variable Properties.....	27
3.4.2 Control Variables Properties	28
4. METHODOLOGY	28
4.1 AI Disclosure Measurement Framework.....	29
4.1.1 AI Factor Design and Categories	29
4.1.2 Prompt Engineering and Validation	31
4.1.3 Cross-Model Reliability Testing.....	33
4.1.4 Composite AI Score Construction	35
4.2 Control Variable Construction	36
4.2.1 ‘Book’ Control Variables	36
4.2.2 Factor Risk Loadings	36
4.2.3 Data Treatment and Quality Controls	37
4.3 Portfolio Analysis Methods	37
4.3.1 Quintile Portfolio Construction	38
4.3.2 Sector-Neutral Analysis	38
4.3.3 Statistical Testing.....	38
4.3.4 Sector-Specific Analysis.....	39
4.4 Regression Framework.....	39
4.4.1 Baseline Specification and Fixed Effects Strategy	40
4.4.2 Progressive Control Inclusion Strategy	41
4.4.3 Individual Factors vs. Composite Score Approach.....	41
4.4.4 Sector-Specific Analysis	42
4.4.5 Identification Challenges and Robustness	42
4.5 Explanatory Case Study Design	44
4.5.1 Methodological Justification.....	44
4.5.2 Case Selection Strategy	44
4.5.3 Data Collection and Analysis.....	45
4.5.4 Validity and Limitations.....	45

5. INTERMEDIATE RESULTS	46
5.1 <i>AI Score Characteristics and Patterns</i>	46
5.1.1 AI Score Characteristics and Patterns	46
5.1.2 Cross-Model Reliability Analysis.....	49
5.2 <i>Portfolio Analysis</i>	52
5.2.1 Individual AI Factor Quintile Sorts	52
5.2.2 Sector-Specific Quintile Sorts.....	55
6. MAIN RESULTS	56
6.1 <i>Individual AI Factor Analysis</i>	57
6.1.1 Build-up Model Demonstration	57
6.1.2 Factor Results	58
6.2 <i>Sector Regression Analysis</i>	60
6.2.1 Cross-Sector Regression Results	60
6.2.2 Individual Factor Analysis by Sector.....	62
6.3 <i>Hypothesis Testing and Theoretical Framework</i>	63
6.3.1 Hypothesis Testing Results and Economic Significance	63
6.3.2 Theoretical Framework for AI Disclosure Paradox	65
6.4 <i>Explanatory Case Studies</i>	67
6.4.1 Case Study Analysis.....	67
6.4.2 Systematic Behavioural Patterns	68
6.5 <i>Practical Implications</i>	70
7. CONCLUSION.....	70
REFERENCES	73
APPENDIX	90
<i>Appendix A: Examples of LLM Scoring Quotes</i>	90
<i>Appendix B: Complete LLM Prompt</i>	92
<i>Appendix C: Case Study Analysis</i>	98
<i>Appendix D: Additional Analysis Regression Tables & Figures</i>	106
<i>Appendix E: Additional Control Variable Summary Statistics</i>	117
<i>Appendix F: Data Validation and Robustness</i>	118

Tables:

TABLE 1: <i>SAMPLE CONSTRUCTION AND FILTERING PROCESS, 2020–2024</i>	25
TABLE 2: <i>SECTOR DISTRIBUTION OF THE FINAL SAMPLE VERSUS THE RUSSELL 3000 INDEX</i>	26
TABLE 3: <i>DESCRIPTIVE STATISTICS FOR BUY-AND-HOLD EXCESS RETURNS</i>	27
TABLE 4: <i>DESCRIPTIVE STATISTICS FOR REGRESSION CONTROL VARIABLES</i>	28
TABLE 5: <i>AI DISCLOSURE SCORING FRAMEWORK</i>	30
TABLE 6: <i>DESCRIPTIVE STATISTICS OF AI DISCLOSURE FACTOR SCORES</i>	46
TABLE 7: <i>DISTRIBUTION OF COMPOSITE AI DISCLOSURE SCORES BY GICS SECTOR</i>	48
TABLE 8: <i>INTER-MODEL RELIABILITY METRICS FOR AI FACTOR SCORES (GPT-4o-MINI VS. GEMINI FLASH 2.5)</i>	50
TABLE 9: <i>PERFORMANCE OF PORTFOLIOS SORTED ON INDIVIDUAL AI DISCLOSURE FACTORS</i>	52
TABLE 10: <i>PERFORMANCE OF PORTFOLIOS SORTED ON THE COMPOSITE AI SCORE ACROSS DIFFERENT HORIZONS</i>	53
TABLE 11: <i>SECTOR-NEUTRAL PORTFOLIO SORTS ON THE COMPOSITE AI SCORE</i>	55
TABLE 12: <i>PROGRESSIVE REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE COMPOSITE AI SCORE</i>	57
TABLE 13: <i>REGRESSIONS OF 12-MONTH EXCESS RETURNS ON INDIVIDUAL AI DISCLOSURE FACTORS</i>	58
TABLE 14: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE COMPOSITE AI DISCLOSURE SCORE</i>	60
TABLE 15: <i>ILLUSTRATIVE EXAMPLES OF AI DISCLOSURE FACTOR SCORING</i>	90
TABLE 16: <i>OVERVIEW OF SELECTED CASE STUDIES: HIGH AI DISCLOSURE AND NEGATIVE RETURNS</i>	98
TABLE 17: <i>PROGRESSIVE REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE TALENT & INVESTMENT FACTOR</i>	106
TABLE 18: <i>PROGRESSIVE REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE COMPOSITE AI SCORE</i>	107
TABLE 19: <i>REGRESSIONS OF 9-MONTH EXCESS RETURNS ON INDIVIDUAL AI DISCLOSURE FACTORS</i>	107
TABLE 20: <i>REGRESSIONS OF 6-MONTH EXCESS RETURNS ON INDIVIDUAL AI DISCLOSURE FACTORS</i>	108
TABLE 21: <i>REGRESSIONS OF 3-MONTH EXCESS RETURNS ON INDIVIDUAL AI DISCLOSURE FACTORS</i>	109
TABLE 22: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE STRATEGIC DEPTH FACTOR</i>	110
TABLE 23: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE DISCLOSURE SENTIMENT FACTOR</i>	111
TABLE 24: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE TALENT & INVESTMENT FACTOR</i>	112
TABLE 25: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE AI WASHING INDEX</i>	113
TABLE 26: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE FORWARD-LOOKING FACTOR</i>	114
TABLE 27: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE RISK - EXTERNAL THREATS FACTOR</i>	114
TABLE 28: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE RISK - NON-ADOPTION FACTOR</i>	115
TABLE 29: <i>SECTOR-LEVEL REGRESSIONS OF 12-MONTH EXCESS RETURNS ON THE RISK - OWN ADOPTION FACTOR</i>	116
TABLE 30: <i>DESCRIPTIVE STATISTICS FOR FAMA-FRENCH-CARHART FACTOR BETAS (3-, 6-, AND 9-MONTH HORIZONS)</i>	117
TABLE 31: <i>PEARSON CORRELATION MATRIX OF KEY VARIABLES</i>	118
TABLE 32: <i>DESCRIPTIVE STATISTICS OF KEY VARIABLES BEFORE WINSORIZATION</i>	119
TABLE 33: <i>DESCRIPTIVE STATISTICS OF KEY VARIABLES AFTER WINSORIZATION</i>	120
TABLE 34: <i>INTER-MODEL RELIABILITY METRICS FOR AI FACTOR SCORES (GPT-4o-MINI VS. GEMINI FLASH 1.5)</i>	122
TABLE 35: <i>INTER-MODEL RELIABILITY METRICS FOR AI FACTOR SCORES (GEMINI FLASH 1.5 VS. GEMINI FLASH 2.5)</i>	123
TABLE 36: <i>DIAGNOSTIC TESTS FOR REGRESSION MODELS</i>	124
TABLE 37: <i>VARIANCE INFLATION FACTORS (VIF) FOR REGRESSION MODELS</i>	124

Figures:

FIGURE 1: <i>GROWTH IN CORPORATE AI DISCUSSION FOLLOWING THE RELEASE OF CHATGPT</i>	7
FIGURE 2: <i>MEAN AI DISCLOSURE FACTOR SCORES BY GICS SECTOR</i>	47
FIGURE 3: <i>DISTRIBUTION OF AI FACTOR GRADES ASSIGNED BY EACH LLM FOR STRATEGIC DEPTH AND DISCLOSURE SENTIMENT</i>	49
FIGURE 4: <i>INTER-MODEL RANK CORRELATION (KENDALL'S TAU) OF AI FACTOR SCORES</i>	50
FIGURE 5: <i>CUMULATIVE PERFORMANCE OF PORTFOLIOS SORTED ON THE COMPOSITE AI DISCLOSURE SCORE</i>	54
FIGURE 6: <i>RELATIONSHIP BETWEEN COMPOSITE AI SCORE AND SUBSEQUENT 12-MONTH EXCESS RETURN</i>	55
FIGURE 7: <i>BINNED SCATTERPLOTS OF 12-MONTH EXCESS RETURNS AGAINST AI FACTOR GRADES</i>	106
FIGURE 8: <i>DISTRIBUTION OF AI FACTOR GRADES ASSIGNED FOR TALENT & INVESTMENT AND AI WASHING</i>	120
FIGURE 9: <i>DISTRIBUTION OF AI FACTOR GRADES ASSIGNED FOR RISK - OWN ADOPTION AND RISK - EXTERNAL TREATS</i>	121
FIGURE 10: <i>DISTRIBUTION OF AI FACTOR GRADES ASSIGNED FOR RISK - NON-ADOPTION AND FORWARD-LOOKING</i>	121
FIGURE 11: <i>INTER-MODEL RANK CORRELATION (SPEARMAN'S RHO) OF AI FACTOR SCORES</i>	122
FIGURE 12: <i>INTER-MODEL LINEAR CORRELATION (PEARSON'S R) OF AI FACTOR SCORES</i>	122

Abstract

This paper examines whether the language public firms use to discuss Artificial Intelligence (AI) in 10-K ¹reports, viewed as a signal of their underlying AI strategy, influences subsequent market performance. Using a panel of 799 Russell 3000 constituents from 2020 to 2024, the study applies three Large Language Models (LLMs), GPT 4o mini and two versions of Gemini Flash, to score eight disclosure dimensions, which include strategic depth, sentiment, three risk categories, forward-looking statements, AI washing, and talent & investment. Model scores are averaged to create an AI factor suite and an aggregate index. Cross-sectional regressions relate these measures to three-, six-, nine- and twelve-month buy-and-hold excess returns while controlling for firm characteristics, Fama-French-Carhart six-factor betas, and fixed effects. The results contradict signalling theory, as higher scores on seven of the eight AI dimensions predict significantly lower future stock returns. This negative effect is particularly strong in traditional sectors, suggesting investors interpret intensive AI talk as a signal of underlying business weakness rather than strategic strength.

¹ Form 10-K is the official annual report required by the U.S. Securities and Exchange Commission (SEC) for publicly traded companies. This study's analysis focuses on two key narrative sections: Item 1A (Risk Factors) and Item 7 (Management's Discussion & Analysis).


1. Introduction

The Rise of AI in Company Talk and Investor Attention

Artificial Intelligence (AI) is rapidly becoming central to corporate strategy and operations (PwC, 2024; Mishra et al., 2022). It increasingly affects how companies make long-term plans and interact with investors and regulators (Campbell, et al., 2025; Cockburn et al., 2018). This trend is clear in corporate reporting, especially in U.S. Securities and Exchange Commission (SEC) 10-K filings, where terms like “machine learning,” “Large Language Model (LLM),” and “Generative AI (GenAI)” appear more often (Mishra, 2024; Alston & Bird, 2024; Heyliger et al., 2023). Figure 1 shows the increasing use of AI-related language in these documents (Accenture, 2024).

Figure 1:

Growth in Corporate AI Discussion Following the Release of ChatGPT



Note: This figure shows the quarterly number of companies mentioning AI-related terms (purple bars, left axis) and the total number of AI mentions (dashed line, right axis) in earnings call transcripts from Q1 2022 to Q3 2023. Source: Accenture Technology Vision 2024.

The growing interest in AI matches its description as a general-purpose technology that is powerful, constantly improving, and able to drive innovation across many sectors (Bresnahan & Trajtenberg, 1995; Cockburn et al., 2018).

However, this trend also poses difficulties in interpretation. The rapid increase in AI references may reflect not only real adoption but also impression management. Firms might use AI-heavy narratives to influence stakeholder perceptions (Merkl-Davies & Brennan,

2007). This is consistent with attention-driven investment theories, which argue that investors are more likely to allocate capital to firms that attract media or trading attention (Barber & Odean, 2008). When companies see that AI-related stories get rewarded in the market, others might do the same. This pattern can create a feedback loop typical of the technology hype cycle. This occurs when initial excitement leads to very high expectations, and eventually, more realistic applications (Fenn, 1995; Manning & Napier, 2024).

A specific form of impression management involves “AI washing”. This describes the use of misleading or shallow language about AI to attract investors (Seele & Schultz, 2022; Al Haddi, 2024). Investors must now distinguish between real AI strategies and vague statements (Uslaner & Coquin, 2025). Growing regulatory concern has led to increased enforcement. SEC Chair Gary Gensler (2024) stated: “We’ve seen time and again that when new technologies come along, they can create buzz from investors as well as false claims by those purporting to use those new technologies. Investment advisers should not mislead the public by saying they are using an AI model when they are not. Such AI washing hurts investors.” (U.S. Securities and Exchange Commission, 2024). Legal reviews are warning companies about these risks and advise them to verify their AI claims to prevent regulatory measures (Heyliger et al., 2023).

However, regulatory action against the worst cases of AI washing does not remove the wider pressures organisations face to manage their public image. Evaluating the credibility of AI disclosures is challenging because genuine integration relies on intangible assets (e.g., proprietary code, specialised staff) that are difficult for outsiders to verify, creating a condition of information asymmetry (Akerlof, 1970; Greene, 2025; Barrios et al., 2025).

This problem is intensified by the “black box” nature of many AI systems, where even its own developers may not fully understand model outputs (Burrell, 2016; Fleisher, 2022).

If markets fail to distinguish real AI adoption from empty, promotional language, capital may be misallocated to firms skilled in storytelling rather than operational substance (Cooper et al., 2001; Shiller, 2000). This behaviour reflects earlier periods of mispricing, such as during the dot-com era and the blockchain boom, when firms received valuation premiums for their use of fashionable terminology (Cooper et al., 2001; Jain & Jain, 2019). As a result, the current AI enthusiasm could risk distorting the efficient allocation of capital, driven by investor overreaction (Barberis et al., 1998).

These dynamics raise the following important question: Do AI-related corporate disclosures signal real technological progress, or are they just attention-seeking statements with no real substance?

The Measurement Challenge: From Basic Word Counts to Language Models That Think

Annual reports, especially 10-K filings, have become longer and more complex (Lesmy et al., 2019), with a rising number of AI-related references (Accenture, 2024). This makes manual analysis impossible at scale (Loughran & McDonald, 2016; Bojić et al., 2025). Early Natural Language Processing (NLP) approaches, such as keyword counting and bag-of-words methods, lacked the ability to detect meaning or nuance (Juluru et al., 2021; Loughran & McDonald, 2011). Even more advanced tools like FinBERT had difficulty with irony, sarcasm, and shifting terminology (Araci, 2019; Weitzel et al., 2016; Alexandris, 2024). LLMs, such as GPT-4, provide a more subtle understanding of the context in financial texts, therefore offering improved analysis. (Barrios et al., 2025; Li et al., 2023). However, concerns remain about their transparency and the potential for fabricated or biased outputs (Burrell, 2016; Prabhakar, 2025; Bommasani et al., 2021). To address known limitations of LLMs in financial analysis, including systematic biases documented by Chen et al. (2025), this study employs categorical scoring, ensemble methods, and structured prompt engineering to ensure reliable measurement.

This challenge resembles efforts to detect greenwashing in Environmental, Social, and Governance (ESG) reports (Seele & Schultz, 2022; Moodaley & Telukdarie, 2023). If markets are to reflect available information accurately, tools that help distinguish substance from exaggeration are crucial (Fama, 1970).

Research Agenda and Questions

This study investigates the relationship between corporate communication about AI and subsequent stock market performance. The main research question is:

Do companies that communicate more clearly or positively about AI in their 10-K reports receive superior stock market returns, or do investors discount such statements if not backed by clear evidence of action?

This core question is supported by two secondary inquiries.

First, *which aspects of AI-related disclosure matter most to investors?* This includes elements like tone, strategic depth, risk language, future focus, and specific investment information. The analysis not only looks at whether AI is mentioned, but also at how it is discussed. Research shows that narrative features such as tone and clarity affect investor reactions (Feldman et al., 2009; Loughran & McDonald, 2011), and that changes in disclosure sentiment over time provide insights into future company performance (Brown & Tucker, 2010). Specifically, disclosures that indicate credible investment or future strategy

may be more significant than vague or overly positive mentions. On the other hand, disclosures that point out risk can either lower perceived value or boost credibility, depending on the context (Kravet & Muslu, 2013; Hope et al., 2016).

Second, *does the industry context affect how AI disclosures are understood?* This study tests whether informative AI communication depends on industry norms and expectations. In technology sectors, using AI might be seen as standard and may not draw much market attention. However, when non-technology firms announce AI adoption with a clear purpose, it might signal real innovation or change. Previous research shows that the economic value of AI investments relies on industry-specific skills and complementary assets (Cockburn et al., 2018; Brynjolfsson et al., 2021). Additionally, asset pricing studies suggest comparing firms within their industry to better estimate disclosure effects (Asness et al., 2000).

These questions test whether financial markets effectively process complex, qualitative information. The study examines the semi-strong form of the Efficient Market Hypothesis (Fama, 1970; Malkiel, 2003), which suggests that all publicly available information should quickly be reflected in prices. If markets efficiently process AI disclosure, systematic relationships between disclosure quality and returns should reflect underlying firm fundamentals rather than narrative effects. LLMs are used to assess this language.

To test these hypotheses, the study uses LLMs to measure six aspects of AI-related disclosure: strategic depth, sentiment, risk framing, forward-looking statements, investment focus, and signs of AI washing (Seele & Schultz, 2022; Al Haddi, 2024). These scores are then connected to future stock returns using panel regressions with standard controls, including time- and industry-fixed effects. This method is based on well-known research showing that text features in regulatory filings predict both investor perception and market outcomes (Loughran & McDonald, 2011; Feldman et al., 2009; Harris, 2024).

This study examines whether markets reward firms for clear AI disclosures, or are influenced by persuasive yet unsupported narratives. If disclosure quality is properly reflected in prices, stock returns should not systematically vary based on language. However, if the tone, structure, or focus predicts market performance, it would raise concerns about the impact of corporate storytelling on capital distribution.

Data and Methodology Overview

This study investigates the relationship between AI-related disclosures and future stock returns. It uses a sample of 799 firms from the Russell 3000 index, covering the years 2020 to 2024. This index covers about 98% of the investable U.S. equity market and includes only

actively traded companies (FTSE Russell, 2025). This timeframe captures a period of rapid growth in corporate AI engagement (Accenture, 2024), as seen by survey data showing that the regular use of generative AI in organisations nearly doubled in less than a year (McKinsey, 2024). This period allows for the examination of how markets react to different types of AI stories (Mahat et al., 2024). Given market shocks and macroeconomic uncertainty during this period (Grigoryev, 2023), the models incorporate control variables to help distinguish AI signals from broader market noise.

The core of the analysis is the textual content of 10-K filings. Two parts of each company's 10-K filing are used:

Item 1A: Risk Factors. This section requires firms to disclose important risks, including those associated with AI. Previous studies have shown that risk disclosures in “Item 1A” give useful information about specific threats to firms and relate to how investors view risk (Kravet & Muslu, 2013; Beatty et al., 2018).

Item 7: Management's Discussion & Analysis (MD&A). In this section management reflects on its past performance and future goals. It frequently contains references to AI-related plans, Research and Development (R&D) priorities, and operational transformations. MD&A language has been shown to affect investor expectations and capture narrative shifts not reflected in quantitative data (Feldman et al., 2009; Brown & Tucker, 2010; Muslu et al., 2015).

An ensemble approach (averaging the outputs of the three models), is used to reduce model-specific bias and improve predictive accuracy (Breiman, 1996; Seni & Elder, 2010). This multi-model approach lowers the risk of single-model bias, similar to how inter-coder reliability is established in manual content analysis to ensure coding consistency across human raters (Artstein & Poesio, 2008).

Also, the analysis uses advanced prompt engineering techniques, including role assignment, Chain-of-Thought (CoT) reasoning, and structured output validation, to ensure reliable extraction of structured information from corporate narratives (OpenAI, 2025; Wei et al., 2022; Dunivin, 2025).

Six dimensions measure AI-related disclosure in 10-K filings. This approach is based on prior research in financial text analysis, technological innovation, and impression management. Each factor is measured using ranked scores that allow for comparison across firms and time. These categories serve as the basis for analysing how qualitative AI communication relates to stock returns.

1. AI Strategic Depth

Measures AI's centrality to business strategy. Research indicates that firms focusing on technological integration and skill development often show greater innovation and long-term success (Cockburn et al., 2018; Mishra et al., 2022; Babina et al., 2024).

2. AI Disclosure Sentiment

Captures the overall tone of AI discussion, from optimistic to cautious. The tone of corporate language has been shown to influence investor behaviour and market reactions (Loughran & McDonald, 2011; Feldman et al., 2009). Positive language may boost short-term sentiment, while negative tone can signal risk (Kearney & Liu, 2014).

3. AI Risk Disclosure (*includes three sub-metrics*)

Measures firm references to risks related to:

(a) Own Adoption (implementation risk),

- Firms that reveal internal risks, like integration issues, legal risks, or technical limits, indicate possible operational problems but also show some transparency in handling AI adoption (Alston & Bird, 2024; Amershi et al., 2019).

(b) External Threats (AI used by competitors or third parties)

- Disclosures that recognise competitive threats from external AI developments show an understanding of industry changes, in line with market shifts driven by innovation (Christensen, 1997; Cockburn et al., 2018).

(c) Non-Adoption (the risk of falling behind due to non-adoption)

- When companies mention the risk of falling behind due to slow or not adopting AI, they stress the need to invest in technology to stay competitive in changing markets (Hava, 2025; Brynjolfsson & McAfee, 2014).

Such risk disclosures have been shown to influence investor interpretation and reduce information asymmetry (Beatty et al., 2018; Hope et al., 2016)

4. Forward-Looking AI Statements

This score assesses language about future AI investments, product plans, or strategic developments. Forward-looking MD&A content is known to shape

investor expectations (Muslu et al., 2015), but vague use of future-tense language may be viewed with scepticism (Düger, 2020). Clear and specific plans are more likely to be seen as valuable signals (Barrios et al., 2025).

5. AI Washing Index

Assesses whether AI is presented using vague, promotional language without evidence of action. This measure draws on research in greenwashing and dot-com-era hype cycles (Cooper et al., 2001; Seele & Schultz, 2022; Al Haddi, 2024).

6. AI Talent & Investment Focus

References to AI-specific hiring, funding, or partnerships. Prior work shows that disclosures about R&D and human capital investment are associated with innovation and signalling firm commitment (Babina et al., 2024; Spence, 1973).

These scores are then used as independent variables in regression models. The aim is to test whether they help predict stock returns after the filing date. The primary dependent variable is excess returns, calculated at 3, 6, 9, and 12 months after the 10-K date. Returns are adjusted for the risk-free rate to capture true outperformance.

To ensure the effect of AI disclosure is not confounded with other known drivers of stock returns, the models include a complete set of control variables. These cover firm size, valuation (like price-to-book), profitability, and standard risk factor loadings from the Fama-French-Carhart models (Fama & French, 1993; Carhart, 1997).

The analysis tests whether these AI narratives provide information beyond classic financial metrics (Loughran & McDonald, 2016).

To explain the patterns observed, the study combines statistical analysis with detailed case studies (Creswell & Plano Clark, 2018). Case studies examine companies with the highest AI disclosure scores and worst stock returns. These cases show how management pressures, misleading signals, and company self-selection create the observed patterns (Jensen & Meckling, 1976; Spence, 1973; Roy, 1951).

Key Findings Preview

The results reveal a counterintuitive pattern: firms with stronger AI disclosures tend to experience lower returns, with effects concentrated in traditional industries rather than technology sectors. Initially, it might seem that companies with detailed or confident AI talk in their 10-K filings would do better. However, the data do not support this assumption (Barrios et al., 2025). In fact, between 2020 and 2024, most aspects of AI disclosure, such as

“Strategic Depth”, “Disclosure Sentiment”, and “Forward-Looking” statements, are negatively linked to future performance.

Three theoretical mechanisms explain this pattern. First, is **self-selection bias** where firms facing operational difficulties use extensive AI narratives to mask underlying weaknesses (Roy, 1951; Spence, 1973). Case studies of companies like Omnicell and 3D Systems show extensive AI disclosure coinciding with earnings disappointments and insider selling, suggesting defensive impression management rather than genuine strategy communication (Jensen & Meckling, 1976).

Second, **information asymmetry** problems arise because AI capabilities depend on unverifiable intangible assets like proprietary algorithms and specialised staff (Akerlof, 1970; Greene, 2025). When firms cannot credibly demonstrate real AI integration, markets rationally discount all AI claims, creating adverse selection where only struggling firms engage in extensive disclosure.

Third, **sector-specific legitimacy concerns** where AI adoption signals potential threats rather than opportunities (Suchman, 1995). The strongest negative effects occur in the Consumer Discretionary, Healthcare, and Industrials sectors. In these sectors AI raises concerns about job displacement and operational risks rather than competitive advantages.

The pattern reflects market sophistication, not irrationality. Investors correctly identify that extensive AI disclosure correlates with business vulnerability rather than strategic strength, validating concerns about narrative-driven capital misallocation during technology hype cycles (Cooper et al., 2001; Fama, 1970).

Contributions and Implications

This study contributes to finance, accounting, and information systems research in several ways. First, it shows that LLMs can convert complex narrative language in 10-Ks into structured data. Rather than counting keywords, the LLM method scores six dimensions of AI disclosure using prompts tested across multiple models for reliability.

Second, the framework can be adapted for other topics like ESG or cybersecurity, where narrative analysis matters, helping extract insights from large volumes of unstructured text.

Third, the empirical results show that firms with stronger AI narratives often have lower future returns, challenging assumptions that markets reward AI-focused statements. Instead, investors may discount excessive or vague language lacking evidence of real

capability. This aligns with signalling theory and self-selection bias (Spence, 1973; Jacobs et al., 2009).

For investors, frequent or overly positive AI statements might signal impression management rather than real progress. The scoring method may help identify inflated disclosures.

For companies, more AI talk is not always better. General optimism without specific details can backfire. As SEC Director Gurbir Grewal (2024) warned, exaggerated AI claims can reduce investor trust (U.S. Securities and Exchange Commission, 2024). Honest, detailed disclosures linked to real investments support long-term value better.

Policy Implications: The findings suggest markets can differentiate between genuine AI developments and hype, with consistent negative reactions to excessive disclosure in traditional sectors. This supports lighter regulatory approaches that aim to prevent obvious AI washing. It suggests that strict disclosure rules are unnecessary. The market's subtle reaction indicates investors understand complex AI information without needing heavy regulation.

These results show market efficiency in processing complex narrative information (Fama, 1970).

2. Literature Review and Hypotheses Development

This section reviews research on narrative disclosure analysis, LLM applications in financial text analysis, and market reactions to AI-related corporate communications. It also identifies gaps that motivate the study's hypotheses.

Narrative Disclosures and Market Implications

Research shows that markets can be influenced by psychological factors and narratives rather than fundamental analysis (Shiller, 2000). The MD&A and Risk Factors sections of 10-K filings contain qualitative, non-financial information. Early studies find only modest market reactions to 10-K filings once earnings information is known (Easton & Zmijewski, 1993). However, recent evidence shows textual content provides valuable information about firm prospects that is reflected in stock prices (You & Zhang, 2009).

Key informative textual properties include narrative tone, readability, and the extent of year-over-year modifications (Feldman et al., 2009; Li, 2008; Brown & Tucker, 2010). MD&A sentiment affects investor responses, with tone changes providing information beyond standard financial measures like earnings surprises (Feldman et al., 2009). Negative

words are associated with lower stock returns, confirming that textual sentiment is priced by the market (Loughran & McDonald, 2011).

Language clarity also affects market responses. Simple language signals transparency, while complex language may indicate firms hiding poor performance (Li, 2008; Loughran & McDonald, 2014; Garel et al., 2019). You and Zhang (2009) show investors underreact to complex 10-K filings, suggesting language complexity affects information processing.

Firms with larger year-over-year MD&A modifications experience stronger stock reactions, though these changes do not predict analyst forecast revisions (Brown & Tucker, 2010). This indicates that professional analysts may not find the same value in these textual changes. Companies with minimal textual changes exhibit weaker market reactions, suggesting investors ignore repetitive content but price new information (Cohen et al., 2020).

The Risk Factors section (Item 1A) became mandatory in 2005 to summarise management's view of significant negative factors (Nelson & Pritchard, 2016). Kravet and Muslu (2013) find that increases in risk sentences, associate with higher volatility and trading volume. Hope et al. (2016) show specific risk factors generate stronger market reactions and improve analyst risk assessment.

This evidence supports examining the MD&A and Risk Factors sections in this study. The documented market reactions to sentiment, readability, textual changes, and risk disclosures suggest that these narrative sections contain useful information that is not immediately reflected in stock prices. This delayed price incorporation challenges the strong form of market efficiency and shows that analysis of qualitative disclosures can reveal information advantages (Fama, 1970; Malkiel, 2003).

Advanced NLP by LLMs

The use of LLMs for textual analysis is motivated by their advantages over traditional NLP models, such as FinBERT. These models offer advantages in scalability and their effective adaptability (Bommasani et al., 2021). Advanced LLMs can process large amounts of textual data from "Risk Factors" and "Management's Discussion and Analysis" sections across thousands of 10-K filings with considerable speed and consistency. This scalability enables analysis of entire market segments that would be impractical with manual coding or traditional computational methods (Bojić et al., 2025). Compared to humans, LLMs keep consistent evaluation criteria across large datasets without fatigue or subjective drift (Wang & Wang, 2025).

Furthermore, LLMs offer superior text understanding compared to traditional methods that rely on keyword counting or basic sentiment scoring (Zhang et al., 2023; Gupta et al., 2024). LLMs can identify and interpret complex language that characterises corporate AI engagement (Thapa et al., 2025). This ability is important for distinguishing between companies with genuine AI integration and those making superficial mentions (Barrios et al., 2025; Dunivin, 2025). Additionally, LLMs can analyse several disclosure aspects at once within a single framework. They detect linguistic patterns missed by traditional methods (Wang et al., 2025).

Finally, while LLMs offer clear advantages, they also come with known biases and limitations. While they show strong language abilities, they still encounter challenges in understanding subtle wording or jargon (Weitzel et al., 2016; Bojić et al., 2025). They can also reflect and amplify biases, such as excessive optimism or following trends in stock predictions (Chen et al., 2025), and may yield inconsistent outcomes when given slightly different prompts (Wei et al., 2022). Since their output is greatly influenced by how prompts design (Wei et al., 2022; Dunivin, 2025), best practice involves using clear, example-based instructions and fixed scoring systems to minimise variation (Dhinakaran & Jolley, 2024; Anthropic, 2025). This study will implement these optimal prompt engineering techniques. To address bias and enhance consistency, this study combines results from three different LLMs. This approach utilises various model perspectives to increase reliability (Bojić et al., 2025).

AI/Textual Factors and Financial Performance

Within the 10-K narratives, several dimensions of AI-related disclosure are created that could possibly affect investor expectations and, therefore, stock returns. AI disclosure informativeness depends on firm characteristics relative to industry peers, as markets interpret the same disclosure differently based on company context (Barrios et al., 2025).

Research indicates that AI reporting practices vary significantly across sectors, with technology firms showing more frequent AI-related disclosures than traditional industries (Bonsón et al., 2021). Each dimension is grounded in prior literature on how textual or thematic elements relate to firm performance (Babina et al., 2024):

1. AI Strategic Depth:

AI Strategic Depth refers to how much a company's 10-K report highlights AI as a central part of its business strategy rather than just a surface-level mention.

A strong strategic focus on AI may signal a firm's commitment to developing technological capabilities for a sustainable competitive advantage. (Cockburn et al., 2018). This strategic focus matches evidence showing that companies making large investments in new technologies, like AI-driven digital change, often see improved long-term performance (Mishra et al., 2022; Cui, 2025).

Research shows that companies making long-term AI investments in hiring and projects tend to achieve higher growth and innovation (Babina et al., 2024). However, this positive relationship follows the "Productivity J-Curve" effect described by Brynjolfsson et al. (2021).

This theory explains why companies adopting new technologies like AI often experience an initial decline in productivity and performance before the long-term benefits become evident. The decline occurs because companies must invest significantly in intangible assets. These include training, process redesign, and organisational change, which are hard to measure but necessary for adopting AI successfully.

Furthermore, Mishra et al. (2022) show that a strong AI focus in 10-K reports, which indicates management's emphasis on AI initiatives, is associated with improved operating efficiency and profitability. Thus, by integrating AI into their core business, companies could hint at future success. The market might reward these organisations, seeing them as leaders in AI.

2. Disclosure Sentiment

Disclosure Sentiment is defined by the tone or sentiment surrounding AI discussions that may influence investors' interpretations. If management speaks about AI in optimistic terms (opportunity, growth, competitive strength) versus in cautious or pessimistic terms (concerns, costs, uncertainties), it could impact stock expectations. Indeed, extensive evidence suggests that the sentiment expressed in corporate disclosures correlates with market reactions and future returns (Kearney & Liu, 2014). Positive language in narratives often accompanies higher immediate returns or improved investor sentiment, whereas negative language can predict stock price declines or signal risk (Loughran & McDonald, 2011; Kearney & Liu, 2014).

In the context of AI, a positive tone may signal management's confidence in AI-related projects, whereas a negative tone could reflect operational or strategic challenges (Kearney & Liu, 2014). Consistent with general findings that textual tone in MD&A and earnings calls affect investor behaviour (Henry, 2008; Feldman et al., 2009), it is anticipated that the sentiment of AI disclosures (measured, for example, via a financial sentiment model) will

relate to stock performance. However, sentiment alone may be less informative than substantive content; therefore, it is tested alongside other aspects of AI disclosure.

3. *AI-Related Risk Disclosures*

In discussions about AI-related risks, several dimensions are considered:

(3a) Implementation Risk (Own AI Adoption)

Adopting new technologies like AI involves significant implementation risks, such as integration problems, cybersecurity threats, and regulatory issues (Alston & Bird, 2024; Amershi et al., 2019). These obstacles can reduce productivity and lead to project failure if not effectively considered (Brynjolfsson et al., 2021)

A company's open disclosure of AI implementation challenges might warn investors of potential setbacks, possibly lowering short-term expectations. However, it can also signal transparency and proactive management of AI-related threats, which means that firms are managing these risks responsibly.

(3b) External AI Threats

Next, companies may describe risks from AI developments in the external environment. External risks include the risk of industry disruption by AI, new AI-enabled competitors, or shifts in customer behaviour due to AI. Such disclosures reflect the concept of technological disruption risk, where companies risk becoming outdated if they fail to keep pace with technological advancements (Christensen, 1997). As discussed in Chapter 1, general-purpose technologies like AI are known to reshape industries (Cockburn et al., 2018). Firms that do not adopt AI could lose market share, as they risk being outcompeted by more agile competitors, a principle of technological disruption theory (Christensen, 1997). When a company warns about external AI threats, it indicates that management is aware of potential competitive disruption. Management awareness could mean two things: it might make investors cautious about the firm's future, but it also shows the firm is conscious of its situation.

(3c) Competitive Risk of Non-Adoption

Additionally, there is a risk of falling behind if AI is not adopted quickly. Companies increasingly recognise that not using AI could mean lagging behind competitors who do, making AI adoption essential rather than optional (Barrios et al., 2025; Hava, 2025). This type of competitive risk acts as the opposite of external threats. Economic theory and evidence suggest that firms investing in emerging technologies tend to achieve superior long-term results, while those slow to adopt technology are more likely to underperform

(Brynjolfsson & McAfee, 2014; Christensen, 1997). Furthermore, recent findings indicate the existence of an “AI premium,” whereby firms with higher levels of AI investment experience better stock performance, suggesting that capital markets reward AI capability (Yang et al., 2024). Therefore, when firms mention the risk of falling behind in AI developments, it shows they see technology adoption as important for staying competitive in the future.

In summary, AI-related risk disclosures, whether about implementation problems or competitive threats, provide useful information. While markets may initially respond negatively due to the uncertainty these risks suggest (Beatty et al., 2018), detailed and specific disclosures over time show that firms are actively dealing with potential challenges. Studies indicate that such disclosures are linked to lower information asymmetry and a reduced cost of capital (Campbell et al., 2018; Hope et al., 2016).

4. Forward-Looking AI Statements

Forward-looking AI statements reflect how firms discuss expected investments, product launches, and long-term strategic plans involving AI. Research indicates that companies with stock prices that poorly reflect future earnings information are more likely to issue forward-looking disclosures in MD&A sections (Muslu et al., 2015). Also, excessive use of future-tense terms such as “will” or “shall” is associated with lower future stock returns, suggesting that markets may interpret such language as overly optimistic or lacking in substance (Düger, 2020). This suggests that companies prioritise talk over taking action.

Confirming that firms use forward-looking AI statements to describe their planned investments, product developments, and long-term AI strategies. If these statements are seen as credible and informative, they are linked to higher stock prices (Muslu et al., 2015; Barrios et al., 2025). However, if they lack substance, they can be viewed negatively (Düger, 2020).

5. “AI-Washing” Index (Hype vs. Substance)

AI washing represents a form of impression management where firms exaggerate or make unsubstantiated claims about their AI capabilities to appear more technologically advanced than they actually are. The concept is based on greenwashing literature, where companies overstate their environmental commitments (Delmas & Burbano, 2011). Seele and Schultz (2022) introduced the term “machinewashing” to describe this phenomenon, defining it as the deliberate misrepresentation of AI and machine learning capabilities to stakeholders. This practice allows firms to profit from the AI hype while avoiding the substantial investments required for genuine AI implementation.

The regulatory environment shows increasing concern about these practices. The SEC recently charged two investment advisers for falsely claiming to use AI and has warned that exaggerated AI claims may mislead investors and attract regulatory attention (U.S. Securities and Exchange Commission, 2024; Uslaner & Coquin, 2025). The use of buzzwords can lead to market mispricings, as seen during the dot-com bubble when firms adding ".com" to their names experienced unjustified stock price jumps not explained by fundamentals (Cooper et al., 2001).

This AI-washing index identifies firms that frequently mention AI-related terms without providing evidence of real engagement. Companies with high-hype, low-substance AI claims may experience short positive market reactions that fade when expectations aren't met, leading to underperformance. Meanwhile, firms providing clear, low-hype disclosures backed by actual AI actions may achieve more lasting stock performance gains.

6. AI Talent & Investment Focus:

The final metric captures how firms describe investments in AI talent, infrastructure, and capabilities. Actions such as hiring data scientists, acquiring AI startups, or funding AI-related projects often indicate serious efforts toward implementation (Babina et al., 2024). Research shows that human capital and research and development (R&D) investment are essential to realise value from new technologies. Babina et al. (2024) show that AI-related hiring is associated with increased innovation and firm growth. These disclosures may show more than simple interest. According to signalling theory, visible and costly actions, like hiring specialised talent, act as valid signs of a firm's quality and long-term potential (Spence, 1973).

Collectively, the study focuses on six AI disclosure dimensions: strategic depth, tone, risk categories, forward-looking content, the balance between hype and substance, and talent and investment. It evaluates them through 10-K text analysis and investigates their individual and combined relationships with future stock returns.

Case Study Applications in Corporate Disclosure Research

Case studies help explain complex disclosure patterns that statistics alone cannot capture (Yin, 2018). Combining statistical analysis with detailed company examples have proven effective in finance research for understanding market reactions (Eisenhardt, 1989).

Recent studies show the value of this approach. Beattie and Smith (2013) use statistics followed by case studies to explain how graphics in reports affect firm value. Linsley and Shrives (2006) combine numbers with specific company examples to understand disclosure

strategies. Similarly, Linsley and Shrives (2006) combine quantitative risk disclosure analysis with specific company examples to understand different disclosure strategies.

This approach works particularly well when markets react in unexpected ways to company disclosures, helping researchers identify the behavioural reasons behind statistical patterns (Patton, 2015).

Hypotheses Development

Building on this literature, three key hypotheses are created about AI disclosures and stock returns:

H1 (AI Disclosure and Returns): *Firms with higher quality and extent of AI-related disclosures in their 10-K will experience better future stock returns on average than those with minimal or poor-quality AI disclosures.*

This hypothesis builds on research showing that markets can underreact to complex textual information, (You & Zhang, 2009) and that narrative disclosures (like future statements or detailed reports) contain information not immediately reflected in stock prices (Brown & Tucker, 2010; Muslu et al., 2015).

H2 (Strategic vs. Sentiment Dimensions): *Strategic and talent/investment-focused AI disclosure dimensions will be more predictive of returns (and future firm performance) than sentiment or hype-related measures.*

Research supports this hypothesis by showing that disclosures focused on implementation, like AI hiring, infrastructure investment, or detailed strategies, provide more information than vague, sentiment-heavy language (Babina et al., 2024; Mishra et al., 2022; Spence, 1973). This suggests that, tangible evidence of implementation is a more credible signal for long-term value than a positive tone or excessive buzzwords. Objective signs of AI adoption may lead to a return premium, while hype-driven sentiment might only result in temporary valuation increases (Düger, 2020; Seele & Schultz, 2022).

H3 (Sectoral Differences): *The impact of AI disclosures on returns will be more substantial for firms outside the traditional tech sector.*

Prior work shows that firm-level signals, such as valuation or profitability, are more informative when considered relative to industry peers (Asness et al., 2000). Bonsón et al. (2021) find that AI reporting is already common in the technology sector, which reduces differences and makes these disclosures less surprising to investors. In contrast, disclosures from non-tech firms are less common and may serve as stronger signals of innovation. Other research confirms that AI adoption impacts sectors in different ways. This impact is stronger

in industries where companies must make major organisational changes (Brynjolfsson et al., 2021; Cockburn et al., 2018). This indicates that AI disclosures provide more valuable information in industries where organisational changes are rare.

Research Gaps and Contributions

This study contributes to corporate disclosure, AI economics, and natural language processing research in several ways. First, it provides the first large-scale analysis of AI-related disclosures.

This extends previous research that used keyword frequency, by analysing full texts with LLMs to evaluate different aspects of AI disclosure (Mishra et al., 2022; Barrios et al., 2025). It then examines the relationship between these disclosure dimensions and stock returns.

Methodologically, the study introduces structured LLM prompts to evaluate AI disclosure quality in context. Instead of relying on word counts or sentiment analysis, the models act like intelligent readers that assess whether a firm's AI communication is meaningful, future-oriented, or exaggerated (Dunivin, 2025). This approach builds on findings that disclosure detail and framing provide more insight than tone or frequency alone (Cao et al., 2024). The ensemble methodology combining three LLMs (OpenAI's GPT-4o-Mini, Gemini Flash 1.5, and Gemini Flash 2.5) reduces model-specific bias and advances text-based financial research (Bojić et al., 2025).

Empirically, the study challenges conventional signalling theory by documenting systematic negative relationships between AI disclosure intensity and future returns. This reveals that extensive AI communication signals business vulnerability rather than strategic strength, particularly in traditional industries. The findings indicate market efficiency in processing complex narrative information, contradicting views on mispricing influenced by narratives.

The AI-Washing Index distinguishes real engagement from hyped claims by examining consistency between disclosure and observable actions like AI hiring or R&D spending. This addresses regulatory concerns about overstated AI claims while showing markets can identify these inconsistencies (U.S. Securities and Exchange Commission, 2024; Uslaner & Coquin, 2025).

Theoretically, the study extends signalling theory by showing disclosure can signal weakness when firms venture beyond core competencies, validates information asymmetry theory in AI contexts, and supports legitimacy theory explanations for sector-specific market

reactions (Spence, 1973; Akerlof, 1970; Suchman, 1995). The results provide evidence of market sophistication in processing qualitative information rather than systematic mispricing.

3. Data

3.1 Sample Selection and Filtering Process

The sample construction follows a sequential filtering process to maintain data quality and methodological consistency (Regulwar et al., 2023). As explained in the Introduction, the study focuses on corporate AI disclosures over 2020-2024 to analyse the recent increase in AI engagement (Accenture, 2024; McKinsey, 2024). The initial sample included all 2,629 constituents of the iShares Russell 3000 index (BlackRock, 2025; FTSE Russell, 2025). After applying a four-stage filtering process, the final sample consists of 799 firms (see Table 1).

Stage 1: Balanced Panel Requirement. To create a balanced panel and ensure consistency over time, only firms with complete 10-K filings available for all five fiscal years (2020-2024) are retained. This reduces the sample from 2,629 to 1,635 firms, eliminating companies that delisted, merged, or had incomplete filing histories during the study period.

Stage 2: Section Extraction Success. A successful extraction of both Item 1A (Risk Factors) and Item 7 (Management's Discussion & Analysis) is required for each filing. These sections serve as the primary textual input for the LLM analysis (Nelson & Pritchard, 2016; Brown & Tucker, 2010; Feldman et al., 2009; Kravet & Muslu, 2013; Harris, 2024). Firms missing either section in any year are excluded, yielding 1,445 firms.

Stage 3: Technical Processing Constraint. To ensure reliable, untruncated analysis by LLMs, firms whose combined Risk Factors and MD&A text exceeds 200,000 characters in any given year are excluded (OpenAI, 2023). This threshold represents the upper limit for consistent LLM processing given current model context sizes. While this filter affects only a small number of firms with exceptionally lengthy filings (primarily large conglomerates), it is necessary to maintain methodological consistency across all observations. This reduces the sample to 828 firms.

Stage 4: Return Data Availability. Firms with insufficient trading history or missing prices during the study period are excluded. This results in a final sample of 799 unique firms corresponding to 3,995 firm-year observations.

The requirement for complete data across all five years introduces survivorship bias, favouring more stable, established firms that remained listed and filed regularly throughout

the period (Brown et al., 1992; Carhart et al., 2002). Although this reduces the ability to apply findings to smaller or less stable firms, it improves the reliability of the panel analysis by keeping a consistent group of firms over time. This eliminates confounding effects that could arise from a changing sample composition (Aptech, 2019).

Table 1:
Sample Construction and Filtering Process, 2020–2024

Stage	Description	Companies	Firm-Years
Russell 3000 Universe	iShares ETF constituents	2,629	13,145
Complete 10-K	All 5 – Years (2020-2024)	1,635	8,175
Section Extraction Success	Valid Risk Factors & MD&A	1,445	7,225
Text Length Sample	<200k characters combined	828	4,140
Final Sample	With returns data	799	3,995

Note: This table details the sequential filtering process applied to Russell 3000 constituents to construct the final sample for fiscal years 2020–2024. "Firm-Years" represents the number of firm-by-fiscal-year observations. The final sample consists of 799 unique firms and 3,995 firm-year observations.

3.2 Data Sources and Collection

This study collects the data through two different data sources: (i) raw 10-K text directly from the SEC and (ii) numerical series accessed through WRDS

- **Textual data.** Annual 10-K filings are downloaded through the SEC’s EDGAR API. For each filing, *Item 1A: Risk Factors* and *Item 7: MD&A* are programmatically extracted. When amended filings (10-K/A) existed, the most recent version is retained to ensure reliance on the final information set available to investors (Nelson & Pritchard, 2016).
- **Market and accounting data.**
 - *Center for Research in Security Prices (CRSP).* Provided daily stock returns used to construct the study’s dependent variables.
 - *Compustat North America Fundamentals Quarterly.* Supplied firm-level accounting items (e.g., book value of equity, total assets) for traditional control variables.
 - *Kenneth French’s Data Library.* Delivered daily factor returns for the Fama–French–Carhart models (Fama & French, 1993, 2015; Carhart, 1997) and the

daily risk-free rate, necessary for calculating excess returns and estimating factor exposures.

Separating textual data (EDGAR) from numerical data (CRSP/Compustat/French Library via WRDS) clarifies the research design by distinguishing raw inputs from the analysis that follows.

3.3 Sample Characteristics

Table 2 shows the industry composition, categorised according to the Global Industry Classification Standard (GICS). The sample represents a broad cross-section of sectors, with Industrials (20.7%), Consumer Discretionary (15.4%), and Information Technology (13.6%) forming the three largest groups.

Comparing the final sample to the full Russell 3000 index reveals some compositional differences. The sample slightly overrepresents Industrials (+5.7 percentage points) and Consumer Discretionary (+4.4 percentage points), while underrepresenting Financials (-5.3 percentage points) and Health Care (-8.4 percentage points). Firms in highly regulated sectors (Financials, Health Care) likely may have more complex or longer reports that are removed during text processing. However, all 11 GICS sectors remain well represented, ensuring that the results are not dominated by one industry (Harris, 2024).

Table 2:
Sector Distribution of the Final Sample versus the Russell 3000 Index

Sector	Companies	Sample (%)	Cumulative %	Russell 3000 (%)
Industrials	165	20.7%	20.7%	15.0%
Consumer Discretionary	123	15.4%	36.0%	11.0%
Information Technology	109	13.6%	49.7%	12.0%
Financials	101	12.6%	62.3%	17.9%
Health Care	70	8.8%	71.1%	17.2%
Real Estate	62	7.8%	78.8%	6.1%
Consumer Staples	47	5.9%	84.7%	4.2%
Materials	46	5.8%	90.5%	4.4%
Energy	33	4.1%	94.6%	4.8%
Communication	26	3.3%	97.9%	4.5%

Sector	Companies	Sample (%)	Cumulative %	Russell 3000 (%)
Utilities	17	2.1%	100.0%	2.5%

Note: This table compares the Global Industry Classification Standard (GICS) sector distribution of the final sample (N=799 firms) against the Russell 3000 index. "Sample (%)" is the percentage of firms per sector in the final sample. "Russell 3000 (%)" is the index composition as of year-end 2024. The sample period is 2020–2024.

The requirement for complete and manageable filings over five years skews the sample towards larger, more established firms. This is because such firms are more likely to meet these criteria. However, the broad sector coverage offers a reliable basis for generalisable conclusions about how capital markets interpret corporate AI narratives (Kravet & Muslu, 2013).

3.4 Descriptive Statistics

3.4.1 Dependent Variable Properties

Table 3:
Descriptive Statistics for Buy-and-Hold Excess Returns

Variable	N	Mean	Std	Min	P25	Median	P75	Max
3m Excess Returns	3325	0.0003	0.185	-0.465	-0.116	-0.008	0.106	0.625
6m Excess Returns	3270	0.0235	0.2541	-0.575	-0.129	0.003	0.156	0.975
9m Excess Returns	3220	0.0451	0.3399	-0.667	-0.178	0.016	0.215	1.389
12m Excess Returns	3174	0.0412	0.376	-0.73	-0.191	0.001	0.213	1.672

Note: This table presents summary statistics for the dependent variables: buy-and-hold excess returns, calculated over 3-, 6-, 9-, and 12-month horizons following the 10-K filing date. Returns are winsorized at the 1% and 99% levels. N represents the number of firm-year observations for the 2020–2024 sample period.

Table 3 shows expected patterns across investment horizons. Mean excess returns increase from essentially zero at 3 months to approximately 4% over 12 months, reflecting equity risk premium (Damodaran, 2025). Return volatility similarly increases from 18.5% to 37.6%, consistent with risk accumulation over time (Fama & French, 1993). The wide range of outcomes (-73% to +167% for 12-month returns) provides necessary variation for cross-sectional analysis.

3.4.2 Control Variables Properties

Table 4:
Descriptive Statistics for Regression Control Variables

Variable	N	Mean	Std	Min	P25	Median	P75	Max
Panel A: Book Control Variables								
Log Market Cap	3989	22.1377	1.6426	18.908	20.935	22.015	23.235	26.405
Price-to-Book	3985	3.9108	4.0677	0.347	1.338	2.352	4.66	16.546
ROA	3984	0.0122	0.0236	-0.08	0.002	0.011	0.022	0.087
Panel B: Factor Beta Controls (12-Month)								
Market Beta (MKTRF)	3843	0.9708	0.4706	-0.587	0.75	0.983	1.223	2.124
Size Beta (SMB)	3843	0.6748	0.6995	-1.125	0.227	0.649	1.102	2.448
Value Beta (HML)	3843	0.1546	0.7864	-2.027	-0.261	0.13	0.554	2.517
Profitability Beta (RMW)	3843	0.1418	0.8669	-2.622	-0.236	0.197	0.599	2.445
Investment Beta (CMA)	3843	0.0365	1.1002	-3.212	-0.522	0.05	0.619	3.191
Momentum Beta (UMD)	3843	-0.0296	0.615	-1.879	-0.322	-0.014	0.307	1.608

Note. This table presents summary statistics for the control variables used in the regression analyses for the 2020–2024 sample. All variables are winsorized. N is the number of firm-year observations. Panel A reports book controls from Compustat. Panel B reports estimated factor loadings from the Fama-French-Carhart six-factor model for the 12-month horizon.

Table 4 shows sensible characteristics for both book and factor controls. The dynamic factor loading procedure is very effective, achieving 96% estimation coverage with R-squared ranging from 39.7% to 41.7%. The mean market beta is 0.97, close to the theoretical expectation of 1.0 (Fama & French, 1993), with 85% of estimates statistically significant. The SMB factor shows a positive mean (0.67), indicating a slight small-cap tilt consistent with Russell 3000 coverage.

The summary statistics for factor betas over the other time horizons (3-, 6-, 9-Month Horizons) can be found in Table 30 in Appendix E.

4. Methodology

This study’s methodology is built on a mixed-methods framework. This chapter first details the quantitative methods used to test the relationship between AI disclosures and stock

returns, including the LLM-based measurement framework, portfolio analysis, and regression specifications. It then outlines the explanatory case study design, a qualitative approach used to interpret, explain, and provide depth to the statistical results.

4.1 AI Disclosure Measurement Framework

This study uses LLMs to convert qualitative AI disclosures into quantitative metrics. This LLM-based approach overcomes the limitations of traditional textual analysis methods, which struggle with context and nuance (Loughran & McDonald, 2011, 2016; Araci, 2019). Traditional keyword-counting methods may falsely emphasise AI mentions with buzzwords rather than meaningful discussions using different terminology (Barrios et al., 2025; Hansen et al., 2018). LLMs evaluate how firms discuss AI, by looking at context, tone, and content to distinguish useful disclosure from empty AI talk (Gupta et al., 2024; Li et al., 2023).

LLMs, however, come with their own known limitations (Wei et al., 2022; Bojić et al., 2025). Therefore, this study applies structured prompting and categorical scoring to guide output (Dhinakaran & Jolley, 2024; Anthropic, 2025), and uses an ensemble of three LLMs (OpenAI's GPT-4o-mini, Google's Gemini Flash 1.5 and 2.5) to reduce model-specific bias. Averaging across models provides a more reliable signal, similar to inter-coder reliability in manual content analysis (Artstein & Poesio, 2008), and follows best practice in ensemble modelling to improve consistency and accuracy (Wei et al., 2022).

4.1.1 AI Factor Design and Categories

The study uses structured, categorical outputs (A–E scales) rather than continuous scores to improve classification reliability. Articles on "LLM-as-a-judge" indicate that models perform more consistently on well-defined classification tasks with clear categories compared to generating nuanced continuous scores (Dhinakaran & Jolley, 2024; Anthropic, 2025).

Initial tests with continuous scoring (1-5 scales) created interpretability problems. The difference between scores of 4.3 and 3.8 is not clearly meaningful and subjective. The methodology therefore uses categorical buckets with clear grade definitions (A-E scales with detailed criteria). This categorical approach offers advantages supported by recent research (Bojić et al., 2025). Clearly defined categorical buckets improve inter-rater reliability, leading to higher agreement scores when comparing outputs across different LLMs or against human annotators (Wang et al., 2025).

The eight disclosure metrics are organised into six broader themes. First, "Strategic Depth" measures the degree to which AI is central to the firm's long-term plan (Mishra et al.,

2022; Cockburn et al., 2018). Second, “Disclosure Sentiment” assesses the tone firms use when describing AI initiatives (Loughran & McDonald, 2011; Feldman et al., 2009). Third, AI-related risk is split into three subtypes: risks from the firm’s own adoption (Amershi et al., 2019), external threats (Christensen, 1997), and the risk of falling behind by not adopting (Brynjolfsson & McAfee, 2014). Fourth, “Forward-Looking” statements capture the extent to which organisations mention future AI projects (Muslu et al., 2015). Fifth, the “AI Washing Index” assesses whether the language is hyped or substantiated (Seele & Schultz, 2022; Al Haddi, 2024). Lastly, “Talent & Investment” indicates if firms talk about hiring, teams, or financial commitment to AI (Babina et al., 2024; Spence, 1973)."

Table 5:
AI Disclosure Scoring Framework

Factor	Scale	Grades
1. Strategic Depth (Mishra et al., 2022; Cockburn et al., 2018)	A-E (5-1)	A=Core Strategic, B=Significant Integration, C=Emerging, D=Superficial, E=No Mention
2. Disclosure Sentiment (Loughran & McDonald, 2011; Feldman et al., 2009)	A-E (5-1)	A=Clearly Positive, B=Mostly Positive, C=Neutral, D=Cautious, E=Negative/N/A
3.a Risk - Own Adoption (Alston & Bird, 2024; Amershi et al., 2019)	A-C (5-3)	A=Detailed Discussion, B=General Mention, C=No Mention
3.b Risk - External Threats (Christensen, 1997; Cockburn et al., 2018)	A-C (5-3)	A=Detailed Discussion, B=General Mention, C=No Mention
3.c Risk - Non-Adoption (Brynjolfsson & McAfee, 2014)	A-C (5-3)	A=Explicitly Discussed, B=Implicitly Suggested, C=No Mention
4. Forward-Looking (Muslu et al., 2015)	A-D (5-2)	A=Specific Plans, B=General Intent, C=Implicit Only, D=No Mention
5. AI Washing Index (Seele & Schultz, 2022; Al Haddi, 2024)	A-E (5-1)	A=Substantive, B=Mostly Substantive, C=Mixed,

Factor	Scale	Grades
		D=Mostly Hype, E=Pure Hype/N/A
6. Talent & Investment (Babina et al., 2024; Spence, 1973)	A-D (5-2)	A=Explicit Focus, B=General Mention, C=Implicit Only, D=No Mention

Note:. This table outlines the framework for scoring corporate AI disclosures. Letter grades are converted to numeric scores for analysis, as shown in parentheses.

Each theme has its grading system. “Strategic Depth”, “Disclosure Sentiment”, and the “AI Washing Index” use a five-level scale ranging from A (5) to E (1). Meanwhile, the three risk types adopt a simpler three-point scale, from A (5) to C (3), due to their more focused nature. “Forward-Looking” statements and “Talent & Investment” use four-point scales, ranging from A (5) to D (2), aiming to strike a balance between detail and clarity.

This approach ensures scores reflect disclosure substance rather than mere presence, addressing concerns about the gap between corporate AI claims and actual capabilities (Barrios et al., 2025; Barrios et al., 2025). Table 15 in Appendix A provides illustrative examples of actual corporate disclosure language from 10-K reports (Item 1A Risk Factors and Item 7 MD&A), showing how different types of AI-related text receive different scores. The examples show the qualitative differences between high-scoring disclosures (Grade A, indicating complete AI integration) and lower-scoring ones, such as generic or superficial AI mentions. They also show how the LLM scoring method separates substantive from superficial AI communications.

4.1.2 Prompt Engineering and Validation

The quality of LLM-based analysis depends heavily on prompt design quality (OpenAI, 2025; Prompting Guide, 2025). This study follows a prompt engineering approach across four key dimensions.

Categorical Output Design

The categorical structure is implemented using Python’s Literal types for strict validation:

```
AIStrategicDepthBucket = Literal[
    "A (Core Strategic Enabler)",
    "B (Significant Operational/Product Integration)",
    "C (Emerging/Exploratory Mentions)",
    "D (Superficial/Generic Mentions OR Risk of Non-Adoption Only)",
```

"E (No Mention of Own AI Strategy/Adoption)"]

For example, classifying a company's AI "Strategic Depth" as 'A (Core Strategic Enabler)' offers a clearer qualitative distinction than a numerical score. Research indicates that LLMs achieve higher consistency on well-defined classification tasks compared to nuanced continuous scoring (Wang & Wang, 2025).

Core Prompt Design Principles

Role Assignment and Context Provision: Each prompt begins by assigning the LLM a specific analytical persona for its evaluation:

"You are an expert financial analyst and AI strategist, specializing in evaluating corporate disclosures for insights into a company's engagement with Artificial Intelligence (AI). Your task is to meticulously analyse the provided text from a company's 10-K filing..."

This approach helps the model focus on key financial, strategic, and risk elements for each metric, following established practices in financial text analysis (Loughran & McDonald, 2016). Assigning roles is a crucial part of structured prompts for qualitative coding using LLMs (Dunivin, 2025). Sufficient context is provided to ensure understanding of the domain, including specifying the type of document and offering guidance on the analytical perspective.

Clear and Exact Instructions: Prompts are designed for maximum clarity, with explicit definitions provided for each categorical grade. For classification tasks, detailed criteria are included directly in the prompts, consistent with the "LLM-as-a-judge" concept (Dhinakaran & Jolley, 2024; Anthropic, 2025). For example, the AI "Strategic Depth" evaluation includes:

"A = AI is deeply integrated into core business strategy with multiple concrete, specific examples of application and impact. B = [Specific criteria]... E = AI is mentioned superficially or not at all, with no connection to core strategy."

Chain-of-Thought Reasoning: The methodology incorporates Chain-of-Thought (CoT) prompting, where the model is instructed to provide its reasoning before delivering a final classification. Chain-of-thought prompting has shown effectiveness in complex analytical tasks (Wei et al., 2022; Li et al., 2023). CoT extracts more robust reasoning pathways within the model, leading to more accurate and reliable results, particularly for complex judgment tasks (Wei et al., 2022). This approach provides a transparent audit trail for the classification process, allowing for qualitative checks on the model's decision-making logic (Dunivin, 2025).

Structured Output Requirements: Critical output formatting requirements specify the exact JSON structure:

```
```json
{
 "ai_strategic_depth": {
 "bucket_assessment": "A (Core Strategic Enabler)",
 "supporting_evidence": ["quote1", "quote2"],
 "chain_of_thought_reasoning": "Your reasoning for this factor."
 }
}
```

This explicit output structure prevents ambiguous responses and ensures information extraction in financial tasks (Wang & Wang, 2025). The structured format follows guidance from OpenAI (2025) and established prompting guides.

#### *Quality Control and Validation Mechanisms*

**Schema Enforcement for Structured Output:** To ensure consistent formatting and complete information, the study uses Pydantic models and the instructor library to define and enforce expected output structures for each prompt (Liu, 2024; Pydantic, 2025). For each AI metric, these models specify fields, data types, and allowed values using Literal types. This approach transitions from ‘hoping for good outputs’ to a more deterministic method of structured data extraction (Liu, 2024).

**Temperature Control:** All LLM generation uses a temperature of zero to ensure reproducibility by eliminating randomness and selecting the highest probability tokens (Schmalbach, 2025)

The complete prompt template, including all specific instructions and categorical definitions, is provided in Appendix B. This template uses the {document\_text} placeholder, which contains the combined "Risk Factors" and "Management’s Discussion and Analysis" sections from each company’s 10-K filing.

#### *4.1.3 Cross-Model Reliability Testing*

The same firm filings are evaluated by all three models using identical prompts to assess reliability. This cross-validation follows established inter-coder reliability methods (Artstein & Poesio, 2008). The reliability assessment uses multiple statistical measures to capture different aspects of agreement between the LLMs’ ordinal classifications:

### *Kendall's Tau Correlation Coefficient*

Kendall's tau measures rank correlation between ordinal classifications and is calculated as:

$$\tau = \frac{C - D}{C + D + T}$$

Where  $C$  represents concordant pairs (when both models rank firms in the same order),  $D$  represents discordant pairs (when models rank firms in opposite order), and  $T$  represents tied pairs (when one or both models give the same grade). Kendall's tau is appropriate for evaluating agreement between the LLMs' ordinal (A–E) classifications because it assesses the similarity of data ordering and remains robust to non-normally distributed data and tied ranks (Kendall, 1938; Goodwin & Leech, 2006). Both characteristics are common in categorical rating scales.

### *Spearman's Rank Correlation Coefficient*

Spearman's rho provides an alternative rank-based correlation measure, calculated as:

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}$$

Where  $d_i$  represents the difference in ranks between models for each firm,  $n$  is the number of firms being compared (Spearman, 1904). If GPT-4o ranks a firm 15th out of 100 firms and Gemini ranks the same firm 20th, then  $d_i = 5$  for that firm. Spearman's rho calculates how consistently the two models rank firms across the entire sample. Spearman's rho complements Kendall's tau by providing a different perspective on rank agreement, though it is more sensitive to outliers.

### *Cohen's Weighted Kappa*

Cohen's weighted kappa adjusts for chance agreement and accounts for the ordinal structure by allowing partial credit for near-miss disagreements:

$$\kappa_w = \frac{p_o - p_e}{1 - p_e}$$

Where  $p_o$  is the observed weighted agreement, and  $p_e$  is the expected agreement by random chance. The weights are defined as  $w_{ij} = 1 - \frac{|i - j|^2}{(k - 1)^2}$  where  $i$  and  $j$  represent grade categories and  $k$  is the total number of categories.

This weighted approach is particularly suitable for ordinal AI disclosure grades because disagreements between near categories (e.g., A vs B) receive partial credit, while disagreements between distant categories (e.g., A vs E) are penalised more heavily. This

reflects the underlying assumption that near-miss disagreements are less problematic than fundamental disagreements about disclosure quality (Cohen, 1960, 1968).

#### *Exact Match and Within-One-Grade Agreement*

Exact match rates calculate the percentage of observations where all models assign identical grades. Within-one-grade agreement measures the percentage of cases where models' ratings differ by at most one category level (e.g., A vs B, or B vs C). These measures provide clear benchmarks for practical agreement levels.

Section 5.1.2 presents the reliability results and cross-model agreement patterns using these statistical measures.

#### *4.1.4 Composite AI Score Construction*

A Cumulative AI Score addresses potential multicollinearity concerns and creates a single measure of overall AI disclosure quality. Multicollinearity occurs when explanatory variables are highly correlated, which can make regression results unstable when using multiple AI factors simultaneously (Brooks, 2019). This composite measure includes information across all eight disclosure dimensions while keeping the benefits of the multi-model validation approach described earlier.

The Cumulative AI Score follows a two-step calculation process. First, for each firm-year observation and each AI factor, the numeric scores from the three LLMs are averaged to create a consensus score for that dimension. This averaging reduces model-specific biases while keeping the useful information from all three. Second, the eight averaged factor scores are summed to create the Cumulative AI Score.

Calculated, for firm  $i$  in year  $t$ :

$$\text{Cumulative AI Score}_{i,t} = \sum_{k=1}^8 \frac{\text{Score}_{i,t,k}^{GPT} + \text{Score}_{i,t,k}^{Flash1.5} + \text{Score}_{i,t,k}^{Flash2.5}}{3}$$

Where  $k$  indexes the eight AI disclosure factors. The resulting composite measure ranges from 16 (the lowest possible AI disclosure across all dimensions) to 38 (the highest possible AI disclosure). This approach retains the information from each individual factor. At the same time, it provides a single composite measure that mitigates the multicollinearity that would arise from including the correlated AI factors simultaneously in a regression (Brooks, 2019).

## 4.2 Control Variable Construction

This section details the construction of all non-AI variables used in the analysis. Data sources and collection procedures are described in Chapter 3; this section focuses on variable construction methodology.

### 4.2.1 'Book' Control Variables

The regressions control for firm characteristics that affect stock returns to isolate AI disclosure effects. Following Fama and French (1992), the analysis includes controls for firm size, value, and profitability.

The traditional control variables are constructed as follows. Firm size equals the natural logarithm of market capitalisation, calculated at the quarterly filing date using shares outstanding and closing prices. The price-to-book ratio measures value, computed as market capitalisation divided by shareholders' equity from the most recent quarterly filing. Return on assets (ROA) measures profitability, calculated as trailing twelve-month net income divided by total assets. These variables follow definitions used in asset pricing studies (Fama & French, 1992).

### 4.2.2 Factor Risk Loadings

Time-varying factor loadings account for each firm's changing exposure to systematic risk using the Fama-French-Carhart six-factor model (Fama & French, 1993, 2015; Carhart, 1997). This approach measures how risk profiles change over time.

The six-factor model includes: Market Risk Premium (MKT-RF), the market return minus the risk-free rate; Size factor (SMB), excess returns of small-cap over large-cap stocks; Value factor (HML), excess returns of high book-to-market over low book-to-market stocks (Fama & French, 1993); Profitability factor (RMW), excess returns of firms with robust over weak profitability; Investment factor (CMA), excess returns of conservative over aggressive investment firms (Fama & French, 2015); and Momentum factor (UMD), excess returns of past winning over losing stocks (Carhart, 1997).

For each firm-year and investment horizon  $h \in \{3,6,9,12\}$  months, time-varying factor loadings are estimated using the following regression:

$$(R_{i,t} - RF_t) = \alpha_i + (\beta_{MKT}MKT_t - RF_t) + \beta_{SMB}SMB_t + \beta_{HML}HML_t + \beta_{RMW}RMW_t + \beta_{CMA}CMA_t + \beta_{UMD}UMD_t + \varepsilon_{i,t}$$

Where,  $R_{i,t}$  is the return on stock  $i$  in period  $t$ ,  $RF_t$  refers to the risk-free rate in period  $t$ ,  $\alpha_i$  is the firm-specific intercept. The  $\beta_{MKT}$ ,  $\beta_{SMB}$ ,  $\beta_{HML}$ ,  $\beta_{RMW}$ ,  $\beta_{CMA}$ ,  $\beta_{UMD}$  are the factor loadings,  $MKT_t$ ,  $SMB_t$ ,  $HML_t$ ,  $RMW_t$ ,  $CMA_t$ ,  $UMD_t$  refer to the factor returns in period  $t$ , and  $\varepsilon_{i,t}$  = firm-specific return error term.

The analysis employs rolling windows to prevent look-ahead bias whilst measuring time-varying risk characteristics. Look-ahead bias occurs when future information is inadvertently used to make decisions about past periods, creating unrealistically favourable results (Roberts & Whited, 2013). For each firm and horizon, loadings are estimated using six-month windows with approximately 130 daily observations. A minimum of 100 valid daily returns is required within each window. To prevent look-ahead bias, estimation windows always end before the 10-K filing date. For example, when calculating factor loadings for the three-month return period following a 10-K filing, the estimation uses daily data from six months before the filing date up to the filing date itself.

#### 4.2.3 Data Treatment and Quality Controls

Buy-and-hold excess stock returns serve as dependent variables, measured over 3-, 6-, 9-, and 12-month horizons following each 10-K filing date. Excess returns equal stock returns minus the daily risk-free rate, compounded over each horizon.

Winsorization addresses outliers using variable-specific thresholds (Adams et al., 2019). Daily returns below 0.1% are dropped to address thin trading effects (Lesmond et al., 1999). Daily returns are winsorized at  $\pm 100\%$  (Ince & Porter, 2006), with final return measures bounded at -95% to +500% (Hou et al., 2020). Price-to-book ratios are winsorized at 5th/95th percentiles due to negative book equity issues, factor betas at 2nd/98th percentiles, and other variables at 1st/99th percentiles (Boudt et al., 2020).

These horizon-specific factor loadings control for time-varying systematic risk exposure, ensuring results reflect relationships rather than outlier effects.

### 4.3 Portfolio Analysis Methods

This section describes the portfolio analysis used to examine relationships between AI disclosure and stock performance before applying multivariate controls. The analysis uses standard quintile sorts and sector-neutral techniques to separate firm-level effects from industry patterns.

#### 4.3.1 Quintile Portfolio Construction

The study constructs value-weighted portfolios using standard sorting procedures (Fama & French, 1993; Daniel et al., 1997). Each year, following 10-K filing dates, firms are ranked by their AI disclosure scores and sorted into five equal-sized portfolios (quintiles). Q1 contains firms with the lowest AI scores (predominantly E grades), while Q5 contains firms with the highest scores (frequent A grades).

Portfolios are rebalanced annually based on new AI scores from the latest filings. Buy-and-hold excess returns are tracked over 3-, 6-, 9-, and 12-month horizons following each rebalancing date. This approach follows established asset pricing methodology for characteristic-based portfolio studies (Jegadeesh & Titman, 1993).

A high-minus-low (HML) spread portfolio measures return differences between Q5 and Q1, following the zero-investment strategy approach of Fama and French (1993). The HML return tests whether stronger AI disclosure predicts superior stock performance. Statistical significance is assessed using Newey-West standard errors to address potential serial correlation (Newey & West, 1987).

#### 4.3.2 Sector-Neutral Analysis

Industry composition may confuse the AI-return relationship since both AI adoption and return patterns vary across sectors. Technology firms typically show higher AI engagement and different risk profiles than utilities or materials companies. Following Asness et al. (2000), the analysis uses sector-neutral sorting to isolate firm-level effects.

The procedure ranks firms within each GICS sector separately. Due to varying sector sizes, terciles replace quintiles to ensure adequate sample sizes across all industries. The top third of each sector forms the "High AI" portfolio, while the bottom third creates the "Low AI" portfolio. This approach ensures both portfolios have similar industry composition, removing sector bias from performance comparisons.

#### 4.3.3 Statistical Testing

Return differences are tested using appropriate statistical methods for the data structure. For HML portfolios with time-series returns, standard t-tests assess whether mean returns differ significantly from zero.

$$t = \frac{\bar{R}}{\frac{\sigma_{HML}}{\sqrt{T}}}$$

Where  $\bar{R}$  is the average HML return,  $\sigma_{HML}$  is its standard deviation, and  $T$  is the number of periods.

For sector-neutral portfolios comparing cross-sectional groups, Welch's t-test accounts for potentially unequal variances and sample sizes (Welch, 1947):

$$t = \frac{\bar{x}_{Low} - \bar{x}_{High}}{\sqrt{\frac{s_{High}^2}{n_{High}} + \frac{s_{Low}^2}{n_{Low}}}}$$

Where  $\bar{x}_{High}$ ,  $s_{High}^2$ ,  $n_{High}$  are the mean return, variance of returns, and number of firms in the 'High AI' portfolio, and  $\bar{x}_{Low}$ ,  $s_{Low}^2$ ,  $n_{Low}$  are the corresponding values for the 'Low AI' portfolio.

This test is more robust than the Student's t-test for comparisons between groups with different characteristics (Ruxton, 2006; Delacre et al., 2017).

#### 4.3.4 Sector-Specific Analysis

The analysis examines each GICS sector individually to identify whether results reflect broad patterns or sector-specific effects. This approach follows Roberts and Whited's (2013) guidance for cross-sectional studies and helps identify uneven effects across industries. The same portfolio construction methodology applies within each sector, enabling comparison of AI disclosure effects across different economic environments.

### 4.4 Regression Framework

The regression analysis aims to answer the primary research question: Do corporate AI disclosures predict subsequent stock returns? The core empirical challenge in establishing this link is endogeneity. The following sections detail the regression specifications, which are designed to address key endogeneity concerns, principally simultaneity and omitted variable bias.

Omitted variable bias may arise because both AI disclosure and future returns reflect an unobserved firm characteristic, such as skilled management, which influences genuine AI investment and effective strategic decisions. If the model ignores this trait, it produces a biased estimate of the effect of disclosure (Brooks, 2019).

Reverse causality, also known as simultaneity, means that expected performance can influence disclosure. Managers expecting weak results may use misleading AI-related disclosures to reassure investors. Conversely, managers expecting strong results might

disclose less information to protect their strategic plans. In this case, the link between disclosure and returns does not mean one causes the other (Brooks, 2019).

The empirical strategy therefore uses a fixed-effects panel regression to mitigate these biases while testing how investors interpret disclosures in their presence. By isolating the effect of AI disclosure on future returns, the framework allows for a robust test of whether markets perceive such communication as a credible signal of strength or as a negative indicator of underlying weakness.

#### 4.4.1 Baseline Specification and Fixed Effects Strategy

The study uses panel regressions to test the hypotheses, controlling for sector- and time-fixed effects. The baseline specification is:

$$ExRet_{i,t+h} = \alpha + \beta AIScore_{i,t} + \sum_s \delta_s \times Sector_{i,s} + \sum_y \gamma_y \times Year_y + \varepsilon_{i,t+h}$$

Where  $ExRet_{i,t+h}$  is the buy-and-hold excess return of firm  $i$  over the risk-free rate, measured over a horizon  $h \in \{3,6,9,12\}$  months following the 10-K filing date.  $AIScore_{i,t}$  represents one of the eight LLM-derived AI scores for that firm-year (converted to numeric 1-5 or 1-3 scale as appropriate). The model includes sector dummies ( $\delta_s$ ) that absorb any constant return differences across industries (GICS sectors) and time dummies ( $\gamma_y$ ) that absorb year-specific shocks. All models use standard errors clustered at the firm level to address both heteroskedasticity and within-firm correlation across the multiple years in the panel (Petersen, 2009). Firm-level clustering is appropriate since the same firms appear multiple times in the dataset, which can cause residuals to be correlated within firms across time periods (Cameron & Miller, 2015).

The coefficient of interest  $\beta$  measures the excess return associated with a one-grade increase in the AI disclosure score after controlling for broad industry and year effects. This baseline specification tests whether there is a predictive relationship between AI disclosure dimensions and future returns.

Although industry- and year-fixed effects control for stable firm characteristics and common annual shocks, they do not account for unobserved factors that vary over time within a firm, such as sudden changes in management quality or risk appetite. If these changing traits influence both AI disclosure and subsequent returns, the ‘AIScore’ coefficient could be biased (Brooks, 2019; Roberts & Whited, 2013).



#### *4.4.2 Progressive Control Inclusion Strategy*

The study uses a stepwise control inclusion strategy to ensure omitted variables do not drive observed AI disclosure effects. Controls are progressively added to the baseline model while monitoring the stability of the AI factor coefficient. A significant shift in the coefficient following the inclusion of a particular control suggests that the AI factor may proxy for that omitted variable. Stable coefficients provide confidence that the AI factor reflects an independent effect (Roberts & Whited, 2013).

Two control specifications are used. The "Book Controls" model adds standard firm attributes: firm size (logarithm of market capitalisation), valuation (price-to-book ratio), and profitability (return on assets). These are standard predictors in cross-sectional return studies (Fama & French, 1992) and account for size, value, and quality effects that may correlate with AI disclosure tendencies.

The "Factor Controls" model includes the six Fama-French-Carhart factor exposures (Fama & French, 1993; Carhart, 1997). These betas are estimated using daily returns over the six months preceding the 10-K filing date, as detailed in Section 4.2. Factors are added incrementally: first the three-factor model (MKT-RF, SMB, HML), then the five-factor version (adding RMW and CMA), and finally the sixth factor (UMD). This addresses the possibility that firms with high AI scores are simply growth stocks with low HML exposure or strong momentum characteristics.

#### *4.4.3 Individual Factors vs. Composite Score Approach*

Regressions are run separately for each of the eight AI disclosure dimensions rather than simultaneously. Including all eight factors together would raise multicollinearity concerns as individual scores tend to correlate (Brooks, 2019; see Section 4.1.5). Eight individual models are estimated, one for each AI factor under each control specification. This produces a coefficient for each AI metric, allowing examination of significance and stability across models.

The Cumulative AI Score addresses multicollinearity whilst capturing overall AI disclosure (Brooks, 2019). Two main composite models are estimated: one regressing excess return on the Composite AI Score with book controls and fixed effects, and another including six-factor betas alongside fixed effects. The Composite AI Score ranges between 16 and 38, with the coefficient interpreted as the return impact per one-point increase in AI disclosure quality.

#### 4.4.4 Sector-Specific Analysis

To examine whether AI disclosure effects vary consistently across sectors, separate regressions are estimated for each of the 11 GICS sectors. This approach allows for sector-specific AI effect estimation without the complications of interaction terms, providing clearer coefficient interpretation (Grace-Martin, 2009). Both book controls and factor control specifications are estimated separately for each sector  $s$ :

**(1) Book Controls with the AI Scores:**

$$ExRet_{i,t+h} = \alpha_s + \beta_{1s}AIScore_{i,t} + \sum_y \gamma_y \times Year_y + \beta_{2s}Size_{i,t} + \beta_{3s}Value_{i,t} \\ + \beta_{4s}Profitability_{i,t} + \varepsilon_{i,t+h}$$

**(2) Factor Controls with the AI Scores:**

$$ExRet_{i,t+h} = \alpha_s + \beta_{1s}AIScore_{i,t} + \sum_y \gamma_y \times Year_y + \beta_{2s}MktRF_{i,t} + \beta_{3s}SMB_{i,t} \\ + \beta_{4s}HML_{i,t} + \beta_{5s}RMW_{i,t} + \beta_{6s}CMA_{i,t} + \beta_{7s}UMD_{i,t} + \varepsilon_{i,t+h}$$

Where, the subscript  $s$  denotes sector-specific coefficients,  $i$  indexes firms,  $t$  indexes time, and  $\gamma_t$  represents year-fixed effects. All models use standard errors clustered at the firm level.

#### 4.4.5 Identification Challenges and Robustness

This section details the robustness of the regression framework, addressing potential identification challenges and validating the statistical integrity of the models.

##### *Endogeneity and Measurement Error Mitigation*

The primary threat to causal inference in this study is endogeneity. As outlined in the introduction to Section 4.4, the fixed-effects panel design is the principal strategy for mitigating bias from omitted time-invariant variables and reverse causality. In addition, measurement error in the primary independent variables, the AI scores, are a potential concern. This is addressed through the multi-model ensemble approach described in Section 4.1. By averaging scores from three different LLMs, the framework reduces individual model bias and improves the reliability of the AI disclosure metrics (Artstein & Poesio, 2008; Seni & Elder, 2010).

### *Diagnostic Testing and Specification*

Statistical robustness is ensured through several diagnostic procedures. Tests for equal error variance (Breusch-Pagan and White tests) reject homoskedasticity at the 1% level (Breusch & Pagan, 1979; White, 1980), justifying the use of clustered standard errors. Standard errors are clustered at the firm level to address both heteroskedasticity and within-firm correlation across the multiple years in the panel (Petersen, 2009). The Durbin-Watson statistic (1.91) suggests no significant autocorrelation in residuals (Durbin & Watson, 1950, 1951; Brooks, 2019).

The potential for multicollinearity among the control variables is assessed via Variance Inflation Factors (VIFs) and a Pearson correlation matrix. The correlation matrix (see Appendix F, Table 31) shows that correlations between control variables are generally low to moderate. All VIFs (see Appendix F, Table 37) remain well below the conventional threshold of 10, confirming that multicollinearity does not inflate the standard errors of the estimates (Brooks, 2019).

The influence of outliers is managed through a layered winsorisation strategy, as detailed in Section 4.2.3. While the normality of residuals is a formal requirement for statistical inference in small samples, it becomes unnecessary in large samples due to the Central Limit Theorem. With 3,995 firm-year observations, the CLT provides a strong theoretical basis for the validity of the reported p-values (Fischer, 2011; Lumley et al., 2002).

### *Justification Against Alternative Methodologies*

Other regression designs, such as Instrumental Variables (IV) or Difference-in-Differences (DiD), rely on external shocks or rules that clearly distinguish treated from control firms. Corporate AI disclosure occurs gradually and strategically over time. No reliable factor changes disclosure without also impacting returns, and no distinct policy change creates the clear threshold required by the DiD method. Therefore, a panel model with firm and year-fixed effects and controls is the most appropriate method, as it directly accounts for stable, unobserved firm heterogeneity (Brooks, 2019).

In summary, this framework follows widely accepted standards in empirical asset pricing (Fama & French, 1993; Jegadeesh & Titman, 1993) and supports the argument that observed patterns reflect genuine relationships.

## 4.5 Explanatory Case Study Design

This qualitative approach is chosen to provide explanatory depth to the statistical patterns identified in the quantitative analysis. The primary purpose of the case studies is to investigate the causal mechanisms driving the negative relationship between high AI disclosure intensity and subsequent stock returns by examining the specific context and language of selected firms.

### *4.5.1 Methodological Justification*

Case studies allow investigation of contemporary market behaviour in its real-world context (Yin, 2018). The AI disclosure paradox represents complex market behaviour that cannot be fully understood through statistical analysis alone. Case studies examine the specific language, timing, and contextual factors that quantitative measures cannot capture, providing the depth needed to understand causal mechanisms (Eisenhardt, 1989).

This approach follows established practice in accounting and finance research. Beattie and Smith (2013) use a similar design to understand how graphics in corporate reports affect firm value, combining large-scale quantitative analysis with a detailed examination of specific company disclosures. Similarly, Linsley and Shrives (2006) follow quantitative risk disclosure analysis with case examples to show the nature and quality of different disclosure strategies.

### *4.5.2 Case Selection Strategy*

The study uses extreme case sampling. This sampling strategy selects cases that represent the most significant patterns of the phenomenon under investigation (Patton, 2015).

From the three sectors showing statistically significant negative relationships (Health Care, Industrials, Consumer Discretionary), the study identifies companies with the highest AI disclosure scores that subsequently experienced the most severe negative returns. These are termed ‘AI Hyper’ firms.

This extreme case selection strategy serves two critical purposes. First, it maximises the potential for identifying clear causal mechanisms by examining instances where the phenomenon is most pronounced (Patton, 2015). Second, it follows Eisenhardt’s (1989) recommendation that theory-building case research should select cases that are likely to replicate or extend the theoretical framework being developed.

The selection criteria are:

- Sectoral Focus: Cases drawn only from sectors showing statistically significant negative AI disclosure effects
- Disclosure Intensity: Companies ranking in the top quartile of AI disclosure scores within their sector
- Performance Impact: Firms experiencing returns in the bottom decile of their sector following high AI disclosure
- Temporal Alignment: Clear temporal sequence where high AI disclosure precedes negative returns

#### *4.5.3 Data Collection and Analysis*

Case study data consists of direct quotes from 10-K filings. This follows the approach used by Linsley and Shrives (2006) in their analysis of risk reporting practices. Each case examines the specific language and framing around AI disclosure. The analysis also considers other external factors such as earnings results, market conditions, and corporate events. These factors may affect how markets interpret disclosure strategies.

The analysis applies established theoretical frameworks to explain the patterns found in disclosure language. These frameworks include agency theory (Jensen & Meckling, 1976), signalling theory (Spence, 1973), and legitimacy theory (Suchman, 1995). This theory-testing approach follows Eisenhardt's (1989) framework for using case studies to refine and extend existing theoretical models in new contexts.

#### *4.5.4 Validity and Limitations*

The extreme case sampling strategy improves validity by selecting instances where the theoretical relationship should be most clearly observable (Yin, 2018). However, this approach limits how broadly findings apply to firms with moderate AI disclosure levels. The cases aim to show and explain what drives the statistical relationship. They do not confirm how common these patterns are across all firms.

The use of multiple cases across different sectors supports external validity by showing that similar patterns occur across varied industry contexts (Yin, 2018). The application of established theoretical frameworks enhances reliability by providing consistent analytical criteria for interpreting disclosure patterns (Eisenhardt, 1989).

## 5. Intermediate Results

This section provides the foundation for the regression analysis by validating the LLM-derived AI measures, showing their characteristics, and examining the basic relationships between AI disclosure and stock returns.

### 5.1 AI Score Characteristics and Patterns

This subsection examines the AI disclosure scores produced by the LLM methodology described in Chapter 4. The analysis first characterises the patterns that emerge across firms and industries and then validates the reliability of the scoring process.

#### 5.1.1 AI Score Characteristics and Patterns

**Table 6:**  
*Descriptive Statistics of AI Disclosure Factor Scores*

Variable	Obs	Mean	Std	Min	Q25	Median	Q75	Max	Skew	Kurt	CV
AI Washing Index	3995	1.43	0.73	1.00	1.00	1.00	1.67	5.00	1.75	2.10	0.51
Disclosure Sentiment	3995	2.16	0.80	1.00	1.67	2.33	2.33	4.67	0.71	0.25	0.37
Forward-Looking	3995	2.31	0.59	2.00	2.00	2.00	2.33	5.00	1.99	3.29	0.26
Strategic Depth	3995	1.72	1.01	1.00	1.00	1.00	2.00	5.00	1.33	0.59	0.59
Talent & Investment	3995	2.33	0.54	2.00	2.00	2.00	2.33	5.00	1.95	3.66	0.23
Risk - External	3995	3.32	0.41	3.00	3.00	3.33	3.33	5.00	1.47	1.97	0.12
Risk – Non-Adoption	3995	3.49	0.52	3.00	3.00	3.33	3.67	5.00	1.10	0.56	0.15
Risk - Own Adoption	3995	3.19	0.45	3.00	3.00	3.00	3.00	5.00	2.58	5.74	0.14
AI Cumulative Score	3995	19.95	4.01	16.00	17.00	18.33	21.67	38.00	1.41	1.34	0.20


**Note:** This table presents summary statistics for the eight LLM-derived AI disclosure factors and the composite ‘AI Cumulative Score’ for the 2020–2024 sample (N=3,995 firm-years). Scores are the average from three LLMs. The ‘AI Cumulative Score’ is the sum of the eight factor scores. CV is the coefficient of variation.

Table 6 shows that substantive AI discussion in 10-K filings is relatively uncommon during 2020-2024. “Strategic Depth” and “AI Washing Index” exhibit low means (1.72 and 1.43 respectively) with median scores at minimum levels, indicating most firms made only superficial AI mentions. The risk factors require careful interpretation since they use a 3-5 scale where 3 represents "No Mention" rather than low performance. The apparently high

averages (3.19-3.49) therefore indicate minimal risk disclosure rather than relatively high scores.

These patterns provide initial evidence relevant to Hypothesis H3’s prediction of sectoral variation, with technology sectors showing markedly different AI disclosure behaviour than traditional industries. Despite generally limited substantive AI engagement, meaningful variation exists for empirical analysis. The “AI Cumulative Score” ranges from 16 to 38 with a standard deviation of 4.01, while individual factors show coefficients of variation from 0.12 to 0.59.

**Figure 2:**  
*Mean AI Disclosure Factor Scores by GICS Sector*



**Note:** This heatmap displays the mean scores for each AI disclosure factor by GICS sector, averaged over the 2020–2024 sample period. Sectors are sorted by the ‘AI Cumulative Score’. From highest (1) to lowest (11): Information Technology, Communication, Industrials, Health Care, Consumer Discretionary, Financials, Consumer Staples, Utilities, Energy, Materials, and Real Estate. Colours are scaled within each column (factor), with darker blue indicating a higher score.

Figure 2 reveals clear industry patterns. Technology-intensive sectors show highest AI engagement, with Information Technology leading (“AI Cumulative Score of 23.4) and showing strength across multiple dimensions. Communication Services follows (21.2) with particularly high awareness of competitive AI risks. Traditional sectors like Real Estate (17.7)

and Materials (18.0) show consistently low engagement across most dimensions, though all sectors converge on standardised risk reporting practices.

**Table 7:**  
*Distribution of Composite AI Disclosure Scores by GICS Sector*

Sector	Obs	N	Mean	Std	Min	Q25	Median	Q75	Max	Skew	Kurt
Information Technology	545	109	23.44	4.86	16.00	19.33	23.00	27.00	38.00	0.58	-0.50
Communication	130	26	21.16	4.12	16.33	18.33	19.67	22.58	32.00	1.16	0.31
Industrials	825	165	20.08	3.79	16.00	17.33	18.67	21.67	33.67	1.27	0.67
Health Care	350	70	20.08	4.04	16.00	17.00	18.33	22.33	32.33	1.20	0.33
Consumer Discretionary	615	123	19.57	3.89	16.00	17.00	18.00	20.67	36.33	1.60	1.92
Financials	505	101	19.32	2.99	16.00	17.33	18.33	20.33	31.00	1.49	2.02
Consumer Staples	235	47	19.24	3.44	16.00	16.83	17.67	20.67	32.33	1.55	1.99
Utilities	85	17	18.74	2.50	16.00	17.33	18.00	19.00	27.00	1.70	2.21
Energy	165	33	18.64	3.05	16.00	17.00	17.33	19.00	30.00	1.80	2.40
Materials	230	46	18.03	2.22	16.00	16.67	17.33	18.33	29.00	2.43	6.26
Real Estate	310	62	17.68	2.47	16.00	16.33	16.67	18.00	27.67	2.32	4.83
Overall	3995	799	19.95	4.01	16.00	17.00	18.33	21.67	38.00	1.41	1.34

**Note:** This table shows the distribution of the composite ‘AI Cumulative Score’ by GICS sector for the 2020–2024 sample. The table is sorted by the mean score. ‘Obs’ represents firm-year observations; ‘N’ is the number of unique firms per sector.

Table 7 confirms that Information Technology not only leads in average performance but also shows the greatest within-sector variation (standard deviation = 4.86), suggesting differences between AI-focused and traditional technology firms. Traditional sectors exhibit both lower means and reduced variation, with most firms clustering near minimum scores while technology sectors drive the overall sample distribution.

The substantial variation in AI disclosure both within and across sectors provides the empirical foundation for testing whether markets consistently price this qualitative information, as examined in the portfolio analysis that follows.




### 5.1.2 Cross-Model Reliability Analysis

The credibility of the empirical analysis depends on the quality of the LLM-derived AI disclosure scores. This subsection examines agreement patterns across the three models to validate the ensemble approach.

**Figure 3:**

*Distribution of AI Factor Grades Assigned by Each LLM for Strategic Depth and Disclosure Sentiment*



**Note:** This figure shows the percentage of firm-years (N=3,995) assigned each grade (A–E) by the three LLMs for the ‘Strategic Depth’ and ‘Disclosure Sentiment’ factors.


Figure 3 reveals both agreement and systematic differences across models. “Strategic Depth” shows broadly similar patterns, though GPT-4o assigns more top grades (A-B) while Gemini 2.5 is more conservative. “Disclosure Sentiment” shows more pronounced differences, with Gemini 1.5 assigning far more D grades than other models. These patterns underscore the value of the ensemble approach, which balances individual model biases to produce more reliable scores.

The remaining factors show additional patterns of interest (see Appendix F, Figures 8–10). Risk factors exhibit high agreement across models, with most firms receiving C grades (“No Mention”) across all three LLMs, reflecting the limited AI risk disclosure noted in Table 6. “Forward-Looking” statements and “Talent & Investment” show moderate variation but generally consistent patterns across models.

However, the AI Washing Index shows notable differences: Gemini 1.5 assigns nearly universal E grades, classifying almost all corporate AI content as “Pure Hype” or “Not Applicable,” while GPT-4o and Gemini 2.5 show more balanced distributions. This pattern

may reflect Gemini 1.5’s limitations as an older model in capturing nuanced distinctions between genuine AI substance and AI washing.

**Figure 4:**  
*Inter-Model Rank Correlation (Kendall’s Tau) of AI Factor Scores*



**Note:** This heatmap shows the Kendall’s tau ( $\tau$ ) rank correlation for AI factor scores between pairs of LLMs. Darker tiles indicate stronger agreement. The sample covers 3,995 firm-years from 2020–2024.

Figure 4 shows strong inter-model agreement for objective dimensions. “Strategic Depth” achieves Kendall’s tau values of 0.59-0.66 across model pairs, while “Risk - Own Adoption” shows even stronger agreement ( $\tau = 0.55$ -0.71). These factors have well-defined scoring criteria, making high agreement both expected and reassuring (Artstein & Poesio, 2008). “Forward-Looking” statements and “Risk - Non-Adoption” show moderate agreement ( $\tau = 0.42$ -0.54), while the “AI Washing Index” exhibits the weakest consistency, particularly between GPT-4o and Gemini 1.5 ( $\tau = 0.096$ ), reflecting the subjective nature of distinguishing "hype" from "substance."

Additional correlation analysis using Spearman’s rho (Appendix F, Figure 11) and Pearson’s r (Appendix F, Figure 12) confirms these patterns. The consistent ranking across different correlation measures confirms the reliability of the ensemble approach and validates the decision to average scores across models (Krippendorff, 2018).

**Table 8:**  
*Inter-Model Reliability Metrics for AI Factor Scores (GPT-4o-Mini vs. Gemini Flash 2.5)*

Factor	Kendall’s tau	Spearman’s rho	Cohen’s weighted kappa	Exact Match	Within 1	N
Strategic Depth	0.56	0.61	0.65	65.60	87.30	3995
Disclosure Sentiment	0.28	0.31	0.23	26.80	41.50	3995

Factor	Kendall's tau	Spearman's rho	Cohen's weighted kappa	Exact Match	Within 1	N
Risk - Own Adoption	0.55	0.57	0.65	82.70	98.60	3995
Risk - External Threats	0.34	0.35	0.37	70.70	97.30	3995
Risk - Non-Adoption	0.42	0.45	0.33	52.10	90.70	3995
Forward-Looking	0.49	0.51	0.50	78.90	86.70	3995
AI Washing Index	0.52	0.54	0.53	66.40	72.70	3995
Talent & Investment	0.37	0.39	0.36	79.50	89.60	3995

**Note:** This table reports inter-model reliability metrics between GPT-4o-Mini and Gemini Flash 2.5 for each AI factor. Kendall's tau and Spearman's rho are rank correlation coefficients. Cohen's weighted kappa measures chance-corrected agreement for ordinal data. 'Exact Match' is the percentage of identical grades; 'Within 1' is the percentage of grades within one level. N=3,995 firm-year assessments for 2020–2024.

Table 8 shows acceptable reliability across most factors. Exact match rates exceed 65% for most dimensions, while within-one-grade agreement surpasses 87% across nearly all factors. “Risk - Own Adoption” achieves the strongest agreement (82.7% exact matches, 98.6% within one grade), while “Disclosure Sentiment” shows the weakest (26.8% exact matches, 41.5% within one grade). These agreement levels compare favourably to manual content analysis studies (Goodwin & Leech, 2006), supporting the validity of using averaged scores for empirical analysis.

Comparison with additional model pairs reveals interesting patterns (see Appendix F, Tables 34 and 35). The GPT-4o-Mini vs. Gemini Flash 1.5 comparison shows notably lower agreement for the AI Washing Index ( $\tau = 0.10$  vs. 0.52), confirming the substantial disagreement identified earlier. However, the Gemini Flash 1.5 vs. Gemini Flash 2.5 comparison shows stronger agreement for Strategic Depth ( $\tau = 0.66$ ) and Risk - Own Adoption ( $\tau = 0.71$ ), suggesting that models within the same family achieve higher consistency.

The reliability assessment confirms that the ensemble approach produces sufficiently consistent measures for examining AI disclosure effects on stock returns (Lombard et al.,

2002). By averaging scores across models, this study mitigates individual model biases and creates more robust measures than any single model could provide, as explored in the portfolio analysis that follows.

## 5.2 Portfolio Analysis

The portfolio analysis examines whether the AI disclosure patterns identified in Section 5.1 translate into systematic return differences, providing initial tests of the research hypotheses before controlling for firm characteristics in Chapter 6.

### 5.2.1 Individual AI Factor Quintile Sorts

**Table 9:**  
*Performance of Portfolios Sorted on Individual AI Disclosure Factors*

AI Factor	Q1 (Low)	Q2	Q3	Q4	Q5 (High)	High-Low
AI Washing Index	3.92	8.1	5.79	3.54	-0.78	-4.70**
Disclosure Sentiment	6	5.55	5.46	3.96	-0.39	-6.39***
Forward-Looking	4.19	7.95	5.96	3.42	-0.94	-5.12**
Strategic Depth	5.43	7.81	3.37	4.72	-0.75	-6.17***
Talent & Investment	4.73	7.65	4.36	6.15	-2.32	-7.04***
Risk - External	4.97	6.08	2.82	4.06	2.66	-2.31
Risk - Non - Adoption	5.19	6.65	5.37	4.24	-0.87	-6.06***
Risk - Own Adoption	5.01	6	5.25	4.81	-0.49	-5.50**
AI Cumulative Score	7.75	4.78	4.08	5.8	-1.97	-9.72***

**Note:** This table reports annualized 12-month excess returns (%) for value-weighted portfolios sorted annually into quintiles based on individual AI factor scores (Q1 = Low, Q5 = High). ‘High-Low’ is the return on a zero-investment portfolio (long Q5, short Q1). Newey-West (1987) t-statistics are in parentheses. The sample period is 2020–2024. \*, \*\*, \*\*\* represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Using standard quintile sorting methodology (Fama & French, 1993), Table 9 shows a consistent negative relationship between AI disclosure intensity and subsequent returns. “Strategic Depth” shows a Q5-Q1 spread of -6.17% (significant at 1%), while “Talent & Investment” has -7.04%. The “AI Cumulative Score” produces the largest spread of -9.72%, which is highly significant, suggesting that AI disclosure is associated with substantial underperformance.

This pattern **challenges Hypothesis H1** and traditional signalling theory predictions, where credible disclosure should benefit quality firms (Spence, 1973). The consistent negative relationship across AI dimensions suggests that H1 may be rejected, though formal hypothesis testing requires controlling for firm characteristics and market factors in the regression analysis that follows. These patterns align with technology bubble literature showing negative market reactions to fashionable terminology during hype cycles (Cooper et al., 2001; Ofek & Richardson, 2003).

**Regarding Hypothesis H2**, the individual factor results provide **mixed preliminary evidence**. Strategic dimensions like “Talent & Investment” (-7.04%) and “Strategic Depth” (-6.17%) show substantial negative effects, but the sentiment-based “Disclosure Sentiment” factor (-6.39%) also exhibits a large negative coefficient that falls within the range of strategic factor effects. While the largest individual effect comes from “Talent & Investment” (a strategic factor), formal testing in the regression analysis will be needed to determine whether strategic dimensions are systematically stronger predictors than sentiment measures, as predicted by H2.

**Table 10:**


*Performance of Portfolios Sorted on the Composite AI Score Across Different Horizons*

Time Horizon	Q1 (Low)	Q2	Q3	Q4	Q5 (High)	High-Low
3 Month	1.65	0.68	-0.06	0.24	-2.32	-3.97***
6 Month	3.18	2.85	3.2	3.7	-0.93	-4.11***
9 Month	6.91	4.51	5.79	6.1	-0.28	-7.19***
12 Month	7.75	4.78	4.08	5.8	-1.97	-9.72***

**Note:** This table reports annualized excess returns (%) over 3-, 6-, 9-, and 12-month horizons for value-weighted portfolios sorted annually into quintiles based on the ‘AI Cumulative Score’ (Q1 = Low, Q5 = High). ‘High-Low’ is the return on a zero-investment portfolio (long Q5, short Q1). Newey-West (1987) t-statistics are in parentheses. The sample period is 2020–2024. \*, \*\*, \*\*\* represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 10 shows this underperformance is not temporary but grows steadily over time, from -3.97% at 3 months to -9.72% at 12 months.


**Figure 5:**  
*Cumulative Performance of Portfolios Sorted on the Composite AI Disclosure Score*



**Note:** This figure plots the cumulative value-weighted excess returns for portfolios sorted annually on the ‘Composite AI Disclosure Score’. Q1 (green line) is the lowest-scoring quintile; Q5 (orange line) is the highest. The sample period is 2020–2024.

Figure 5 visualises the year-by-year average excess returns for the lowest (Q1) and highest (Q5) “AI Cumulative Score” portfolios. In every year except 2024, Q5 underperforms Q1, with especially large gaps in 2020 and 2022. This time-series pattern reinforces the notion of consistent underperformance among high-AI firms.

Figure 6 plots 12-month excess returns against the “AI Cumulative Score” at the firm level. The negative slope and significant Pearson correlation ( $\rho = -0.087$ ,  $p < 0.01$ ) show that higher AI disclosure scores are associated with lower future returns on average. Despite considerable noise, the downward trend line supports the portfolio sorting results and confirms the negative relationship exists across the full distribution of firms rather than being driven by extreme quintiles alone.

**Figure 6:***Relationship Between Composite AI Score and Subsequent 12-Month Excess Return*

**Note:** This scatterplot shows the relationship between the ‘AI Cumulative Score’ (x-axis) and subsequent 12-month excess return (y-axis) for each firm-year observation (2020–2024). The red line is the OLS fit.

### 5.2.2 Sector-Specific Quintile Sorts

**Table 11:***Sector-Neutral Portfolio Sorts on the Composite AI Score*

Sector	High-Low Spread (%)	T-Stat	P-Value	N Firms
Communication	13.04	1.19	0.237	66
Consumer Discretionary	-12.24***	-2.71	0.007	331
Consumer Staples	-10.95*	-1.76	0.08	136
Energy	8.64	0.81	0.418	93
Financials	1.32	0.33	0.743	274
Health Care	-9.71*	-1.76	0.08	201
Industrials	-12.57***	-3.39	0.001	471

Sector	High-Low Spread (%)	T-Stat	P-Value	N Firms
Information Technology	-3.72	-0.8	0.424	307
Materials	-9.09	-1.44	0.152	138
Real Estate	-1.38	-0.33	0.745	178
Utilities	-2.25	-0.35	0.725	44

**Note:** This table reports the annualized 12-month excess return for a sector-neutral ‘High-Low’ portfolio. Firms are sorted into terciles based on their ‘AI Cumulative Score’ within each GICS sector annually. T-statistics are from a Welch’s t-test. ‘N Firms’ is the number of firm-year observations per sector. The sample period is 2020–2024. \*, \*\*, \*\*\* represent statistical significance at the 10%, 5%, and 1% levels, respectively.

The sector-neutral analysis in Table 11 shows that the negative effect persists within several sectors (Asness et al., 2000), indicating that industry composition does not fully explain the patterns. Consumer Discretionary (-12.24%) and Industrials (-12.57%) show the strongest within-sector effects, both significant at 1%. Information Technology shows a smaller, insignificant spread (-3.72%), while Communication Services shows a positive (though insignificant) spread (13.04%).

These results provide **preliminary support for Hypothesis H3**, which predicted systematic variation across industry sectors. The dramatic differences between traditional sectors and technology sectors suggest that industry context matters critically for interpreting AI disclosures, consistent with competency-based disclosure theory where market reactions depend on perceived alignment between firm capabilities and strategic commitments (Barrios et al., 2025).

The portfolio analysis reveals a clear puzzle: Firms with extensive AI communication experience poor subsequent stock performance. This challenges expectations about technology disclosure based on prior research showing positive associations between AI communication and firm performance (Mishra et al., 2022; Babina et al., 2024). The relationship appears robust across multiple AI dimensions, persists over various horizons, and exists within industry peer groups, motivating the multivariate regression analysis that follows.

## 6. Main Results

This chapter tests whether the negative relationship between AI disclosure and future returns shown in Chapter 5 remains after controlling for known stock performance drivers.



The analysis progresses from testing individual AI factor regressions to analysing sector-specific data. It demonstrates that the intermediate results reveal meaningful economic relationships rather than false patterns caused by missing control variables.

The regression framework tests the three core hypotheses from Chapter 2. Hypothesis H1 predicted a positive relationship between AI disclosure quality and future returns, consistent with signalling theory (Spence, 1973). Hypothesis H2 suggested that strategic dimensions would be stronger predictors than sentiment-based measures, following disclosure theory literature (Muslu et al., 2015). Hypothesis H3 anticipated variation across industry sectors, based on competency-based signalling research (Asness et al., 2000). The results challenge conventional assumptions about corporate technology disclosure whilst highlighting the importance of industry context.

## 6.1 Individual AI Factor Analysis

### 6.1.1 Build-up Model Demonstration

**Table 12:**  
*Progressive Regressions of 12-Month Excess Returns on the Composite AI Score*

Variable	Base	Model 1	Model 2	Model 3
Cum AI Score	-0.0114*** (0.0020)	-0.0108*** (0.0020)	-0.0102*** (0.0020)	-0.0094*** (0.0021)
Log Market Cap		-0.0060 (0.0046)	-0.0018 (0.0050)	-0.0064 (0.0051)
Price-to-Book			-0.0044** (0.0019)	-0.0054*** (0.0019)
ROA				1.1734*** (0.4014)
$R^2$	0.0459	0.0465	0.0483	0.0526
Adj. $R^2$	0.0414	0.0417	0.0431	0.0472
Observations	3,174	3,174	3,174	3,171

**Note:** This table presents panel regressions of 12-month excess returns on the ‘AI Cumulative Score’. The columns show the progressive inclusion of book controls. All models include GICS sector- and year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. \*, \*\*, \*\*\* represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 12 shows four regression specifications with progressive control inclusion. The “AI Cumulative Score” coefficient declines modestly from -0.0114 in the Base model to

-0.0094 in the full Model 3 specification, a 17% decline that indicates stability across control specifications.

A one-unit increase in the “AI Cumulative Score” implies a 0.94 percentage point decrease in 12-month excess returns in the full specification. Given the standard deviation of approximately 4.01, a one-standard-deviation increase implies a 3.77 percentage point decrease in annual returns.

This stability contradicts traditional signalling theory predictions (Spence, 1973; Ross, 1977), which suggest that quality firms should benefit from credible disclosure. Instead, the consistently negative coefficients align with dot-com era evidence of technology buzzword penalties (Cooper et al., 2001) and contemporary AI washing concerns highlighted by regulators (U.S. Securities and Exchange Commission, 2024).

### 6.1.2 Factor Results

**Table 13:**  
*Regressions of 12-Month Excess Returns on Individual AI Disclosure Factors*

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
AI Washing Index	-0.0456*** (0.0104)	-0.0516*** (0.0103)
Disclosure Sentiment	-0.0351*** (0.0097)	-0.0395*** (0.0095)
Forward-Looking	-0.0570*** (0.0128)	-0.0653*** (0.0125)
Strategic Depth	-0.0302*** (0.0077)	-0.0336*** (0.0076)
Talent & Investment	-0.0568*** (0.0136)	-0.0672*** (0.0135)
Risk - External	-0.0093 (0.0207)	-0.0118 (0.0205)
Risk - Non - Adoption	-0.0516*** (0.0162)	-0.0579*** (0.0165)
Risk - Own Adoption	-0.0435** (0.0216)	-0.0552*** (0.0214)
AI Cumulative Score	-0.0094*** (0.0021)	-0.0106*** (0.0020)

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
Average $R^2$	0.0502	0.036
Average Observations	3,171	3,031

**Note:** This table presents coefficients from separate panel regressions of 12-month excess returns on each AI factor (18 total regressions: 9 factors  $\times$  2 specifications). Column (1) includes book controls (Log Market Cap, P/B, ROA). Column (2) includes the six Fama-French-Carhart factor exposures. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. \*, \*\*, \*\*\* represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 13 shows seven of eight individual factors exhibit negative coefficients significant at the 5% level or better in both specifications, with only Risk External showing insignificant effects.

These results **reject Hypothesis H1**. With seven of eight AI disclosure dimensions showing significant negative effects after controlling for firm characteristics and market factors, the multivariate analysis confirms that higher quality AI disclosure consistently associates with lower future returns, contradicting signalling theory predictions (Spence, 1973; Ross, 1977).

“Forward-Looking” statements show the largest negative effects, with coefficients of -0.0570 and -0.0653. A one-unit increase in “Forward-Looking” AI language associates with a 5.70 to 6.53 percentage point decrease in 12-month excess returns. “Talent & Investment” shows similar large negative effects (-0.0568 and -0.0672), while “Strategic Depth” shows coefficients of -0.0302 and -0.0336.

“Disclosure Sentiment” shows coefficients of -0.0351 and -0.0395, indicating that more positive AI discussion tone associates with lower returns. The “AI Washing Index” shows effects of -0.0456 and -0.0516.

“Risk Non-Adoption” shows strong negative effects (-0.0516 and -0.0579), while “Risk Own Adoption” shows coefficients of -0.0435 and -0.0552. “Risk External Threats” remains the sole exception, with small, insignificant coefficients.

These patterns directly contradict theoretical predictions. The severe penalties for “Forward-Looking” statements contradict predictions that credible forward-looking disclosures should enhance firm value (Muslu et al., 2015). The negative “Talent & Investment” contradict signalling theory expectations that visible investments should signal firm quality (Spence, 1973). Instead, suggesting investor concerns about agency costs and overinvestment in fashionable technologies (Jensen, 1986; Malmendier & Tate, 2015).

“Strategic Depth” penalties contradict research suggesting AI strategic integration should improve performance (Mishra et al., 2022).

These findings provide **partial support for Hypothesis H2**, though not in the expected direction. Strategic factors generally emerge as stronger predictors than sentiment-based measures, with Forward-Looking statements and Talent & Investment showing the largest negative effects, while Disclosure Sentiment exhibits relatively smaller coefficients. However, formal statistical tests comparing coefficient magnitudes would be needed to confirm whether these differences are statistically significant. The pattern suggests that markets distinguish between different types of AI communication and react most negatively to substantive strategic commitments rather than superficial AI mentions, indicating that investors view concrete AI strategies as particularly risky signals in the current market environment.

## 6.2 Sector Regression Analysis

The analysis examines whether AI disclosure effects vary across industry sectors, testing Hypothesis H3’s prediction of industry-specific variation. This approach follows research showing that firm-level signals are more informative when considered relative to industry peers (Asness et al., 2000), and that AI adoption patterns differ significantly across sectors based on technological infrastructure and competency requirements (Cockburn et al., 2018; Brynjolfsson et al., 2021).

### 6.2.1 Cross-Sector Regression Results

**Table 14:**  
*Sector-Level Regressions of 12-Month Excess Returns on the Composite AI Disclosure Score*

Sector	(Book Controls Model)			(Fama French 6 Factor Model)			
	Coefficient	P-Value	R <sup>2</sup>	Coefficient	P-Value	R <sup>2</sup>	N
Communication	-0.08% (1.29)	0.951	0.125	-1.11% (1.18)	0.346	0.261	102
Consumer Discretionary	-1.68%*** (0.49)	5.55E-04	0.077	-1.85%*** (0.48)	1.17E-04	0.064	489
Consumer Staples	-0.34% (0.76)	0.657	0.095	-0.64% (0.69)	0.358	0.045	187
Energy	1.18% (1.80)	0.510	0.092	1.63% (1.60)	0.308	0.168	131

	(Book Controls Model)			(Fama French 6 Factor Model)			
Financials	0.30% (0.78)	0.698	0.121	0.43% (0.69)	0.538	0.124	403
Health Care	-1.66%*** (0.58)	0.004	0.039	-0.91% (0.56)	0.106	0.045	276
Industrials	-1.76%*** (0.39)	6.79E-06	0.060	-1.72%*** (0.40)	1.53E-05	0.057	658
Information Technology	-0.13% (0.53)	0.805	0.035	-0.46% (0.53)	0.383	0.032	432
Materials	-0.74% (1.13)	0.511	0.120	-0.018	0.063	0.089	184
Real Estate	-0.66% (0.79)	0.405	0.262	-0.14% (0.75)	0.849	0.304	244
Utilities	-0.45% (0.82)	0.581	0.265	-0.008	0.097	0.242	68

**Note:** This table presents coefficients for the ‘AI Cumulative Score’ from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors  $\times$  2 specifications). The ‘Book Controls Model’ includes firm size, P/B, and ROA. The ‘Fama French 6 Factor Model’ includes the six factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. \*, \*\*, \*\*\* represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 14 shows three sectors with significant negative coefficients: Industrials (-1.76%,  $p < 0.001$ ), Consumer Discretionary (-1.68%,  $p < 0.001$ ), and Health Care (-1.66%,  $p = 0.004$ ), whilst Information Technology shows no significant relationship (-0.13%,  $p = 0.805$ ).

**Traditional Sector Penalties:** Industrials show the strongest negative relationship, with coefficients of -0.0176 and -0.0172. A one-unit increase in the “AI Cumulative Score” implies a 1.76 to 1.72 percentage point decrease in annual returns within this sector. Given the within-sector standard deviation of AI scores (approximately 3.8 for Industrials), a one-standard-deviation increase implies roughly 6.7 percentage point lower returns.

Consumer Discretionary show similar strong patterns with coefficients of -0.0168 and -0.0185. Combined with within-sector AI score variation (standard deviation  $\approx 3.9$ ), this translates to approximately 6.5-7.2 percentage point return differences between high and low AI disclosure firms within the sector.

These negative reactions in traditional sectors may reflect self-selection bias, where firms under operational pressure use AI-heavy language to signal strength while potentially lacking fundamental AI capabilities (Roy, 1951; Hava, 2025).

**Technology Sector Neutrality:** Information Technology shows virtually no relationship between AI disclosure and returns, with small coefficients (-0.0013 and -0.0046) that are statistically insignificant. This neutrality indicates that within the technology sector, AI disclosure carries minimal informational content for investors. This aligns with market efficiency theory, where AI disclosures provide minimal new information in sectors with rich information environments and sophisticated investor bases (Fama, 1970).

These sector regressions provide **support for Hypothesis H3**, showing systematic variation across industry contexts. The large differences in both coefficient magnitudes and statistical significance across sectors confirm that industry context critically matters for interpreting AI disclosures. Traditional sectors show highly significant negative effects, while Information Technology shows no significant relationship ( $p=0.805$ ). This pattern strongly supports competency-based disclosure theory, where market reactions depend on perceived alignment between firm capabilities and strategic commitments (Barrios et al., 2025).

#### *6.2.2 Individual Factor Analysis by Sector*

Sector-specific analysis of individual AI factors reveals the factors driving aggregate patterns. This analysis tests each of the eight AI factors separately within each sector (8 factors  $\times$  11 sectors  $\times$  2 specifications = 176 total regressions). Detailed results appear in Appendix D, Tables 22-29.

“Strategic Depth” emerges as a particularly strong negative predictor in traditional sectors. Consumer Discretionary show coefficients of -6.61% and -7.27% for “Strategic Depth”, indicating that a one-unit increase in strategic AI positioning implies a 6.61 to 7.27 percentage point return penalty. Industrials show similar patterns with coefficients of -6.57% and -6.13%. These magnitudes are much larger than the effects shown in Table 14, suggesting strategic AI positioning is viewed as especially risky when undertaken by firms outside their core technological competencies, consistent with competency-based signalling theory (Barrios et al., 2025).

Market reactions become even more severe for future-oriented AI commitments in traditional sectors. Industrials show the largest negative coefficients for “Forward-Looking” statements (-10.85% and -10.38%), while Consumer Discretionary follows with coefficients of -8.76% and -10.06%. These severe penalties contradict theoretical predictions that credible “Forward-Looking” disclosures should enhance firm value (Muslu et al., 2015), instead suggesting markets view such commitments as evidence of overoptimism in traditional sectors.

“Talent & Investment” discussions trigger particularly strong negative reactions in sectors lacking technical infrastructure. Consumer Discretionary show coefficients of -9.46% and -10.25%, Health Care show -11.48% and -6.73%, and Industrials show -9.19% and -9.57%. These patterns suggest investor concerns about the costs and execution risks of building AI capabilities in sectors lacking established technical infrastructures, contradicting signalling theory predictions that talent hiring and investments should signal firm commitment and quality (Spence, 1973; Babina et al., 2024).

Acknowledging AI-related risks amplifies rather than reduces negative market reactions in traditional sectors, directly contradicting traditional disclosure theory (Healy & Palepu, 2001). Consumer Discretionary firms discussing AI adoption risks face severe return penalties across different risk dimensions. The “Risk Non-Adoption” factor shows particularly pronounced effects in Consumer Discretionary (-13.70% and -13.91%) and Industrials (-9.35% and -9.11%). “Risk Own Adoption” shows even more severe effects, with Consumer Discretionary coefficients of -15.98% and -18.31%.

Information Technology consistently shows small, insignificant coefficients across all individual factors, reinforcing that AI discourse carries minimal informational content within technology-native sectors. This neutrality extends across all AI disclosure dimensions, supporting theoretical predictions that AI capabilities represent baseline expectations rather than differentiating factors within technology-native industries (Bonsón et al., 2021).

These sector-specific patterns provide strong support for competency-based disclosure theory (Barrios et al., 2025), where market reactions depend critically on perceived alignment between firm capabilities and strategic commitments.

## 6.3 Hypothesis Testing and Theoretical Framework

### 6.3.1 Hypothesis Testing Results and Economic Significance

The clearest and most surprising finding is that firms with higher scores on AI-related disclosure dimensions tend to show significantly lower excess stock returns in the following 12 months. This negative relationship remains strong even after accounting for firm-specific characteristics, factor exposures, and fixed effects for both sector- and year (Roberts & Whited, 2013).

The economic significance of these relationships is substantial and exceeds many traditional asset pricing anomalies. At the composite level, a one-standard-deviation increase in AI disclosure associates with 3.77 percentage point lower annual returns. The classic size

effect confirmed by Fama and French (1992) typically produces annual return differences of 2-4 percentage points, indicating that the observed AI disclosure effects represent economically meaningful relationships. Within traditional sectors, return spreads reach 6.7 percentage points between high and low AI disclosure firms, creating potential opportunities for systematic investment strategies (Barber & Odean, 2008).

**Hypothesis H1 (AI Disclosure and Returns) - Rejected:** The results contradict signalling theory predictions (Spence, 1973; Ross, 1977). Higher AI disclosure associates with lower future returns rather than signalling firm strength. This negative relationship appears across most disclosure categories.

This negative pattern appears across nearly all disclosure categories, contradicting the fundamental premise that firms with superior private information should benefit from credible disclosure.

**Hypothesis H2 (Strategic vs. Sentiment Dimensions) - Partial Support (Unexpected Direction):** Factors tied to deliberate planning and investment, such as “Strategic Depth”, “Talent & Investment”, and “Forward-Looking” language, are indeed the strongest predictors of returns, but their effect is negative rather than positive as predicted by strategic disclosure theory (Muslu et al., 2015) and talent investment research (Babina et al., 2024). The finding that substantive strategic discourse is penalised more heavily than superficial mentions suggests investors distinguish between different types of AI communication and view strategic commitments as particularly risky signals, aligning with research on managerial overconfidence where concrete commitments in uncertain domains may signal overoptimism rather than strategic clarity (Malmendier & Tate, 2015).

**Hypothesis H3 (Sectoral Differences) - Supported:** The strength and direction of the AI disclosure effect vary sharply by industry, with effect magnitudes varying dramatically from significant negative effects in traditional sectors to neutral effects in technology sectors. In traditional sectors like Industrials, Consumer Discretionary, and Health Care, the return penalty for high AI disclosure is especially pronounced, whilst in Information Technology, where AI is already integrated into business models, the effect is close to zero. This pattern strongly supports competency-based theories of corporate disclosure, where market reactions depend on perceived alignment between firm capabilities and strategic commitments (Barrios et al., 2025).



### *6.3.2 Theoretical Framework for AI Disclosure Paradox*

The unexpected negative relationship between AI disclosure and returns requires theoretical explanation, as it contradicts predictions from signalling theory (Spence, 1973; Ross, 1977). This section presents a post-hoc theoretical analysis developed to explain both the overall negative effects and the variation across sectors.

Three related theoretical frameworks explain this paradox. These frameworks work together to show why firms with more extensive AI communication experience worse subsequent returns.

#### *Information Asymmetry and Market Scepticism*

**Self-Selection and Signalling Distortions:** A main reason for the negative relationship is self-selection bias, where firms facing operational or financial difficulties may use language emphasising AI to appear strong or to divert attention from their issues. This is a well-known pattern in markets where information is unevenly shared (Roy, 1951; Spence, 1973). This interpretation aligns with recent evidence suggesting that companies may use disclosure tone to hide weak company fundamentals (Hava, 2025). The self-selection mechanism creates adverse selection problems where investors cannot distinguish genuine AI capabilities from "cheap talk," leading to market-wide discounting of all AI disclosures (Akerlof, 1970).

Evidence from earlier technology bubbles supports this view. During the blockchain boom, underperforming firms rebranded or made vague claims about blockchain use to ride investor excitement, only to later underdeliver (Jain & Jain, 2019). The dot-com bubble showed a similar pattern. During this bubble companies with weak fundamentals used internet-related terminology to inflate valuations (Ofek & Richardson, 2003). The AI disclosures may represent a similar pattern, where struggling firms use technological buzzwords to distract and signal strength.

**Proprietary Cost Considerations:** The strong negative reactions to “Strategic Depth” disclosures could reflect proprietary cost concerns, where detailed AI disclosures reveal competitive information to rivals (Verrecchia, 1983). In highly competitive sectors, investors may expect that strategic AI disclosures will decrease future competitive advantages, leading to negative valuation effects despite potentially positive underlying projects (Admati & Pfleiderer, 2000).

### *Agency Problems and Managerial Behaviour*

**Agency Costs and Overinvestment:** The negative relationship may reflect investor concerns about managerial overinvestment in fashionable but unproven technologies. Agency theory suggests that managers with substantial free cash flow may invest in negative-NPV projects to build empires rather than return cash to shareholders (Jensen, 1986). The particularly strong negative reactions to “Talent & Investment” disclosures support this interpretation, suggesting markets view AI hiring and other investment announcements as potential evidence of value-destroying overinvestment rather than strategic capability building (Malmendier & Tate, 2015).

**Hype Cycle Effects:** The sample period (2020-2024) covered a significant AI hype cycle (PwC, 2024; Manning & Napier, 2024). Following the emergence of generative models, firms raced to mention AI in filings during increasing investor excitement (Manning & Napier, 2024). This period may reflect a pattern similar to earlier technology bubbles, where expectations outpaced results (Fenn, 1995). Such dynamics potentially expose firms with optimistic AI disclosures to negative revaluation as implementation challenges emerge. The progressive strengthening of negative effects over investment horizons suggests markets are learning to increasingly discount AI narratives (Fama, 1970; Barberis et al., 1998).

### *Competency-Based Signalling and Sectoral Context*

**Linking Signalling and Core Competency Theory:** The identified sectoral differences suggest linking signalling theory (Spence, 1973) with core competency theory (Prahalad & Hamel, 1990). Traditional signalling models assume that expensive, hard-to-fake disclosures should be credible signals of firm quality. However, these findings suggest that signalling effectiveness also depends on whether the disclosed strategy aligns with the companies’ core business.

When traditional manufacturing or retail firms make detailed AI disclosures, markets could interpret these as negative signals that management is moving away from proven business areas rather than positive signals of innovation (Barney, 1991). This may explain why Consumer Discretionary and Industrials firms face negative reactions to AI announcements. On the contrary, when technology firms make similar AI disclosures, markets view these as routine updates about expected capabilities rather than meaningful new information, leading to neutral reactions (Teece et al., 1997). This could explain why Information Technology firms show no significant market response to AI disclosures.

**Legitimacy and Stakeholder Concerns:** The sectoral variation also reflects differential stakeholder reactions to AI adoption. Legitimacy theory suggests that firms disclose activities to align with stakeholder expectations and social norms (Suchman, 1995). In sectors where AI raises ethical, employment, or safety concerns, disclosures may signal potential legitimacy threats rather than strategic opportunities (Freeman, 1984). The negative reactions in traditional sectors may partly reflect stakeholder concerns about job displacement, safety risks, or departure from established industry practices.

**Market Efficiency and Information Processing:** The null effects in Information Technology suggest that AI disclosures in this sector provide little new information to already well-informed investors and analysts. This matches market efficiency theory. In sectors with many informed investors and analysts, new information gets processed quickly and accurately (Fama, 1970). The contrast is true with traditional sectors, where information asymmetries are greater and information processing ability weaker. This supports the view that disclosure effectiveness varies with market structure and investor knowledge (Lowry & Shu, 2002).

**Summary:** The overall negative relationship likely comes from a mix of self-selection problems, agency costs, and hype cycle effects. The differences between sectors arise from how well AI fits with a firm's core business, stakeholder reactions, and how much investors know about each industry. This thesis suggests the market's negative reaction to AI disclosures during 2020-2024 has a clear explanation: Rather than viewing AI announcements as positive innovation signals, investors treated them as warning signs.

## 6.4 Explanatory Case Studies

To show how the theory explains the AI disclosure paradox, this section examines the most extreme examples, focusing on companies with the highest AI disclosure scores but large negative returns. Appendix C provides five additional case studies relating to the three major negatively impacted sectors.

### 6.4.1 Case Study Analysis

#### **Case Study 1: Chegg, Inc. (CHGG) - Consumer Discretionary Sector**

- *2023 Form 10-K, Filed: 20 February 2024*
- *AI Score: 36.33 (Highest in their industry)*
- *12-Month Return: -72.99% (Worst performance)*

Chegg provides the most dramatic example of the AI disclosure paradox. The company achieved the highest AI disclosure score while experiencing the worst stock performance across the Consumer Discretionary industry sample. The company's traditional homework assistance model is disrupted by free generative AI tools like ChatGPT, with Q4 2023 earnings showing a 21% year-over-year decline in subscribers and decreasing revenues. Management announced the company is undergoing strategic review to consider all options, including a potential sale.

This context possibly created extreme agency pressures where management needed to maintain stakeholder confidence while acknowledging fundamental competitive threats (Jensen & Meckling, 1976). The extensive AI disclosure represents an attempt to reframe existential disruption as strategic opportunity through impression management (Merkl-Davies & Brennan, 2007).

The company's 2023 10-K filing reveals the underlying vulnerability:

*"Our competition may develop products and technologies that are similar or superior to our technologies or are more cost-effective to develop and deploy... Given the long history of development in the AI sector, other parties may have (or in the future may obtain) patents or other proprietary rights that would prevent, limit, or interfere with our ability to make, use, or sell our own AI products."* (Chegg, Inc., 2024)

This admission shows information so damaging that firms would only disclose it if hiding it posed greater regulatory risks (Spence, 1973; Verrecchia, 1983). The self-selection pattern becomes clear: Chegg felt forced to discuss AI extensively precisely because its business model had become outdated, creating a situation where extensive disclosure signals weakness rather than strength (Akerlof, 1970; Li & Prabhala, 2007).

#### 6.4.2 Systematic Behavioural Patterns

The Chegg case shows specific patterns that extend across all firms with high AI disclosure scores and negative returns. Analysis of additional companies across the three significant sectors reveals consistent patterns that explain the AI disclosure paradox through five interconnected theoretical frameworks.

The case study analysis of AI-intensive firms across Health Care, Industrials, and Consumer Discretionary sectors reveals that extensive AI disclosure consistently occurs under conditions of business pressure rather than strategic strength. Every case shows the same sequence: underlying business challenges create agency pressures for management,

leading to impression management strategies that create negative signalling effects through self-selection bias.

The analysis validates the explanatory power of the integrated theoretical framework across different industry contexts, showing that the AI disclosure paradox results from systematic responses to business pressures rather than random market behaviour:

**Agency Theory and Impression Management** (Jensen & Meckling, 1976; Merkl-Davies & Brennan, 2007): Cases reveal management teams under pressure to maintain stakeholder confidence while facing performance challenges, competitive threats, or strategic failures. The extensive AI discourse represents attempts to reframe negative developments within aspirational technological narratives.

**Signalling Theory Distortions** (Spence, 1973; Ross, 1977): Cases reveal systematic "over-signalling without substance" where extensive AI language attempts to signal strength while concrete evidence reveals weakness, creating contradictory signals that markets correctly interpret as negative.

**Self-Selection and Market for Lemons** (Roy, 1951; Li & Prabhala, 2007; Akerlof, 1970): Companies consistently self-select into extensive AI disclosure when facing business model threats, failed investments, or competitive disadvantage, creating adverse selection where disclosure intensity correlates with underlying vulnerability rather than strategic capability.

**Legitimacy Theory** (Suchman, 1995; Freeman, 1984): In traditional sectors, extensive AI disclosure often signals potential stakeholder concerns about employment displacement or operational risks rather than competitive advantages.

**Proprietary Cost Theory** (Verrecchia, 1983): Cases reveal companies inadvertently disclosing competitive vulnerabilities while attempting to demonstrate AI capabilities.

The consistent negative market reactions across all significant sectors suggest sophisticated investor interpretation of patterns rather than simple punishment of transparency. Markets appear capable of distinguishing genuine strategic communication from impression management driven by agency costs or business model threats. The subsequent performance of these companies generally validates the market's initial negative assessments, confirming that extensive AI communication often accompanies rather than contradicts fundamental business weaknesses (Fama, 1970).

## 6.5 Practical Implications

**For Investors:** The findings provide clear guidance for portfolio construction. High AI disclosure creates significant negative return effects in Consumer Discretionary, Health Care, and Industrials while showing neutral effects in Information Technology and Financials. These patterns create potential for sector-specific investment strategies that avoid high AI disclosure firms in the Consumer Discretionary, Health Care, and Industrials sectors.

**For Management:** The results suggest firms may benefit from exercising restraint when discussing AI initiatives in formal disclosures. The significant negative effects in Consumer Discretionary, Health Care, and Industrials indicate that markets interpret extensive AI communication in these sectors as potential overinvestment or "AI washing" rather than strategic capability (Jensen, 1986; Barrios et al., 2025). Technology and Financial sector firms face less severe penalties, suggesting that AI commitments within core competencies trigger weaker negative reactions.

**For Regulators:** The negative market reactions in specific sectors validate current regulatory concerns about AI disclosure practices. The SEC's recent enforcement actions align with market evidence showing sector-dependent penalties for AI disclosure intensity (U.S. Securities and Exchange Commission, 2024). The sector-specific results suggest regulatory frameworks could account for industry-specific factors when evaluating AI disclosure adequacy.

The findings indicate that markets distinguish between sectors when evaluating AI disclosure of any kind. Even firms making substantive disclosures about AI strategies, talent hiring, AI-related investments, and forward-looking plans face substantial penalties in Consumer Discretionary, Health Care, and Industrials, suggesting that markets view detailed AI commitments in these sectors as potentially value-destroying rather than credible signals of future performance.

## 7. Conclusion

### *Research Focus and Empirical Design*

This thesis examines whether AI-related corporate disclosures predict stock market performance, analysing Russell 3000 constituents from 2020-2024. The study employs LLMs to score AI narratives across six dimensions in 10-K filings, moving beyond keyword

approaches to capture subtle disclosure quality. Results are tested using portfolio sorts and panel regressions controlling for firm characteristics and risk factors.

### *Broad Take-aways of Main Findings*

This study tests whether AI disclosure benefits firms, following established signalling theory. The results show the opposite pattern. Markets penalise firms with higher AI disclosure scores, particularly in Consumer Discretionary, Health Care, and Industrials sectors. This contradicts theoretical predictions. Technology firms show neutral effects, suggesting AI capabilities represent baseline expectations rather than differentiating factors.

### *Theoretical Mechanisms*

Three interconnected theories explain these patterns. **Self-selection bias** occurs when operationally struggling firms use extensive AI narratives to mask underlying weaknesses (Roy, 1951; Spence, 1973). Case studies of Omnicell and 3D Systems show this pattern, with extensive AI disclosure coinciding with earnings disappointments and insider selling.

**Information asymmetry** arises because AI capabilities depend on unverifiable intangible assets, creating adverse selection where markets rationally discount all AI claims (Akerlof, 1970). **Legitimacy concerns** emerge when AI adoption signals stakeholder threats rather than opportunities, especially in relation to job displacement and operational risks in traditional sectors (Suchman, 1995).

### *Market Efficiency Validation*

The findings demonstrate market sophistication rather than dysfunction. Investors correctly identify that extensive AI disclosure correlates with business vulnerability, not strategic strength. This validates efficient market processing of complex narrative information (Fama, 1970), contradicting concerns about narrative-driven capital misallocation.

### *Contributions and Practical Relevance*

Methodologically, the ensemble LLM framework provides a replicable template for analysing complex disclosure topics. Theoretically, the results extend signalling theory by showing how disclosure can signal weakness when firms venture beyond core competencies.

For practitioners, investors should scrutinise AI claims from traditional sectors where negative effects concentrate. Managers should prioritise substance over extensive disclosure, as ambitious AI narratives without verifiable action signal weakness. Regulators gain support for targeted oversight while recognising market efficiency in processing AI information.

### *Limitations and Future Research*

LLM assessments may contain biases requiring human validation. Future research should link disclosures to tangible AI investments, extend analysis beyond US markets, and track whether disclosure penalties persist as AI technology matures beyond the 2020-2024 hype cycle.



## References

- Accenture. (2024, January 9). *Technology vision 2024: Human by design—How AI unleashes the next level of human potential*. <https://www.accenture.com/gb-en/insights/technology/technology-trends-2024>
- Adams, J. C., Hayunga, D., Mansi, S., Reeb, D., & Verardi, V. (2019). Identifying and treating outliers in finance. *Financial Management*, 48(2), 345–384. <https://doi.org/10.1111/fima.12269>
- Admati, A. R., & Pfleiderer, P. (2000). Forcing firms to talk: Financial disclosure regulation and externalities. *The Review of Financial Studies*, 13(3), 479–519. <https://doi.org/10.1093/rfs/13.3.479>
- Akerlof, G. A. (1970). The market for "lemons": Quality uncertainty and the market mechanism. *The Quarterly Journal of Economics*, 84(3), 488–500. <https://doi.org/10.2307/1879431>
- Al Haddi, H. (2024, December 3). *AI washing: The cultural traps that lead to exaggeration and how CEOs can stop them*. California Management Review. <https://cmr.berkeley.edu/2024/12/ai-washing-the-cultural-traps-that-lead-to-exaggeration-and-how-ceos-can-stop-them/>
- Alexandris, C. (2024). GenAI and socially responsible AI in natural language processing applications: A linguistic perspective. In *Proceedings of the 2024 AAAI Spring Symposium Series* (Vol. 3, No. 1). AAAI Press. <https://doi.org/10.1609/aaaiss.v3i1.31230>
- Alston & Bird. (2024, July 26). *Navigating AI-related disclosure challenges: Securities filing, SEC enforcement, and shareholder litigation trends* [Legal advisory]. <https://www.alston.com/en/insights/publications/2024/07/navigating-ai-related-disclosure-challenges>
- Amershi, S., Begel, A., Bird, C., DeLine, R., Gall, H., Kamar, E., Nagappan, N., Nushi, B., & Zimmermann, T. (2019). Software engineering for machine learning: A case study. *2019 IEEE/ACM 41st International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*, 291–300. <https://doi.org/10.1109/ICSE-SEIP.2019.00042>

- Anthropic. (2025, June 13). *How we built our multi-agent research system*. <https://www.anthropic.com/engineering/built-multi-agent-research-system>
- Anthropic. (2025). Test & evaluate: Create strong empirical evaluations. Anthropic Documentation. <https://docs.anthropic.com/en/docs/test-and-evaluate/develop-tests#tone-and-style-customer-service-llm-based-likert-scale>
- Aptech. (2019, March 12). *Introduction to the fundamentals of panel data*. Aptech Blog. <https://www.aptech.com/blog/introduction-to-the-fundamentals-of-panel-data/>
- Araci, D. (2019). *FinBERT: Financial sentiment analysis with pre-trained language models* (arXiv:1908.10063). arXiv. <https://doi.org/10.48550/arXiv.1908.10063>
- Artstein, R., & Poesio, M. (2008). Inter-coder agreement for computational linguistics. *Computational Linguistics*, 34(4), 555–596. <https://doi.org/10.1162/coli.07-034-R2>
- Asness, C. S., Porter, R. B., & Stevens, R. L. (2000). *Predicting stock returns using industry-relative firm characteristics* (SSRN Working Paper No. 213872). <https://doi.org/10.2139/ssrn.213872>
- Babina, T., Fedyk, A., He, A. X., & Hodson, J. (2024). Artificial intelligence, firm growth, and product innovation. *Journal of Financial Economics*, 151, Article 103745. <https://doi.org/10.1016/j.jfineco.2023.103745>
- Barber, B. M., & Odean, T. (2008). All that glitters: The effect of attention and news on the buying behavior of individual and institutional investors. *The Review of Financial Studies*, 21(2), 785–818. <https://doi.org/10.1093/rfs/hhm079>
- Barberis, N., Shleifer, A., & Vishny, R. (1998). A model of investor sentiment. *Journal of Financial Economics*, 49(3), 307–343. [https://doi.org/10.1016/S0304-405X\(98\)00027-0](https://doi.org/10.1016/S0304-405X(98)00027-0)
- Barney, J. B. (1991). Firm resources and sustained competitive advantage. *Journal of Management*, 17(1), 99–120. <https://doi.org/10.1177/014920639101700108>
- Barrios, J. M., Campbell, J. L., Johnson, R. G., & Liu, Y. C. (2025). *Artificially intelligent or artificially inflated? Determinants and informativeness of corporate AI disclosures* [SSRN working paper]. SSRN. <https://doi.org/10.2139/ssrn.5133107>

Barrios, J. M., Fedyk, A., & Fedyk, T. (2025, March 3). *How accurate are corporate AI disclosures?* Columbia Law School Blue Sky Blog.

<https://elsbluesky.law.columbia.edu/2025/03/03/how-accurate-are-corporate-ai-disclosures/>

Beattie, V., & Smith, S. J. (2013). Value creation and business models: Refocusing the intellectual capital debate. *The British Accounting Review*, 45(4), 243–254.

<https://doi.org/10.1016/j.bar.2013.06.001>

Beatty, A., Cheng, M., & Zhang, H. (2018). Are risk-factor disclosures still relevant? Evidence from market reactions to risk-factor disclosures before and after the financial crisis. *Contemporary Accounting Research*, 36(2), 805–838. <https://doi.org/10.1111/1911-3846.12444>

BlackRock. (2025, June 23). *iShares Russell 3000 ETF (IWR): Fund holdings and information*. <https://www.ishares.com/us/products/239714/ishares-russell-3000-etf>

Bojić, L., Zagovora, O., Zelenkauskaitė, A., Vuković, V., Čabarkapa, M., Veseljević Jerković, S., & Jovančević, A. (2025). Comparing large language models and human annotators in latent content analysis of sentiment, political leaning, emotional intensity and sarcasm. *Scientific Reports*, 15, Article 11477. <https://doi.org/10.1038/s41598-025-96508-3>

Bommasani, R., Hudson, D. A., Adeli, E., Altman, R., Arora, S., von Arx, S., ... Liang, P. (2021). *On the opportunities and risks of foundation models* (arXiv:2108.07258). arXiv. <https://doi.org/10.48550/arXiv.2108.07258>

Bonsón, E., Lavorato, D., Lamboglia, R., & Mancini, D. (2021). Artificial intelligence activities and ethical approaches in leading listed companies in the European Union. *International Journal of Accounting Information Systems*, 43, Article 100535. <https://doi.org/10.1016/j.accinf.2021.100535>

Boudt, K., Todorov, V., & Wang, W. (2020). Robust distribution based winsorization in composite indicators construction. *Social Indicators Research*, 149(2), 375–397. <https://doi.org/10.1007/s11205-019-02259-w>

Breiman, L. (1996). Bagging predictors. *Machine Learning*, 24(2), 123–140. <https://doi.org/10.1007/BF00058655>

- Bresnahan, T. F., & Trajtenberg, M. (1995). General purpose technologies: ‘Engines of growth?’. *Journal of Econometrics*, 65(1), 83–108. [https://doi.org/10.1016/0304-4076\(94\)01598-T](https://doi.org/10.1016/0304-4076(94)01598-T)
- Breusch, T. S., & Pagan, A. R. (1979). A simple test for heteroscedasticity and random coefficient variation. *Econometrica*, 47(5), 1287–1294. <https://doi.org/10.2307/1911963>
- Brooks, C. (2019). *Introductory econometrics for finance* (4th ed.). Cambridge University Press.
- Brown, S. J., Goetzmann, W. N., Ibbotson, R. G., & Ross, S. A. (1992). Survivorship bias in performance studies. *The Review of Financial Studies*, 5(4), 553–580. <https://doi.org/10.1093/rfs/5.4.553>
- Brown, S. V., & Tucker, J. W. (2010). Large-sample evidence on firms’ year-over-year MD&A modifications. *Journal of Accounting Research*, 49(2), 309–346. <https://doi.org/10.1111/j.1475-679X.2010.00396.x>
- Brynjolfsson, E., & McAfee, A. (2014). *The second machine age: Work, progress, and prosperity in a time of brilliant technologies*. W. W. Norton.
- Brynjolfsson, E., Rock, D., & Syverson, C. (2021). The productivity J-curve: How intangibles complement general-purpose technologies. *American Economic Journal: Macroeconomics*, 13(1), 333–372. <https://doi.org/10.1257/mac.20180386>
- Burrell, J. (2016). How the machine ‘thinks’: Understanding opacity in machine learning algorithms. *Big Data & Society*, 3(1). <https://doi.org/10.1177/2053951715622512>
- Cameron, A. C., & Miller, D. L. (2015). A practitioner’s guide to cluster-robust inference. *Journal of Human Resources*, 50(2), 317–372. <https://doi.org/10.3368/jhr.50.2.317>
- Campbell, J. L., Cecchini, M., Cianci, A. M., Ehinger, A. C., & Werner, E. M. (2018). Tax-related mandatory risk-factor disclosures, future profitability, and stock returns. *Review of Accounting Studies*, 24(1), 264–308. <https://doi.org/10.1007/s11142-018-9474-y>
- Cao, Y., Liu, M., Qiu, J., & Zhao, R. (2024a). *Information in disclosing emerging technologies: Evidence from AI disclosure* (SSRN Working Paper No. 4987085). <https://doi.org/10.2139/ssrn.4987085>

Cao, Y., Liu, M., Qiu, J., & Zhao, R. (2024, December 19). *What firms disclose about their use of AI*. Columbia Law School Blue Sky Blog.

<https://elsbluesky.law.columbia.edu/2024/12/19/what-are-companies-disclosing-about-their-use-of-ai/>

Carhart, M. M. (1997). On persistence in mutual fund performance. *The Journal of Finance*, 52(1), 57–82. <https://doi.org/10.1111/j.1540-6261.1997.tb03808.x>

Carhart, M. M., Carpenter, J. N., Lynch, A. W., & Musto, D. K. (2002). Mutual fund survivorship. *The Review of Financial Studies*, 15(5), 1439–1463.

<https://ssrn.com/abstract=1300796>

Chegg, Inc. (2024). *Form 10-K*. U.S. Securities and Exchange Commission.

Chen, S., Green, T. C., Gulen, H., & Zhou, D. (2025). *What does ChatGPT make of historical stock returns? Extrapolation and miscalibration in LLM stock-return forecasts* [Working paper]. Mitch Daniels School of Business, Purdue University.

<https://dx.doi.org/10.2139/ssrn.4941906>

Christensen, C. M. (1997). *The innovator's dilemma: When new technologies cause great firms to fail*. Harvard Business School Press.

Cockburn, I. M., Henderson, R., & Stern, S. (2018). *The impact of artificial intelligence on innovation* (NBER Working Paper No. 24449). National Bureau of Economic Research.

<https://doi.org/10.3386/w24449>

Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46. <https://doi.org/10.1177/001316446002000104>

Cohen, J. (1968). Weighted kappa: Nominal scale agreement with provision for scaled disagreement or partial credit. *Psychological Bulletin*, 70(4), 213–220.

<https://doi.org/10.1037/h0026256>

Cohen, L., Malloy, C. J., & Nguyen, Q. (2020). Lazy prices. *The Journal of Finance*, 75(3), 1371–1415. <https://doi.org/10.1111/jofi.12885>

Cooper, M. J., Dimitrov, O., & Rau, P. R. (2001). A rose.com by any other name. *The Journal of Finance*, 56(6), 2371–2388. <https://doi.org/10.1111/0022-1082.00408>

- Corvel Corporation. (2024). *Form 10-K*. U.S. Securities and Exchange Commission.
- Creswell, J. W., & Plano Clark, V. L. (2018). *Designing and conducting mixed methods research* (3rd ed.). SAGE Publications.
- Cui, J. (2025). *AI-driven digital transformation and firm performance in Chinese industrial enterprises: Mediating role of green digital innovation and moderating effects of human-AI collaboration* (arXiv:2505.11558). arXiv. <https://doi.org/10.48550/arXiv.2505.11558>
- Damodaran, A. (2025). *Equity risk premiums (ERP): Determinants, estimation, and implications – The 2025 edition* (SSRN Working Paper No. 4777040). <https://dx.doi.org/10.2139/ssrn.5168609>
- Daniel, K., Grinblatt, M., Titman, S., & Wermers, R. (1997). Measuring mutual fund performance with characteristic-based benchmarks. *The Journal of Finance*, 52(3), 1035–1058. <https://doi.org/10.1111/j.1540-6261.1997.tb02724.x>
- Delacre, M., Lakens, D., & Leys, C. (2017). Why psychologists should by default use Welch’s t-test instead of Student’s t-test. *International Review of Social Psychology*, 30(1), 92–101. <https://doi.org/10.5334/irsp.82>
- Delmas, M. A., & Burbano, V. C. (2011). The drivers of greenwashing. *California Management Review*, 54(1), 64–87. <https://doi.org/10.1525/cm.2011.54.1.64>
- Dhinakaran, A., & Jolley, E. (2024, March 3). Why you should not use numeric evals for LLM as a judge. Arize AI. <https://arize.com/blog-course/numeric-evals-for-llm-as-a-judge/>
- Düger, Y. S. (2020). ‘Willful blindness’ as a dark side of organizational behavior: Can effective leadership overcome this challenge? In *Studies on interdisciplinary economics and business* (Vol. III). Peter Lang GmbH. <https://www.researchgate.net/publication/344491271>
- Dunivin, Z. O. (2025). Scaling hermeneutics: A guide to qualitative coding with LLMs for reflexive content analysis. *EPJ Data Science*, 14, Article 28. <https://doi.org/10.1140/epjds/s13688-025-00548-8>
- Durbin, J., & Watson, G. S. (1950). Testing for serial correlation in least squares regression. I. *Biometrika*, 37(3/4), 409–428. <https://doi.org/10.1093/biomet/37.3-4.409>

- Durbin, J., & Watson, G. S. (1951). Testing for serial correlation in least squares regression. II. *Biometrika*, 38(1/2), 159–178. <https://doi.org/10.1093/biomet/38.1-2.159>
- Easton, P. D., & Zmijewski, M. E. (1993). SEC Form 10-K/10-Q reports and annual reports to shareholders: Reporting lags and squared market-model prediction errors. *Journal of Accounting Research*, 31(1), 113–129. <https://doi.org/10.2307/2491044>
- Eisenhardt, K. M. (1989). Building theories from case study research. *Academy of Management Review*, 14(4), 532–550. <https://doi.org/10.5465/amr.1989.4308385>
- Fama, E. F. (1970). Efficient capital markets: A review of theory and empirical work. *The Journal of Finance*, 25(2), 383–417. <https://doi.org/10.2307/2325486>
- Fama, E. F., & French, K. R. (1992). The cross-section of expected stock returns. *The Journal of Finance*, 47(2), 427–465. <https://doi.org/10.1111/j.1540-6261.1992.tb04398.x>
- Fama, E. F., & French, K. R. (1993). Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics*, 33(1), 3–56. [https://doi.org/10.1016/0304-405X\(93\)90023-5](https://doi.org/10.1016/0304-405X(93)90023-5)
- Fama, E. F., & French, K. R. (2015). A five-factor asset pricing model. *Journal of Financial Economics*, 116(1), 1–22. <https://doi.org/10.1016/j.jfineco.2014.10.010>
- Fama, E. F., & French, K. R. (2025, May 20). *U.S. daily factors, 1963–2025*. Retrieved from Wharton Research Data Services. <https://wrds-www.wharton.upenn.edu/>
- Feldman, R., Govindaraj, S., Livnat, J., & Segal, B. (2009). Management’s tone change, post earnings announcement drift and accruals. *Review of Accounting Studies*, 15(4), 915–953. <https://doi.org/10.1007/s11142-009-9111-x>
- Fenn, J. (1995, August 1). *When to leap on the hype cycle*. Gartner.
- Fischer, H. (2011). *A history of the central limit theorem: From classical to modern probability theory*. Springer. <https://doi.org/10.1007/978-0-387-87857-7>
- Fleisher, W. (2022). Understanding, idealization, and explainable AI. *Episteme*, 19(4), 534–560. <https://doi.org/10.1017/epi.2022.39>
- Freeman, R. E. (1984). *Strategic management: A stakeholder approach*. Pitman.



- FTSE Russell. (2025, May 31). *Russell 3000® Index* [Factsheet]. LSEG.  
<https://research.ftserussell.com/Analytics/FactSheets/Home/DownloadSingleIssue?issueName=US3000USD&isManual=True>
- Garel, A., Gilbert, A., & Scott, A. (2019). *Linguistic complexity and cost of equity capital* (SSRN Working Paper No. 3240292). <https://doi.org/10.2139/ssrn.3240292>
- Goodwin, L. D., & Leech, N. L. (2006). Understanding correlation: Factors that affect the size of r. *The Journal of Experimental Education*, 74(3), 249–266.  
<https://doi.org/10.3200/JEXE.74.3.249-266>
- Grace-Martin, K. (2009, January 19). *Interpreting interactions in regression*. The Analysis Factor. <https://www.theanalysisfactor.com/interpreting-interactions-in-regression/>
- Greene, J. (2025, February 5). *New wave of investor lawsuits targets ‘AI washing’*. Reuters.  
<https://www.reuters.com/legal/government/column-new-wave-investor-lawsuits-targets-ai-washing-2025-02-05/>
- Grewal, G. S. (2024, April 15). *Remarks at Program on Corporate Compliance and Enforcement Spring Conference 2024*. U.S. Securities and Exchange Commission.  
<https://www.sec.gov/newsroom/speeches-statements/gurbir-remarks-pcce-041524>
- Grigoryev, L. M. (2023). The shocks of 2020–2023 and the business cycle. *Journal of the New Economic Association*, 60(4), 8–32. <http://dx.doi.org/10.17323/2949-5776-2023-1-1-8-32>
- Gupta, S., Ranjan, R., & Singh, S. N. (2024). *Comprehensive study on sentiment analysis: From rule-based to modern LLM-based system* (arXiv:2409.09989). arXiv.  
<https://doi.org/10.48550/arXiv.2409.09989>
- Hansen, S., McMahon, M., & Prat, A. (2018). Transparency and deliberation within the FOMC: A computational linguistics approach. *The Quarterly Journal of Economics*, 133(2), 801–870. <https://ideas.repec.org/a/oup/qjecon/v133y2018i2p801-870..html>
- Harris, T. (2024). Managers’ perception of product market competition and earnings management: A textual analysis of firms’ 10-K reports. *Journal of Accounting Literature*, 47(3), 499–524. <https://doi.org/10.1108/JAL-11-2022-0116>



Hava, K. B. (2025). *The AI-induced disclosure pressure model and empirical evidence from Management Discussion and Analysis (MD&A) reporting* (SSRN Working Paper No. 4938415). <http://dx.doi.org/10.13140/RG.2.2.27598.37447>

Healy, P. M., & Palepu, K. G. (2001). Information asymmetry, corporate disclosure, and the capital markets: A review of the empirical disclosure literature. *Journal of Accounting and Economics*, 31(1–3), 405–440. [https://doi.org/10.1016/S0165-4101\(01\)00018-0](https://doi.org/10.1016/S0165-4101(01)00018-0)

Heckman, J. J. (1979). Sample selection bias as a specification error. *Econometrica*, 47(1), 153–161.

Henry, E. (2008). Are investors influenced by how earnings press releases are written? *The Journal of Business Communication*, 45(4), 363–407. <https://doi.org/10.1177/0021943608319388>

Heyliger, A., Rong, J., & Grapsas, R. (2023, July 21). *Quarterly review of corporate governance and public company disclosure developments*. Weil, Gotshal & Manges LLP. <https://governance.weil.com/featured/quarterly-review-of-corporate-governance-and-public-company-disclosure-developments/>

Hope, O.-K., Hu, D., & Lu, H. (2016). The benefits of specific risk-factor disclosures. *Review of Accounting Studies*, 21(4), 1005–1045. <https://doi.org/10.1007/s11142-016-9371-1>

Hou, K., Xue, C., & Zhang, L. (2020). Replicating anomalies. *The Review of Financial Studies*, 33(5), 2019–2133. <https://doi.org/10.1093/rfs/hhy131>

Ince, O. S., & Porter, R. B. (2006). Individual equity return data from Thomson Datastream: Handle with care! *Journal of Financial Research*, 29(4), 463–479. <https://dx.doi.org/10.2139/ssrn.486523>

iRobot Corporation. (2023). *Form 10-K*. U.S. Securities and Exchange Commission.

Jacobs, B., Hartog, J., & Vijverberg, W. (2009). Self-selection bias in estimated wage premiums for earnings risk. *Empirical Economics*, 37(2), 271–286. <https://doi.org/10.1007/s00181-008-0231-0>

Jain, A., & Jain, P. C. (2019). Blockchain hysteria: Adding “blockchain” to company’s name. *Economics Letters*, 181, 178–181. <https://wisconsin->

[uwlax.primo.exlibrisgroup.com/permalink/01UWI\\_LC/1slgdsd/cdi\\_proquest\\_miscellaneous\\_2352979559](http://uwlax.primo.exlibrisgroup.com/permalink/01UWI_LC/1slgdsd/cdi_proquest_miscellaneous_2352979559)

Jegadeesh, N., & Titman, S. (1993). Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of Finance*, 48(1), 65–91.  
<https://doi.org/10.1111/j.1540-6261.1993.tb04702.x>

Jensen, M. C., & Meckling, W. H. (1976). Theory of the firm: Managerial behavior, agency costs and ownership structure. *Journal of Financial Economics*, 3(4), 305–360.  
[https://doi.org/10.1016/0304-405X\(76\)90026-X](https://doi.org/10.1016/0304-405X(76)90026-X)

Jensen, M. C. (1986). Agency costs of free cash flow, corporate finance, and takeovers. *The American Economic Review*, 76(2), 323–329.

Juluru, K., Shih, H.-H., Murthy, K. N. K., & Elnajjar, P. (2021). Bag-of-words technique in natural language processing: A primer for radiologists. *RadioGraphics*, 41(6), 1797–1806.  
<https://doi.org/10.1148/rg.2021210025>

Kearney, C., & Liu, S. (2014). Textual sentiment in finance: A survey of the literature. *Journal of Economic Surveys*, 28(2), 367–397. <https://dx.doi.org/10.2139/ssrn.2213801>

Kendall, M. G. (1938). A new measure of rank correlation. *Biometrika*, 30(1/2), 81–93.  
<https://doi.org/10.1093/biomet/30.1-2.81>

Kravet, T. D., & Muslu, V. (2013). Textual risk disclosures and investors' risk perceptions. *Review of Accounting Studies*, 18(4), 1088–1122. <https://doi.org/10.1007/s11142-013-9228-9>

Krippendorff, K. (2018). Content analysis: An introduction to its methodology (4th ed.). Sage.

Lesmond, D. A., Ogden, J. P., & Trzcinka, C. A. (1999). A new estimate of transaction costs. *The Review of Financial Studies*, 12(5), 1113–1141. <https://doi.org/10.1093/rfs/12.5.1113>

Lesmy, D., Muchnik, L., & Mugeran, Y. (2019). *Do you read me? Temporal trends in the language complexity of financial reporting* (SSRN Working Paper No. 3469073).  
<https://doi.org/10.2139/ssrn.3469073>

- Li, F. (2008). Annual report readability, current earnings, and earnings persistence. *Journal of Accounting and Economics*, 45(2–3), 221–247. <https://doi.org/10.1016/j.jacceco.2008.02.003>
- Li, K., & Prabhala, N. R. (2007). *Self-selection models in corporate finance*. (Robert H. Smith School Research Paper No. RHS 06-020) [Working paper]. SSRN. <https://dx.doi.org/10.2139/ssrn.843105>
- Li, X., Chan, S., Zhu, X., Pei, Y., Ma, Z., Liu, X., & Shah, S. A. R. (2023). *Are ChatGPT and GPT-4 general-purpose solvers for financial text analytics?* (arXiv:2305.05862). arXiv. <https://doi.org/10.48550/arXiv.2305.05862>
- Linsley, P. M., & Shrives, P. J. (2006). Risk reporting: A study of risk disclosures in the annual reports of UK companies. *The British Accounting Review*, 38(4), 387–404. <https://doi.org/10.1016/j.bar.2006.05.002>
- Liu, J. (2024, January 4). *Steering large language models with Pydantic*. Pydantic. <https://pydantic.dev/articles/llm-intro>
- Lombard, M., Snyder-Duch, J., & Bracken, C. C. (2002). Content analysis in mass communication: Assessment and reporting of intercoder reliability. *Human Communication Research*, 28(4), 587–604. <https://doi.org/10.1111/j.1468-2958.2002.tb00826.x>
- Loughran, T., & McDonald, B. (2011). When is a liability not a liability? Textual analysis, dictionaries, and 10-Ks. *The Journal of Finance*, 66(1), 35–65. <https://doi.org/10.1111/j.1540-6261.2010.01625.x>
- Loughran, T., & McDonald, B. (2014). Measuring readability in financial disclosures. *The Journal of Finance*, 69(4), 1643–1671. <https://doi.org/10.1111/jofi.12162>
- Loughran, T., & McDonald, B. (2016). Textual analysis in accounting and finance: A survey. *Journal of Accounting Research*, 54(4), 1187–1230. <https://doi.org/10.1111/1475-679X.12123>
- Lowry, M., & Shu, S. (2002). Litigation risk and IPO underpricing. *Journal of Financial Economics*, 65(3), 309–335. [https://doi.org/10.1016/S0304-405X\(02\)00144-7](https://doi.org/10.1016/S0304-405X(02)00144-7)
- Lumley, T., Diehr, P., Emerson, S., & Chen, L. (2002). The importance of the normality assumption in large public health data sets. *Annual Review of Public Health*, 23, 151–169. <https://doi.org/10.1146/annurev.publhealth.23.100901.140546>

- Mahat, D., Neupane, D., & Shrestha, S. (2024). Quantitative research design and sample trends: A systematic examination of emerging paradigms and best practices. *Cognizance Journal of Multidisciplinary Studies*, 4(2), 20–27.  
<https://doi.org/10.47760/cognizance.2024.v04i02.002>
- Malkiel, B. G. (2003). The efficient market hypothesis and its critics. *Journal of Economic Perspectives*, 17(1), 59–82. <https://doi.org/10.1257/089533003321164958>
- Malmendier, U., & Tate, G. (2015). Behavioral CEOs: The role of managerial overconfidence. *Journal of Economic Perspectives*, 29(4), 37–60.  
<https://doi.org/10.1257/jep.29.4.37>
- Manning & Napier Advisors, LLC. (2024, April 16). *The AI hype cycle*. Manning & Napier.  
<https://www.manning-napier.com/insights/the-ai-hype-cycle>
- McKinsey & Company. (2024, May 30). *The state of AI in early 2024: Gen AI adoption spikes and starts to generate value*.  
<https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2024>
- Merkel-Davies, D. M., & Brennan, N. M. (2007). Discretionary disclosure strategies in corporate narratives: Incremental information or impression management? *Journal of Accounting Literature*, 26, 116–194. <https://ssrn.com/abstract=1089447>
- Mishra, S. (2024, April 11). *AI governance appears on corporate radar*. Harvard Law School Forum on Corporate Governance. <https://corpgov.law.harvard.edu/2024/04/11/ai-governance-appears-on-corporate-radar/>
- Mishra, S., Ewing, M. T., & Cooper, H. B. (2022). Artificial intelligence focus and firm performance. *Journal of the Academy of Marketing Science*, 50(6), 1176–1197.  
<https://doi.org/10.1007/s11747-022-00876-5>
- Moodaley, W., & Telukdarie, A. (2023). Greenwashing, sustainability reporting, and artificial intelligence: A systematic literature review. *Sustainability*, 15(2), Article 1481.  
<https://doi.org/10.3390/su15021481>

Muslu, V., Rebello, M. J., & Xu, D. (2015). Forward-looking MD&A disclosures and the information environment. *Management Science*, 61(5), 973–991.

<https://ssrn.com/abstract=2386967>

Nelson, K. K., & Pritchard, A. C. (2016). Carrot or stick? The shift from voluntary to mandatory disclosure of risk factors. *Journal of Empirical Legal Studies*, 13(2), 266–297.

<https://doi.org/10.1111/jels.12115>

Newey, W. K., & West, K. D. (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3), 703–708.

<https://doi.org/10.2307/1913610>

Ofek, E., & Richardson, M. (2003). DotCom mania: The rise and fall of Internet stock prices. *The Journal of Finance*, 58(3), 1113–1137. <https://doi.org/10.1111/1540-6261.00560>

Omnicell, Inc. (2022). *Form 10-K*. U.S. Securities and Exchange Commission.

OpenAI. (2023). *GPT-4 technical report* (arXiv:2303.08774). arXiv.

<https://doi.org/10.48550/arXiv.2303.08774>

OpenAI. (May 22, 2025) *Prompt engineering.*, from

<https://platform.openai.com/docs/guides/prompt-engineering>

Patton, M. Q. (2015). *Qualitative research & evaluation methods: Integrating theory and practice* (4th ed.). SAGE Publications.

Petersen, M. A. (2009). Estimating standard errors in finance panel data sets: Comparing approaches. *The Review of Financial Studies*, 22(1), 435–480.

<https://doi.org/10.1093/rfs/hhn053>

Prabhakar, A. V. (2025, February 15). *FailSafeQA: Evaluating AI hallucinations, robustness, and compliance in financial LLMs*. Ajith's AI Pulse. <https://ajithp.com/2025/02/15/failsafeqa-evaluating-ai-hallucinations-robustness-and-compliance-in-financial-llms/>

Prahalad, C. K., & Hamel, G. (1990). The core competence of the corporation. *Harvard Business Review*, 68(3), 79–91.

Prompting Guide. (2025, April 24). *Prompting guide* (April 2025).

<https://www.promptingguide.ai>

PwC. (2024). *2025 AI business predictions*. <https://www.pwc.com/us/en/tech-effect/ai-analytics/ai-predictions.html>

Pydantic. (2025, May 24th). *Pydantic documentation (Version 2)*. Retrieved June 22, 2025, from <https://docs.pydantic.dev/>

Regulwar, G. B., Mahalle, A., Pawar, R., Shamkuwar, S., & Kakde, P. R. (2023). Big data collection, filtering, and extraction of features. In *Big data analytics techniques for market intelligence* (pp. 136–158). IGI Global. <https://doi.org/10.4018/979-8-3693-0413-6.ch005>

Roberts, M. R., & Whited, T. M. (2013). Endogeneity in empirical corporate finance. In G. M. Constantinides, M. Harris, & R. M. Stulz (Eds.), *Handbook of the economics of finance* (Vol. 2, pp. 493–572). Elsevier. <https://doi.org/10.1016/B978-0-44-453594-8.00007-0>

Ross, S. A. (1977). The determination of financial structure: The incentive-signalling approach. *The Bell Journal of Economics*, 8(1), 23–40. <https://doi.org/10.2307/3003485>

Roy, A. D. (1951). Some thoughts on the distribution of earnings. *Oxford Economic Papers*, 3(2), 135–146. <https://www.jstor.org/stable/2662082>

Ruxton, G. D. (2006). The unequal variance t-test is an underused alternative to Student's t-test and the Mann–Whitney U test. *Behavioral Ecology*, 17(4), 688–690. <https://doi.org/10.1093/beheco/ark016>

S&P Global Market Intelligence. (2025, May 20). *Compustat North America—Daily updates (1963–2025)*. Wharton Research Data Services. <https://wrds-www.wharton.upenn.edu/>

S&P Global Market Intelligence. (2025, May 20). *Compustat North America fundamentals quarterly: U.S. and Canadian company financials, 1961–2025*. Wharton Research Data Services. <https://wrds-www.wharton.upenn.edu/>

Schmalbach, V. (2025, February 7). *Does temperature 0 guarantee deterministic LLM outputs?* <https://www.vincenteschmalbach.com/p/does-temperature-0-guarantee-deterministic>

Seele, P., & Schultz, M. D. (2022). From greenwashing to machinewashing: A model and future directions derived from reasoning by analogy. *Journal of Business Ethics*, 178(4), 1063–1089. <https://doi.org/10.1007/s10551-022-05054-9>

Seni, G., & Elder, J. F. (2010). *Ensemble methods in data mining: Improving accuracy through combining predictions*. Morgan & Claypool.  
<http://dx.doi.org/10.2200/S00240ED1V01Y200912DMK002>

Shiller, R. J. (2000). *Irrational exuberance*. Princeton University Press.

Spearman, C. (1904). “General intelligence,” objectively determined and measured. *The American Journal of Psychology*, 15(2), 201–292. <https://doi.org/10.2307/1412107>

Spence, A. M. (1973). Job market signaling. *The Quarterly Journal of Economics*, 87(3), 355–374. <https://doi.org/10.2307/1882010>

Suchman, M. C. (1995). Managing legitimacy: Strategic and institutional approaches. *Academy of Management Review*, 20(3), 571–610. <https://doi.org/10.2307/258788>

Teece, D. J., Pisano, G., & Shuen, A. (1997). Dynamic capabilities and strategic management. *Strategic Management Journal*, 18(7), 509–533. <https://www.jstor.org/stable/3088148>

Thapa, S., Shiwakoti, S., Shah, S. B., Adhikari, S., Veeramani, H., Nasim, M., & Naseem, U. (2025). Large language models in computational social science: Prospects, current state, and challenges. *Social Network Analysis and Mining*, 15, Article 4.  
<https://doi.org/10.1007/s13278-025-01428-9>

3D Systems Corporation. (2023). *Form 10-K*. U.S. Securities and Exchange Commission.

TTEC Holdings, Inc. (2024). *Form 10-K*. U.S. Securities and Exchange Commission.

U.S. Securities and Exchange Commission. (2024, March 18). *SEC charges two investment advisers with making false and misleading statements about their use of artificial intelligence* [Press release]. <https://www.sec.gov/news/press-release/2024-36>

Uslaner, J. D., & Coquin, A. (2025, May 30). ‘AI washing’: Regulatory and private actions to stop overstating claims. Reuters Legal News. <https://www.reuters.com/legal/legalindustry/ai-washing-regulatory-private-actions-stop-overstating-claims-2025-05-30/>



- Verrecchia, R. E. (1983). Discretionary disclosure. *Journal of Accounting and Economics*, 5(3), 179–194. [https://doi.org/10.1016/0165-4101\(83\)90011-3](https://doi.org/10.1016/0165-4101(83)90011-3)
- Wang, J., Ding, W., & Zhu, X. (2025). *Financial analysis: Intelligent financial data analysis system based on LLM-RAG* (arXiv:2405.06279). arXiv. <https://doi.org/10.48550/arXiv.2504.06279>
- Wang, J. J., & Wang, V. X. (2025). *Assessing consistency and reproducibility in the outputs of large language models: Evidence across diverse finance and accounting tasks* (arXiv:2503.16974). arXiv. <https://doi.org/10.48550/arXiv.2503.16974>
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Chi, E., Le, Q., & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems* (Vol. 35, pp. 24824–24837). Curran Associates, Inc. <https://doi.org/10.48550/arXiv.2201.11903>
- Weitzel, L., Prati, R. C., & Aguiar, R. F. (2016). The comprehension of figurative language: What is the influence of irony and sarcasm on NLP techniques? In *Sentiment analysis and ontology engineering* (pp. 49–74). Springer. [https://doi.org/10.1007/978-3-319-30319-2\\_3](https://doi.org/10.1007/978-3-319-30319-2_3)
- Welch, B. L. (1947). The generalization of ‘Student’s’ problem when several different population variances are involved. *Biometrika*, 34(1/2), 28–35. <https://doi.org/10.2307/2332510>
- Wernerfelt, B. (1984). A resource-based view of the firm. *Strategic Management Journal*, 5(2), 171–180. <https://www.jstor.org/stable/2486175>
- White, H. (1980). A heteroskedasticity-consistent covariance matrix estimator and a direct test for heteroskedasticity. *Econometrica*, 48(4), 817–838. <https://doi.org/10.2307/1912934>
- Yang, R., Mao, Y., Caporin, M., & Jiménez-Martin, J.-A. (2024). *The artificial intelligence premium* (SSRN Working Paper No. 4850201). <https://dx.doi.org/10.2139/ssrn.4850201>
- Yin, R. K. (2018). *Case study research and applications: Design and methods* (6th ed.). SAGE Publications.



You, H., & Zhang, X. (2009). Financial reporting complexity and investor underreaction to 10-K information. *Review of Accounting Studies*, 14(4), 559–586.

<https://dx.doi.org/10.2139/ssrn.985365>

Zhang, B., Yang, H., Zhou, T., Babar, M. A., & Liu, X.-Y. (2023). Enhancing financial sentiment analysis via retrieval-augmented large language models. *Proceedings of the Fourth ACM International Conference on AI in Finance*, 349–356.

<https://doi.org/10.1145/3604237.3626866>

## Appendix

### Appendix A: Examples of LLM Scoring Quotes

**Table 15:**  
*Illustrative Examples of AI Disclosure Factor Scoring*

Factor	Quote – High (A)	Quote – Generic
<b>1. Strategic Depth</b> (Mishra et al., 2022; Cockburn et al., 2018)	“We believe that AI capabilities are central to products and solutions across our markets and we have a broad technology roadmap and products targeting AI training and inference spanning cloud, edge and intelligent endpoints. We offer products that include capabilities to support AI deployment and we expect this part of our business to grow.”   (AMD 2024-01-31)	“We may face increased competition due to the rapid development and rising use of artificial intelligence (AI) and machine learning technologies. Failure to adopt and incorporate such technologies to improve productivity, manufacturing technology or support functional teams may put us at a long-term competitive disadvantage.”   (HUBB 2025-02-13)
<b>2. Disclosure Sentiment</b> (Loughran & McDonald, 2011; Feldman et al., 2009)	“We have established a leadership position in the emerging market segment for powering high-performance processors used for acceleration of applications associated with artificial intelligence (‘AI’). Our customers in the AI market segment include the leading innovators in processor and accelerator design, as well as early adopters in cloud computing and high performance computing.”   (VICR 2023-02-28)	“There are risks associated with artificial intelligence (‘AI’), any or all of which could adversely affect our business.”   (SPG 2025-02-21)
<b>3.a Risk - Own Adoption</b> (Alston & Bird, 2024; Amershi et al., 2019)	“We believe generative AI will have a significant effect on how we provide services to our clients and how we enhance the productivity of our people. As with any new technology, we are working closely with our clients and technology partners to take advantage of the benefits of AI while being mindful of its limitations, risks, and privacy concerns.” (OMC 2024-02-07)	“We are evaluating the use of artificial intelligence technologies (‘AI’) within our business and we recognize that third parties that provide services to us may independently use AI. The use of AI, which involves reliance on substantial data volumes, introduces risks that may have a material adverse effect on our business.”   (PJT 2024-02-28)
<b>3.b Risk - External</b>	“Our systems and third-party systems with which we interact, as well as	“Any future cybersecurity incidents or other similar disruptions impacting

Factor	Quote – High (A)	Quote – Generic
<b>Threats</b> (Christensen, 1997; Cockburn et al., 2018)	systems those third-parties utilize, are subject to increasingly sophisticated methods of attack, including the deployment of artificial intelligence to find and exploit vulnerabilities, such as “deep fakes,” long-term persistent attacks referred to as advanced persistent threats and the use of the IT supply chain to introduce malware.”  (MMC 2021-02-17)	Core Molding Technologies or significant customers and/or suppliers could adversely affect us.”  (CMT 2025-03-11)
<b>3.c Risk - Non-Adoption</b> (Brynjolfsson & McAfee, 2014)	“Competitors have substantially greater resources to invest in technological improvements and have invested more heavily than us, and will continue to be able to do so, in developing and adopting new technologies, which may put us at a competitive disadvantage.”  (SNV 2024-02-23)	“The risk of cybersecurity attacks may increase as artificial intelligence capabilities improve and are increasingly used to identify vulnerabilities and construct increasingly sophisticated attacks, with the possibility of additional vulnerabilities being introduced through our own use of artificial intelligence and its use by our stakeholders, including vendors and customers.” (NVT 2025-02-18)
<b>4. Forward-Looking</b> (Muslu et al., 2015)	“For retail, we will continue leveraging data science and AI to offer even better digital experiences for associates and consumers, driving conversion and efficiency. We plan to improve our online vehicle transfer experience and expand Skye’s functionality with additional data and new architecture.”  (KMX 2024-04-15)	“We are also continuing to explore the use of machine learning, AI and deep learning in our data and analytics strategies.”  (TRU 2021-02-16)
<b>5. AI Washing Index</b> (Seele & Schultz, 2022; Al Haddi, 2024)	"With the acquisitions of Mipsology SAS and Nod, Inc. in 2023, we expanded our AI software capabilities to accelerate our AI growth strategy centered on an open software ecosystem to help lower the barriers of entry for customers through developer tools, libraries and models."  (AMD, 2025-02-05)	“Some of our computer systems currently, and may in the future, incorporate AI solutions, including machine learning and generative AI tools that collect, aggregate, and analyze data to assist in the development of our services and products and in the use of internal tools that support our business.”  (ABR, 2025-02-21)
<b>6. Talent &amp; Investment</b>	"We expect spending in technology and infrastructure to increase over time	"Our ability to attract and retain executives and other qualified

Factor	Quote – High (A)	Quote – Generic
(Babina et al., 2024; Spence, 1973)	as we continue to add employees and infrastructure. These costs are allocated to segments based on usage. The increase in technology and infrastructure costs... is primarily due to an increase in spending on infrastructure and increased payroll and related costs associated with technical teams..."  (AMZN, 2024-02-02)	personnel... is important to our success. We must also attract and retain a large number of skilled employees... The competition for these employees is intense and we may not be able to hire and retain the number of employees we need in each of our locations."  (CNDT, 2024-02-21)

**Note:** This table provides real examples from corporate 10-K filings demonstrating high-scoring (A-grade) versus typical/generic disclosure language for each AI factor. Company tickers and filing dates are shown in parentheses. These examples illustrate the qualitative differences captured by the grading framework.

## Appendix B: Complete LLM Prompt

The prompt template below is used to instruct the LLM for analysing the 10-K filing excerpts. The {document\_text} placeholder is filled with the combined "Risk Factors" and "Management's Discussion and Analysis" sections for each company.

You are an expert financial analyst and AI strategist, specializing in evaluating corporate disclosures for insights into a company's engagement with Artificial Intelligence (AI). Your task is to meticulously analyse the provided text from a company's 10-K filing (specifically excerpts from "Item 1. Risk Factors" and "Item 7. Management's Discussion and Analysis") and provide a structured JSON output assessing the company's AI preparedness based on the defined factors and criteria.

**\*\*DOCUMENT CONTEXT:\*\***

The provided text below, enclosed in triple backticks (````), contains excerpts from "Item 1. Risk Factors" and "Item 7. Management's Discussion and Analysis (MD&A)" of a single company's 10-K report.

{document\_text}

**\*\*INSTRUCTIONS:\*\***

- Carefully read the entire provided text.
- For each of the AI Assessment Factors listed below, evaluate the text based on the defined Buckets & Criteria.
- Your entire response **MUST** be a single, valid JSON object. Do not include any explanatory text before or after the JSON object. The JSON object **MUST** conform to the Pydantic schema `CompanyAIAalysis`.

4. For each factor, you MUST include in the JSON:

\* `bucket\_assessment`: The exact string for the category you've assigned.

**\*\*CRITICAL:** This value MUST be one of the precise string values listed in the Buckets & Criteria for that factor (e.g., for AI Strategic Depth, it must be exactly "A (Core Strategic Enabler)" or "B (Significant Operational/Product Integration)", etc. Do NOT use abbreviations or variations).

\* `chain\_of\_thought\_reasoning`: A brief explanation (1-2 sentences) of your thought process for arriving at the bucket assessment, referencing the criteria and evidence.

\* `supporting\_evidence`: Key verbatim quotes (or concise summaries of multiple related statements) from the text that justify your assessment. Limit to 2-3 key pieces of evidence. If no direct evidence is found for a factor, you MUST provide an empty list `[]` for this field.

5. **\*\*CRITICAL FOR `ai\_risk\_disclosure` FIELD\*\*:** The `ai\_risk\_disclosure` field in the JSON output MUST be an object containing three specific nested assessment objects as its keys: `risk\_from\_own\_ai\_adoption`, `risk\_from\_external\_ai\_threats`, and `risk\_of\_not\_adopting\_ai\_or\_competitive\_ai\_threat`. Each of these three nested objects must follow the `BaseAssessment` structure (i.e., have `bucket\_assessment`, `chain\_of\_thought\_reasoning`, and `supporting\_evidence`). Refer to the example JSON structure below.

6. If information for any other factor or sub-factor (within `ai\_risk\_disclosure`) is not present in the text, use the appropriate "No Mention" bucket assessment for that factor (e.g., "E (No Mention of Own AI Strategy/Adoption)" or "C (No Mention of Own AI Adoption Risk)"). For `supporting\_evidence` in such cases, provide an empty list `[]`.

\***IMPORTANT**\*: Ensure ALL main assessment factor objects (ai\_strategic\_depth, ai\_disclosure\_sentiment, ai\_risk\_disclosure, forward\_looking\_ai\_statements, ai\_washing\_hype\_index, ai\_talent\_and\_investment\_focus) are ALWAYS present as keys in your JSON output. If a factor has no relevant information, its `bucket\_assessment` should reflect its defined "No Mention" or equivalent lowest grade for that factor's rubric, and `supporting\_evidence` should be an empty list. Do not omit these top-level keys from the JSON.

7. For the `company\_identifier` field, extract the company name from the "Company Name (if known):" line in the DOCUMENT CONTEXT. If it's "Unknown Company" or not clearly identifiable, set it to null.

8. For the `overall\_ai\_preparedness\_summary\_cot` field, you MUST provide a 1-2 sentence summary of the company's overall AI posture based on your analysis of all factors.

9. For the `key\_ai\_related\_terminology\_found` field, you MUST list any explicit AI-related terms found in the text (e.g., "artificial intelligence", "machine learning", "NLP", "generative AI"). If no such terms are found, provide an empty list `[]`.

**\*\*AI ASSESSMENT FACTORS, BUCKETS & CRITERIA:\*\***

**\*\*1. AI Strategic Depth:\*\***

\* **Concept:** How deeply and explicitly AI is integrated into the company's core business strategy and future plans as evidenced by specific examples, resource allocation, or stated strategic importance. Evaluation should consider the scale and impact of AI initiatives, not

just mentions.

\* \*Buckets & Criteria:\*

\* "A (Core Strategic Enabler)": AI is explicitly and repeatedly identified as a core component of the company's overall strategy, with multiple concrete, specific examples of current or planned AI applications that are central to key products, services, or operational efficiency. Evidence of significant investment, dedicated AI leadership, or AI-driven competitive advantages.

\* "B (Significant Operational/Product Integration)": AI is discussed as being integrated into significant aspects of operations or specific product lines with some concrete examples, but may not be portrayed as a universal, core strategic driver for the entire enterprise. Impact is identifiable but perhaps more localised than 'A'.

\* "C (Emerging/Exploratory Mentions)": AI is mentioned in the context of emerging opportunities, ongoing exploration, pilot projects, or general capabilities the company is developing or monitoring. Specific strategic integration or widespread impact is not yet evident.

\* "D (Superficial/Generic Mentions OR Risk of Non-Adoption Only)": AI mentions are high-level, generic, or boilerplate, lacking specific examples or connection to core strategy. Alternatively, AI is only mentioned in the context of the \*risk of not adopting it\* or competitive pressures, without detailing the company's own strategic AI adoption.

\* "E (No Mention of Own AI Strategy/Adoption)": There is no substantive discussion of the company's own AI strategy, adoption, or integration efforts. (Note: General discussion of AI as an external market trend or risk factor without connection to the company's own plans falls here).

## **\*\*2. AI Disclosure Sentiment:\*\***

\* \*Concept:\* The overall tone of the company's \*own statements\* regarding its AI initiatives, capabilities, and outlook. This excludes sentiment about external AI risks not directly tied to the company's actions.

\* \*Buckets & Criteria:\*

\* "A (Clearly Positive)": Disclosures about the company's AI are predominantly optimistic, highlighting successes, opportunities, competitive advantages, and positive impacts. Language used is confident and forward-looking regarding their AI.

\* "B (Mostly Positive/Balanced)": Disclosures are generally positive but may include some neutral statements or acknowledge minor challenges that are framed as manageable. The overall impression is favourable.

\* "C (Neutral/Factual)": Disclosures about AI are presented in a purely factual, objective manner, without significant positive or negative framing. Focus is on description rather than promotion or concern.

\* "D (Cautious/Mixed)": Disclosures express caution, highlight significant uncertainties, or present a mixed picture of benefits and challenges regarding the company's own AI. May include discussion of ongoing struggles or unproven AI ventures.

\* "E (Negative OR Not Applicable - No Own AI Discussion)": Disclosures about the company's own AI are predominantly negative, focusing on failures, significant unresolved problems, or downsides. Also applies if there is no substantive discussion of the company's own AI initiatives to assess sentiment from.

## **\*\*3. AI Risk Disclosure:\*\* (This section requires careful attention to the nested structure)**

\* \*Concept:\* How thoroughly and specifically the company acknowledges and discusses potential risks associated with AI. This factor itself is an object in the JSON, containing the following three sub-assessments:

**\*\*\*3a. `risk\_from\_own\_ai\_adoption`:** Risks arising directly from the company's own development, deployment, or use of AI systems (e.g., AI system failure, data bias, model inaccuracies, ethical concerns, implementation challenges, security of own AI systems, intellectual property issues related to own AI).

**\*\*Buckets & Criteria for 3a:**

**"A (Detailed & Specific Discussion)":** The company provides a detailed and specific discussion of one or more risks related to its own AI adoption, potentially outlining specific scenarios or impacts.

**"B (General Mention of Own AI Adoption Risk)":** The company makes a general mention of risks associated with its own AI adoption, without detailing specific scenarios or impacts extensively.

**"C (No Mention of Own AI Adoption Risk)":** The company does not mention any risks specifically arising from its own AI adoption or systems.

**\*\*\*3b. `risk\_from\_external\_ai\_threats`:** Risks posed by AI used by external actors or general AI advancements in the environment (e.g., AI used in cyberattacks by others against the company, AI-driven misinformation campaigns affecting the company, AI capabilities of competitors if not framed as a 'failure to adopt' risk for the company itself).

**\*\*Buckets & Criteria for 3b:**

**"A (Detailed & Specific Discussion)":** The company provides a detailed and specific discussion of one or more external AI threats.

**"B (General Mention of External AI Threat)":** The company makes a general mention of external AI threats.

**"C (No Mention of External AI Threat)":** The company does not mention external AI threats.

**\*\*\*3c. `risk\_of\_not\_adopting\_ai\_or\_competitive\_ai\_threat`:** Risks associated with the company failing to adopt AI, keep pace with AI developments, or risks from competitors leveraging AI more effectively.

**\*\*Buckets & Criteria for 3c:**

**"A (Explicitly Discussed)":** The company explicitly discusses the risk of falling behind competitors or negative impacts from not adopting/leveraging AI.

**"B (Implicitly Suggested)":** The risk of not adopting AI or competitive AI threats is implicitly suggested (e.g., by discussing competitors' AI advancements without stating it as a direct risk, or general statements about needing to innovate).

**"C (No Mention of Non-Adoption/Competitive AI Risk)":** The company does not mention risks related to its own non-adoption of AI or competitive AI threats.

**\*\*4. Forward-Looking AI Statements:**

**\*Concept:** The extent and specificity of discussion about future plans, intentions, or expectations regarding AI.

**\*\*Buckets & Criteria:**

**"A (Specific & Detailed Future Plans)":** The company provides specific and detailed statements about future AI plans, projects, investments, or expected milestones (e.g., "launching product X with AI feature Y in Q3," "investing Z amount in AI R&D over next year").

**"B (General Future Intent)":** The company expresses general intentions or expectations regarding future AI use or development, without providing specific details, timelines, or quantifiable targets (e.g., "we plan to explore AI," "AI will be important for our future").

**"C (Implicit Future Focus Only)":** Future AI focus is only implied through discussion of ongoing AI initiatives that naturally have a future dimension, but no explicit forward-looking statements are made.

\* "D (No Mention of Future AI Plans)": There is no mention of future plans, intentions, or expectations specifically related to AI.

**\*\*5. "AI Washing" Hype Index:\*\*** (Factor is only applicable if positive claims about own AI are made. If no positive claims, or only risks are discussed, use "E")

\* **\*Concept:** Degree to which a company's positive AI claims appear substantive and grounded in verifiable details versus being exaggerated, vague, or aspirational without clear evidence.

\* **\*Buckets & Criteria:**

\* "A (Substantive & Grounded)": Positive AI claims are consistently backed by specific examples, data, achievements, or clear explanations of functionality and impact. Claims appear realistic and verifiable.

\* "B (Mostly Substantive)": Most positive AI claims are substantive, but there may be occasional instances of slightly vague or aspirational language. Overall, claims seem credible.

\* "C (Mixed - Some Substance, Some Hype)": There's a mix of substantive claims and more hyped, vague, or aspirational statements about AI. Difficult to clearly distinguish hype from reality throughout.

\* "D (Mostly Hype)": Positive AI claims are predominantly vague, aspirational, or use buzzwords without concrete evidence or specific details. Claims seem exaggerated or difficult to substantiate.

\* "E (Pure Hype OR Not Applicable - No Positive AI Claims)": Positive AI claims are entirely generic, buzzword-laden, and lack any substance, or the company makes no positive claims about its own AI to evaluate for hype (e.g., only discusses AI risks or general market trends).

**\*\*6. AI Talent & Investment Focus:\*\***

\* **\*Concept:** Explicit discussion of focus on acquiring, developing, or retaining AI talent, or specific investments (financial, infrastructural) in AI.

\* **\*Buckets & Criteria:**

\* "A (Explicit & Significant Focus)": The company explicitly and significantly discusses efforts to acquire/develop AI talent (e.g., hiring initiatives, training programs) AND/OR details specific, material investments in AI (e.g., R&D spending, acquisitions, infrastructure development).

\* "B (General Mention of AI Talent/Investment)": The company makes general mentions of the importance of AI talent or investment in AI, but without providing specific details on initiatives, scale, or amounts.

\* "C (Implicit Focus Only)": Focus on AI talent or investment is only implicitly suggested (e.g., by discussing AI projects that would logically require talent/investment, but without explicitly stating this focus).

\* "D (No Mention of AI Talent/Investment)": There is no mention of a focus on AI talent acquisition, development, or specific AI investments.

**\*\*EXPECTED JSON OUTPUT STRUCTURE (Example for LLM guidance, actual structure enforced by Pydantic):\*\***

```
```json
{
  "company_identifier": "EXTRACT_COMPANY_NAME_IF_POSSIBLE_ELSE_NULL",
  "overall_ai_preparedness_summary_cot":
  "LLM_GENERATED_1_2_SENTENCE_SUMMARY_OF_OVERALL_AI_POSTURE_B
  ASSED_ON_ALL_FACTORS",
```



```

"ai_strategic_depth": {
  "bucket_assessment": "A (Core Strategic Enabler)",
  "supporting_evidence": ["quote1 from text about core AI strategy", "quote2 detailing
specific AI application"],
  "chain_of_thought_reasoning": "The company clearly states AI is central to its strategy
and provides multiple concrete examples of its application in key products, aligning with
criteria for 'A'."
},
"ai_disclosure_sentiment": {
  "bucket_assessment": "B (Mostly Positive/Balanced)",
  "supporting_evidence": ["quote expressing optimism about AI initiative"],
  "chain_of_thought_reasoning": "The language used when discussing their AI is generally
positive, highlighting benefits, though some statements are factual, fitting 'B'."
},
"ai_risk_disclosure": {
  "risk_from_own_ai_adoption": {
    "bucket_assessment": "B (General Mention of Own AI Adoption Risk)",
    "supporting_evidence": ["The company mentions risks related to 'implementation of new
technologies' including AI."],
    "chain_of_thought_reasoning": "A general risk of adopting new technologies including
AI is mentioned, but not highly specific to AI, fitting 'B'."
  },
  "risk_from_external_ai_threats": {
    "bucket_assessment": "C (No Mention of External AI Threat)",
    "supporting_evidence": [],
    "chain_of_thought_reasoning": "No specific external AI threats like AI-cyberattacks by
others were identified in the text."
  },
  "risk_of_not_adopting_ai_or_competitive_ai_threat": {
    "bucket_assessment": "A (Explicitly Discussed)",
    "supporting_evidence": ["'Failure to innovate or adopt new technologies such as AI could
harm our competitive position.'"],
    "chain_of_thought_reasoning": "The company explicitly states that not adopting AI poses
a competitive risk, aligning with 'A'."
  }
},
"forward_looking_ai_statements": {
  "bucket_assessment": "A (Specific & Detailed Future Plans)",
  "supporting_evidence": ["'We plan to invest $X million in AI R&D in the next fiscal
year.'", "'A new AI-powered feature for product Y is scheduled for Q4.'"],
  "chain_of_thought_reasoning": "Specific investment figures and product launch timelines
for AI are provided, meeting criteria for 'A'."
},
"ai_washing_hype_index": {
  "bucket_assessment": "A (Substantive & Grounded)",
  "supporting_evidence": ["The claims about AI improving efficiency are supported by
specific operational examples mentioned earlier."],
  "chain_of_thought_reasoning": "Positive AI claims are linked to concrete examples and
described functionalities, indicating substance, hence 'A'."
},

```

```

"ai_talent_and_investment_focus": {
  "bucket_assessment": "B (General Mention of AI Talent/Investment)",
  "supporting_evidence": ["“We are focused on attracting skilled personnel in technology areas like AI.”"],
  "chain_of_thought_reasoning": "A general statement about attracting AI talent is made, but lacks specifics on initiatives or scale, fitting ‘B’."
},
"key_ai_related_terminology_found": ["artificial intelligence", "machine learning", "automation"]
}

```

Appendix C: Case Study Analysis

B.1 Introduction and Case Study Overview

This appendix analyses five additional AI Hyper firms that support the Chegg case from the main text. These cases show that the same behavioural patterns appear across all three sectors with significant negative relationships between AI disclosure and stock returns.

Table 16:

Overview of Selected Case Studies: High AI Disclosure and Negative Returns

Company	Sector	AI Score	12m Return	Key Issue	Analysis Location
Chegg (CHGG)	Consumer Discretionary	36.33	-72.99%	Core Business Disrupted by AI	Main Text (6.4.1)
3D Systems (DDD)	Industrials	29.67	-62.04%	Failed AI Investment	Appendix C.2.1
Omnicell (OMCL)	Health Care	30.33	-60.14%	AI Creates New Risks	Appendix C.2.2
Corvel (CRVL)	Health Care	24.67	-64.23%	Regulatory Pressure	Appendix C.2.3
TTEC Holdings (TTEC)	Industrials	31.33	-72.99%	AI Threatens Core Business	Appendix C.3.1
iRobot (IRBT)	Consumer Discretionary	31.33	-72.99%	Losing Key AI Talent	Appendix C.3.2

Note: This table summarises the firms selected for case study analysis. These firms received high ‘AI Cumulative Scores’ but experienced severely negative subsequent 12-month excess returns. ‘Key Issue’ summarises the primary business challenge identified in the analysis.

C.2 Multi-Sector Case Analysis

C.2.1 Case Study: 3D Systems Corp. (DDD) - Industrials Sector

- 2022 Form 10-K, Filed: 16 March 2023
- AI Score: 29.67 (15th highest)
- 12-Month Return: -62.04% (7th worst)

External Context and Failed Investment Pressures

3D Systems' high AI disclosure came with disappointing Q4 2022 earnings. The company missed analyst estimates, showed declining revenues, and reported significant losses. Rising interest rates put extra pressure on unprofitable technology companies to show clear paths to profitability.

This created strong agency pressures for management. They had made large AI investments that failed to generate expected returns. Managers faced the challenge of explaining poor performance while maintaining credibility about their strategy.

Critical Evidence: Failed Investment Signalling

"On November 1, 2021, we acquired Oqton for \$187.8 million. Oqton is a software company that creates an intelligent, cloud-based MOS platform... leveraging artificial intelligence, and machine learning technologies... However, the acquisition's impact on the Company's financial position, results of operations and cash flows has been dilutive."

(3D Systems Corporation, 2023)

Theoretical Analysis

This disclosure shows clear **Failed Investment Signalling** (Spence, 1973; Ross, 1977). Management explicitly admits their \$187.8 million AI acquisition "has been dilutive." At the same time, they use positive language about "intelligent, cloud-based platforms" and "artificial intelligence and machine learning technologies."

This creates what signalling theory calls "contradictory signals" (Spence, 1973). The company uses extensive positive AI language while admitting concrete value destruction. This pattern shows management trying to maintain credibility about AI strategy despite clear execution failure. This contradicts signalling theory's requirement that credible signals must be costly and hard to fake (Ross, 1977).

Agency Theory implications are severe (Jensen & Meckling, 1976). Management faced a substantial investment that destroyed shareholder value. They confronted classic pressures to justify strategic decisions while maintaining stakeholder confidence. The

extensive AI discourse represents impression management designed to frame failure within technological sophistication narratives (Merkel-Davies & Brennan, 2007).

Markets rationally interpreted this pattern as evidence of poor strategic execution and inadequate accountability. The combination of positive AI language and explicit value destruction signals fundamental problems with management's capital allocation decisions. This aligns with research on managerial overinvestment in fashionable technologies (Malmendier & Tate, 2015; Jensen, 1986).

C.2.2 Case Study: Omnicell, Inc. (OMCL) - Health Care Sector

- *2021 Form 10-K, Filed: 25 February 2022*
- *AI Score: 30.33 (7th highest)*
- *12-Month Return: -60.14% (7th worst)*

External Context and Agency Pressures

Omnicell's extensive AI disclosure came with severely disappointing Q4 2021 earnings. The company missed analyst estimates, faced rising supply chain costs, and gave forward guidance well below market expectations. This created classic agency pressures where management needed to offset negative financial news with strategic narratives (Jensen & Meckling, 1976; Healy & Palepu, 2001).

Critical Evidence: Legitimacy Theory and Proprietary Costs

"An effective attack on our solutions could disrupt the proper functioning of our solutions, allow unauthorized access to sensitive and confidential information of our customers (including protected health information)... These risks are likely to increase as we continue to grow our cloud-based offerings, including in support of the industry vision of the Autonomous Pharmacy, and as we receive, store, and process more of our customers' data." (Omnicell, Inc., 2022)

Theoretical Analysis

This disclosure shows how **Legitimacy Theory** explains negative market reactions to AI announcements (Suchman, 1995). Rather than positioning "Autonomous Pharmacy" as a competitive advantage, management frames their most ambitious AI initiative primarily in terms of escalating cybersecurity and privacy risks. This suggests management views AI expansion as increasing rather than decreasing vulnerability. This contradicts traditional signalling theory expectations that strategic investments should signal strength (Spence, 1973).

Proprietary Cost Theory implications are severe (Verrecchia, 1983). By explicitly linking their strategic "Autonomous Pharmacy" vision to increased data security risks, the company provides competitors and stakeholders with specific information about potential vulnerabilities. Markets rationally interpret such admissions as evidence of poor strategic positioning that will harm future competitive advantages (Admati & Pfleiderer, 2000).

The timing amplifies these effects through what **Impression Management Theory** predicts (Merkl-Davies & Brennan, 2007). When extensive AI disclosure accompanies earnings disappointments, sophisticated investors correctly identify the pattern as defensive impression management rather than genuine strategic communication.

C.2.3 Case Study: Corvel Corp. (CRVL) - Health Care Sector

- 2023 Form 10-K, Filed: 24 May 2024
- AI Score: 24.67 (36th highest)
- 12-Month Return: -64.23% (3rd worst)

External Context and Self-Selection

Corvel's extensive AI disclosure occurred alongside significant insider selling by executives. This created contradictory signals about management confidence. This combination shows self-selection bias where defensive disclosure accompanies personal risk mitigation by management (Li & Prabhala, 2007; Heckman, 1979).

Critical Evidence: Regulatory Disruption and Existential Threat

"Our future success depends, in part, on our ability to anticipate and respond effectively to the threat and opportunity presented by digital disruption, 'big data' and data analytics... We may not be successful in anticipating or responding to these developments on a timely and cost-effective basis and our ideas may not be accepted in the marketplace... There has also been increased regulatory scrutiny of the use of 'big data' techniques, machine learning, and artificial intelligence." (Corvel Corporation, 2024)

Theoretical Analysis

This case powerfully shows **Business Model Disruption Theory** (Christensen, 1997). The language reveals management's perception that AI represents both "threat and opportunity." However, the extensive discussion of potential failure ("may not be successful," "ideas may not be accepted") signals genuine uncertainty about the company's ability to adapt to technological change (Teece et al., 1997).

The **Regulatory Uncertainty** dimension adds another layer of risk that **Legitimacy Theory** predicts will create negative stakeholder reactions (Suchman, 1995). Management

explicitly acknowledges that "increased regulatory scrutiny" of AI/ML creates compliance risks and potential operational restrictions. This represents mandatory disclosure of factors that could materially impair the business model (Healy & Palepu, 2001).

Self-Selection Bias becomes apparent when considering concurrent insider selling (Roy, 1951; Li & Prabhala, 2007). Management's words emphasize caution and regulatory risk while their actions show personal lack of confidence in company prospects. This behavioural pattern provides rational justification for negative market reactions. Investors correctly interpret the combination as evidence that management belongs to a cohort facing genuine competitive threats rather than strategic opportunities (Akerlof, 1970).

C.3 Additional Sector Cases

C.3.1 Case Study: TTEC Holdings, Inc. (TTEC) - Industrials Sector

- 2023 Form 10-K, Filed: 29 February 2024
- AI Score: 31.33 (6th highest)
- 12-Month Return: -72.99% (Worst performance)

External Context and Existential Disruption

TTEC represents the most severe case of business model disruption. The company's core human-led customer service offerings face direct replacement by AI automation (Christensen, 1997). The filing accompanied disastrous Q4 2023 earnings with drastically reduced guidance and major client losses.

Critical Evidence: Existential Business Model Threat

"These technology innovations have potential to significantly disrupt our business, reduce business volumes and related revenue unless we are successful in adapting... Deploying AI too rapidly without appropriate controls may result in poor client adoption, liability, and reputational damage." (TTEC Holdings, Inc., 2024)

Theoretical Analysis

This case shows **Existential Disruption Theory** in its purest form (Christensen, 1997; Teece et al., 1997). TTEC faces what we call "mandatory negative signalling." Regulatory requirements force the company to acknowledge that AI "has potential to significantly disrupt our business, reduce business volumes and related revenue" (Healy & Palepu, 2001).

The 10-K report reveals management's challenging strategic position. They must show competence by demonstrating an understanding of AI threats. At the same time, they need to acknowledge that their business model may become unviable. This creates what signalling

theory predicts: when firms attempt to signal contradictory messages, markets interpret the confusion as confirmation of underlying weakness (Spence, 1973; Akerlof, 1970).

Self-Selection Mechanism becomes apparent when considering why TTEC discusses AI so extensively (Roy, 1951; Li & Prabhala, 2007). The company must engage in detailed AI disclosure precisely because AI represents an existential threat requiring regulatory acknowledgment. This creates unavoidable negative signalling where disclosure intensity correlates with business model vulnerability rather than strategic strength.

The market's severely negative reaction (-72.99% return) reflects rational assessment. Companies forced to acknowledge existential AI threats belong to a cohort facing potential extinction rather than transformation opportunity. This aligns with **Market for Lemons** theory where adverse selection leads to systematic undervaluation of firms revealing vulnerability (Akerlof, 1970).

C.3.2 Case Study: iRobot Corp. (IRBT) - Consumer Discretionary Sector

- 2022 Form 10-K, Filed: 14 February 2023
- AI Score: 31.33 (8th highest)
- 12-Month Return: -72.99% (Worst performance)

External Context and Core Competency Loss

iRobot's case shows how AI disclosure can reveal operational vulnerabilities that become critical when external circumstances change. While the failed Amazon acquisition was the direct cause of stock decline, the AI disclosure revealed underlying talent retention problems. These problems explained why the acquisition was necessary rather than opportunistic (Prahalad & Hamel, 1990).

Critical Evidence: Core Competency Erosion

"Competition for hiring these employees is intense, especially with regard to engineers with high levels of experience in... artificial intelligence... we have experienced increased employee turnover... we expect to continue to experience increased employee turnover in the future." (iRobot Corporation, 2023)

Theoretical Analysis

This disclosure shows **Core Competency Erosion Theory** (Prahalad & Hamel, 1990; Barney, 1991). For a technology company, admitting inability to retain AI talent signals fundamental operational weakness. This undermines the firm's core competitive advantages. The extensive discussion of AI strategy combined with explicit admission of talent retention

problems suggests management lacks the human capital necessary to execute their stated strategic vision.

Following **Proprietary Cost Theory** (Verrecchia, 1983), companies typically conceal information about competitive advantages to prevent exploitation by rivals. iRobot's willingness to reveal AI talent retention problems suggests that concealment posed greater legal or regulatory risks than disclosure. This implies that underlying challenges are severe and potentially visible to stakeholders (Admati & Pfleiderer, 2000).

When the Amazon acquisition failed, these disclosed weaknesses became critical determinants of independent viability. Markets possessed clear information suggesting iRobot lacked internal capabilities necessary for success in an AI-driven competitive landscape. This aligns with **Resource-Based View** predictions about the importance of human capital for sustained competitive advantage (Barney, 1991; Wernerfelt, 1984).

Agency Theory implications reflect management's position facing talent exodus (Jensen & Meckling, 1976). The extensive AI disclosure represents impression management designed to show strategic awareness while preparing stakeholders for potential performance impacts from talent losses (Merkl-Davies & Brennan, 2007).

C.4 Summary of Case Study Findings

The analysis of six AI Hyper firms reveals systematic behavioural patterns across all three significant sectors. Each case shows the same fundamental sequence: underlying business pressures create agency incentives for extensive AI disclosure. However, this strategy backfires by revealing vulnerabilities rather than strength.

Health Care Sector: Omnicell and Corvel face regulatory and competitive pressures that force defensive AI disclosure. Omnicell frames AI expansion in terms of increased cybersecurity risks. Corvel explicitly acknowledges potential failure in AI adaptation alongside regulatory scrutiny.

Industrials Sector: 3D Systems and TTEC represent different patterns - failed investment justification versus existential business model threats. 3D Systems explicitly admits their \$187.8M AI acquisition was "dilutive." TTEC acknowledges AI may "significantly disrupt our business, reduce business volumes."

Consumer Discretionary Sector: Chegg (analysed in main text) and iRobot demonstrate challenges related to the decline of their business models and weakening of their core competencies. Chegg admits competitors may develop "superior" technologies. iRobot reveals critical AI talent retention problems.

All cases show self-selection bias where extensive AI disclosure correlates with business vulnerability rather than strategic strength. The consistent negative market reactions validate investor sophistication in distinguishing genuine strategic communication from impression management driven by agency costs.

C.5 Cross-Case Theoretical Synthesis

The detailed analysis of six AI Hyper firms across three sectors reveals that extensive AI disclosure consistently occurs under conditions of business pressure rather than strategic strength. Every case shows the same sequence: underlying business challenges create agency pressures for management, leading to impression management strategies that create negative signalling effects through self-selection bias.

C.5.1 Unified Theoretical Framework

The cases validate five interconnected theoretical mechanisms:

Failed Investment Signalling (Spence, 1973; Ross, 1977): Companies use extensive AI discourse to justify value-destroying investments. This creates contradictory signals that markets correctly interpret as evidence of poor capital allocation.

Business Model Disruption (Christensen, 1997; Teece et al., 1997): This occurs when firms are threatened by AI-driven changes. These firms must disclose regulatory information that reveals serious risks to their survival.

Core Competency Erosion (Prahalad & Hamel, 1990; Barney, 1991): Companies losing essential capabilities use AI narratives to mask deteriorating advantages while inadvertently revealing vulnerabilities through proprietary cost disclosures.

Legitimacy Theory (Suchman, 1995; Freeman, 1984): AI disclosure signals potential stakeholder concerns about ethics, safety, or regulatory compliance that create legitimacy threats rather than competitive advantages.

Self-Selection Bias (Roy, 1951; Li & Prabhala, 2007; Akerlof, 1970): Companies systematically self-select into extensive AI disclosure when facing business pressures. This creates adverse selection where disclosure intensity correlates with underlying weakness rather than strength.

C.5.2 Market Efficiency and Rational Responses

The consistent negative market reactions across all cases suggest sophisticated investor interpretation of behavioural patterns rather than irrational punishment of transparency. Markets appear capable of distinguishing genuine strategic communication from impression


management driven by agency costs (Jensen & Meckling, 1976), business model threats (Christensen, 1997), or operational failures (Prahalad & Hamel, 1990).

The subsequent performance of these companies generally validates initial negative market assessments. This confirms that extensive AI disclosure often accompanies rather than contradicts fundamental business weaknesses. This outcome consistency shows market efficiency in processing complex behavioural signals embedded within corporate disclosures (Fama, 1970).

Appendix D: Additional Analysis Regression Tables & Figures

Figure 7:

Binned Scatterplots of 12-Month Excess Returns against AI Factor Grades



Note: This figure presents binned scatterplots showing the average 12-month excess return for firms grouped by their numeric score for “AI Washing”, “Strategic Depth”, and “Talent & Investment” factors. The sample period is 2020–2024.

Table 17:

Progressive Regressions of 12-Month Excess Returns on the Talent & Investment Factor

Variable	Base	Model 1	Model 2	Model 3
Talent & Investment	-0.0700*** (0.0134)	-0.0657*** (0.0135)	-0.0609*** (0.0136)	-0.0568*** (0.0136)
Log Market Cap		-0.0070 (0.0046)	-0.0030 (0.0050)	-0.0076 (0.0050)
Price-to-Book			-0.0042** (0.0019)	-0.0053*** (0.0019)
ROA				1.2031*** (0.3995)
R^2	0.0446	0.0454	0.047	0.0516
Adj. R^2	0.04	0.0406	0.0418	0.0462

Variable	Base	Model 1	Model 2	Model 3
Observations	3,174	3,174	3,174	3,171

Note: This table presents panel regressions of 12-month excess returns on the ‘Talent & Investment’ factor. The columns show the progressive inclusion of book controls. All models include GICS sector- and year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 18:
Progressive Regressions of 12-Month Excess Returns on the Composite AI Score

Variable	Base	FF3	FF5	FF6
Cum AI Score	-0.0114*** (0.0020)	-0.0107*** (0.0020)	-0.0106*** (0.0020)	-0.0106*** (0.0020)
Market Beta		-0.0013 (0.0189)	-0.0049 (0.0193)	-0.0048 (0.0193)
Size Beta		-0.0041 (0.0111)	-0.0016 (0.0113)	-0.0016 (0.0117)
Value Beta		-0.0175 (0.0110)	-0.0141 (0.0122)	-0.0141 (0.0123)
Profitability Beta			-0.0060 (0.0106)	-0.0060 (0.0107)
Investment Beta			0.0076 (0.0084)	0.0076 (0.0084)
Momentum Beta				0.0005 (0.0134)
R^2	0.0459	0.0386	0.0393	0.0393
Adj. R^2	0.0414	0.0328	0.0329	0.0326
Observations	3,174	3,031	3,031	3,031

Note: This table presents panel regressions of 12-month excess returns on the ‘AI Cumulative Score’. The columns show the progressive inclusion of Fama-French factor controls. All models include GICS sector- and year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 19:
Regressions of 9-Month Excess Returns on Individual AI Disclosure Factors

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
AI Washing Index	-0.0323*** (0.0092)	-0.0376*** (0.0089)

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
Disclosure Sentiment	-0.0254*** (0.0084)	-0.0320*** (0.0081)
Forward-Looking	-0.0377*** (0.0115)	-0.0457*** (0.0113)
Strategic Depth	-0.0202*** (0.0068)	-0.0250*** (0.0066)
Talent & Investment	-0.0480*** (0.0111)	-0.0551*** (0.0107)
Risk External	0.0051 (0.0187)	-0.0025 (0.0185)
Risk Non - Adoption	-0.0470*** (0.0151)	-0.0569*** (0.0148)
Risk Own Adoption	-0.0007 (0.0199)	-0.0451** (0.0192)
AI Cumulative Score	-0.0068*** (0.0018)	-0.0082*** (0.0017)
Average R^2	0.0966	0.0811
Average Observations	3,217	3,159

Note: This table presents coefficients from separate panel regressions of 9-month excess returns on each AI factor (18 total regressions: 9 factors $\times$ 2 specifications). Column (1) includes book controls. Column (2) includes the six Fama-French-Carhart factor exposures. All models include GICS sector- and year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 20:
Regressions of 6-Month Excess Returns on Individual AI Disclosure Factors

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
AI Washing Index	-0.0173** (0.0073)	-0.0216*** (0.0071)
Disclosure Sentiment	-0.0139** (0.0063)	-0.0186*** (0.0061)
Forward-Looking	-0.0246*** (0.0089)	-0.0297*** (0.0087)
Strategic Depth	-0.0111** (0.0053)	-0.0148*** (0.0050)

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
Talent & Investment	-0.0286*** (0.0089)	-0.0328*** (0.0085)
Risk External	-0.0020 (0.0135)	-0.0040 (0.0133)
Risk Non - Adoption	-0.0281** (0.0112)	-0.0347*** (0.0112)
Risk Own Adoption	-0.0003612 (0.014)	-0.0307** (0.0137)
AI Cumulative Score	-0.0040*** (0.0014)	-0.0049*** (0.0013)
Average R^2	0.0838	0.0757
Average Observations	3,267	3,251

Note: This table presents coefficients from separate panel regressions of 6-month excess returns on each AI factor (18 total regressions: 9 factors $\times$ 2 specifications). Column (1) includes book controls. Column (2) includes the six Fama-French-Carhart factor exposures. All models include GICS sector- and year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 21:
Regressions of 3-Month Excess Returns on Individual AI Disclosure Factors

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
AI Washing Index	-0.0115** (0.0055)	-0.0137** (0.0054)
Disclosure Sentiment	-0.0112** (0.0047)	-0.0128*** (0.0046)
Forward-Looking	-0.0154** (0.0065)	-0.0179*** (0.0063)
Strategic Depth	-0.0084** (0.0039)	-0.0100*** (0.0038)
Talent & Investment	-0.0209*** (0.0069)	-0.0240*** (0.0066)
Risk External	-0.0050 (0.0089)	-0.0024 (0.0090)
Risk Non - Adoption	-0.0241*** (0.0081)	-0.0261*** (0.0082)

AI Factor	Book Control Variables Model	Fama-French 6-Factor Model
Risk Own Adoption	-0.0087 (0.0100)	-0.0112 (0.0099)
AI Cumulative Score	-0.0029*** (0.0011)	-0.0033*** (0.0010)
Average R^2	0.1397	0.1304
Average Observations	3,322	3,317

Note: This table presents coefficients from separate panel regressions of 3-month excess returns on each AI factor (18 total regressions: 9 factors $\times$ 2 specifications). Column (1) includes book controls. Column (2) includes the six Fama-French-Carhart factor exposures. All models include GICS sector- and year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 22:
Sector-Level Regressions of 12-Month Excess Returns on the Strategic Depth Factor

Sector	(Controls Model)			(FF6 Model)			N
	Coefficient	P-Value	R	Coefficient	P-Value	R	
Communication	-0.15% (5.47)	0.978	0.125	-2.90% (5.13)	0.572	0.256	102
Consumer Discretionary	-6.61%*** (1.90)	4.94E-04	0.076	-7.27%*** (1.87)	1.05E-04	0.063	489
Consumer Staples	-1.15% (2.65)	0.665	0.095	-1.99% (2.61)	0.446	0.044	187
Energy	8.98% (7.67)	0.241	0.105	10.66% (6.95)	0.125	0.184	131
Financials	1.36% (2.82)	0.628	0.121	2.15% (2.48)	0.385	0.125	403
Health Care	-4.91%** (2.01)	0.014	0.032	-2.59% (1.94)	0.183	0.042	276
Industrials	-6.57%*** (1.45)	6.14E-06	0.06	-6.13%*** (1.46)	2.78E-05	0.055	658
Information Technology	0.11% (1.85)	0.954	0.035	-0.84% (1.85)	0.648	0.03	432
Materials	-0.73% (4.09)	0.858	0.118	-5.39% (3.61)	0.136	0.084	184

	(Controls Model)			(FF6 Model)		
Real Estate	-1.45%	0.693	0.26	-0.17%	0.96	0.304 244
	(3.68)			(3.47)		
Utilities	-1.21%	0.707	0.264	-3.50%	0.184	0.237 68
	(3.22)			(2.63)		

Note: This table presents coefficients for the ‘Strategic Depth’ factor from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors × 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 23:
Sector-Level Regressions of 12-Month Excess Returns on the Disclosure Sentiment Factor

	(Controls Model)			(FF6 Model)			
Sector	Coefficient	P-Value	R	Coefficient	P-Value	R	N
Communication	-1.04%	0.87	0.125	-2.93%	0.601	0.255	102
	(6.37)			(5.61)			
Consumer Discretionary	-7.85%***	9.92E-04	0.074	-8.68%***	2.52E-04	0.061	489
	(2.38)			(2.37)			
Consumer Staples	-3.06%	0.33	0.099	-3.92%	0.204	0.05	187
	(3.14)			(3.09)			
Energy	6.69%	0.328	0.098	7.80%	0.236	0.173	131
	(6.84)			(6.58)			
Financials	-0.99%	0.774	0.121	-0.71%	0.826	0.123	403
	(3.44)			(3.22)			
Health Care	-7.36%**	0.013	0.035	-3.97%	0.181	0.043	276
	(2.98)			(2.97)			
Industrials	-6.55%***	5.73E-04	0.049	-6.28%***	0.001	0.045	658
	(1.90)			(1.93)			
Information Technology	0.93%	0.722	0.036	-0.38%	0.887	0.03	432
	(2.62)			(2.67)			
Materials	-0.76%	0.858	0.118	-3.52%	0.363	0.081	184
	(4.25)			(3.87)			
Real Estate	-1.66%	0.449	0.261	0.13%	0.948	0.304	244
	(2.19)			(1.96)			

	(Controls Model)			(FF6 Model)			
Utilities	-1.67%	0.61	0.264	-2.62%	0.39	0.228	68
	(3.27)			(3.05)			

Note: This table presents coefficients for the ‘Disclosure Sentiment’ factor from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors $\times$ 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 24:
Sector-Level Regressions of 12-Month Excess Returns on the Talent & Investment Factor

	(Controls Model)			(FF6 Model)			
Sector	Coefficient	P-Value	R	Coefficient	P-Value	R	N
Communication	3.55%	0.563	0.127	-5.91%	0.415	0.261	102
	(6.14)			(7.26)			
Consumer Discretionary	-9.46%***	0.006	0.072	-10.25%***	0.004	0.057	489
	(3.42)			(3.51)			
Consumer Staples	-7.37%	0.124	0.102	-8.35%	0.104	0.052	187
	(4.79)			(5.14)			
Energy	-11.22%	0.544	0.091	-3.61%	0.831	0.16	131
	(18.50)			(16.96)			
Financials	1.39%	0.734	0.121	2.16%	0.571	0.123	403
	(4.09)			(3.80)			
Health Care	-11.48%**	0.011	0.034	-6.73%	0.142	0.044	276
	(4.52)			(4.58)			
Industrials	-9.19%***	4.85E-04	0.052	-9.57%***	2.64E-04	0.05	658
	(2.64)			(2.62)			
Information Technology	-0.78%	0.803	0.035	-2.99%	0.334	0.032	432
	(3.12)			(3.09)			
Materials	-11.66%	0.154	0.128	-14.35%	0.114	0.095	184
	(8.17)			(9.07)			
Real Estate	-1.93%	0.779	0.26	0.31%	0.963	0.304	244
	(6.88)			(6.72)			
Utilities	2.80%	0.669	0.264	-5.44%	0.108	0.234	68
	(6.55)			(3.38)			

Note: This table presents coefficients for the ‘Talent & Investment’ factor from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors × 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 25:
Sector-Level Regressions of 12-Month Excess Returns on the AI Washing Index

Sector	(Controls Model)			(FF6 Model)			N
	Coefficient	P-Value	R	Coefficient	P-Value	R	
Communication	-2.90% (6.31)	0.645	0.126	-7.08% (5.62)	0.208	0.264	102
Consumer Discretionary	-6.88%** (2.73)	0.012	0.066	-8.13%*** (2.76)	0.003	0.053	489
Consumer Staples	-2.09% (3.86)	0.588	0.096	-3.09% (3.90)	0.428	0.045	187
Energy	3.96% (8.31)	0.633	0.09	7.09% (8.23)	0.389	0.166	131
Financials	4.00% (6.63)	0.546	0.123	6.78% (6.16)	0.271	0.129	403
Health Care	-6.76%** (2.81)	0.016	0.032	-0.126828	0.084	0.045	276
Industrials	-8.02%*** (1.87)	1.81E-05	0.055	-7.83%*** (1.86)	2.58E-05	0.052	658
Information Technology	-1.79% (2.44)	0.465	0.037	-3.12% (2.37)	0.188	0.034	432
Materials	-0.146118	0.091	0.123	-7.67%*** (2.97)	0.01	0.089	184
Real Estate	-1.15% (3.76)	0.76	0.26	0.36% (3.48)	0.919	0.304	244
Utilities	-1.87% (3.34)	0.576	0.265	-5.16%** (2.07)	0.012	0.243	68

Note: This table presents coefficients for the ‘AI Washing Index’ from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors × 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 26:
Sector-Level Regressions of 12-Month Excess Returns on the Forward-Looking Factor

Sector	(Controls Model)			(FF6 Model)			N
	Coefficient	P-Value	R	Coefficient	P-Value	R	
Communication	-1.02% (6.66)	0.878	0.125	-3.56% (6.21)	0.567	0.255	102
Consumer Discretionary	-8.76%*** (3.22)	0.007	0.068	-10.06%*** (3.17)	0.002	0.054	489
Consumer Staples	-2.17% (4.30)	0.614	0.096	-4.48% (4.25)	0.292	0.046	187
Energy	-7.87% (12.38)	0.525	0.092	-2.01% (11.20)	0.858	0.159	131
Financials	3.61% (4.94)	0.465	0.122	2.66% (4.18)	0.524	0.124	403
Health Care	-8.43%** (3.30)	0.011	0.032	-0.193848	0.072	0.046	276
Industrials	-10.85%*** (2.71)	6.07E-05	0.057	-10.38%*** (2.69)	1.15E-04	0.052	658
Information Technology	-1.29% (3.02)	0.669	0.036	-3.41% (3.06)	0.264	0.033	432
Materials	-6.93% (8.43)	0.411	0.121	-10.40% (8.34)	0.212	0.085	184
Real Estate	-4.47% (6.00)	0.456	0.261	-1.61% (5.86)	0.784	0.304	244
Utilities	-3.40% (4.33)	0.433	0.267	-7.05%** (2.94)	0.016	0.244	68

Note: This table presents coefficients for the ‘Forward-Looking’ factor from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors × 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 27:
Sector-Level Regressions of 12-Month Excess Returns on the Risk - External Threats Factor

Sector	(Controls Model)			(FF6 Model)			N
	Coefficient	P-Value	R	Coefficient	P-Value	R	

	(Controls Model)			(FF6 Model)			
Communication	-10.29% (10.95)	0.348	0.131	-15.06% (10.66)	0.158	0.269	102
Consumer Discretionary	0.08% (7.92)	0.992	0.053	-0.39% (7.92)	0.96	0.033	489
Consumer Staples	5.17% (10.21)	0.613	0.097	1.27% (9.18)	0.89	0.042	187
Energy	19.72% (13.61)	0.147	0.1	9.63% (14.26)	0.499	0.162	131
Financials	1.14% (4.19)	0.786	0.121	0.06% (4.63)	0.989	0.123	403
Health Care	-11.97% (7.28)	0.1	0.024	-3.53% (7.25)	0.626	0.038	276
Industrials	1.45% (4.50)	0.747	0.035	0.59% (4.45)	0.895	0.032	658
Information Technology	-3.56% (4.84)	0.463	0.036	-1.32% (5.17)	0.799	0.03	432
Materials	-4.98% (13.35)	0.709	0.12	-11.96% (13.50)	0.376	0.085	184
Real Estate	-4.98% (5.75)	0.387	0.263	-3.01% (5.35)	0.574	0.305	244
Utilities	-6.71% (9.97)	0.501	0.273	-4.37% (8.85)	0.621	0.227	68

Note: This table presents coefficients for the ‘Risk - External Threats’ factor from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors $\times$ 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 28:

Sector-Level Regressions of 12-Month Excess Returns on the Risk - Non-Adoption Factor

	(Controls Model)			(FF6 Model)			
Sector	Coefficient	P-Value	R	Coefficient	P-Value	R	N
Communication	-5.78% (11.63)	0.619	0.127	-13.76% (10.47)	0.189	0.268	102

	(Controls Model)			(FF6 Model)			
Consumer Discretionary	-13.70%*** (3.93)	4.89E-04	0.074	-13.91%*** (3.84)	2.89E-04	0.057	489
Consumer Staples	10.07% (10.37)	0.332	0.104	8.47% (10.08)	0.401	0.05	187
Energy	-1.59% (13.77)	0.908	0.088	-0.93% (13.19)	0.944	0.159	131
Financials	-1.01% (3.39)	0.766	0.121	-0.80% (3.60)	0.823	0.123	403
Health Care	-0.571938	0.06	0.028	-4.08% (5.66)	0.47	0.039	276
Industrials	-9.35%*** (3.17)	0.003	0.048	-9.11%*** (3.30)	0.006	0.044	658
Information Technology	2.34% (4.51)	0.603	0.036	0.05% (4.81)	0.991	0.03	432
Materials	3.82% (8.54)	0.655	0.119	-1.82% (9.13)	0.842	0.078	184
Real Estate	-4.96% (4.47)	0.267	0.263	-2.85% (4.76)	0.55	0.305	244
Utilities	-4.62% (10.16)	0.649	0.266	-7.00% (9.30)	0.452	0.231	68

Note: This table presents coefficients for the ‘Risk - Non-Adoption’ factor from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors × 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 29:
Sector-Level Regressions of 12-Month Excess Returns on the Risk - Own Adoption Factor

	(Controls Model)			(FF6 Model)			
Sector	Coefficient	P-Value	R	Coefficient	P-Value	R	N
Communication	11.79% (14.11)	0.404	0.135	2.09% (12.22)	0.865	0.253	102
Consumer Discretionary	-15.98%*** (6.14)	0.009	0.066	-18.31%*** (6.25)	0.003	0.052	489

	(Controls Model)			(FF6 Model)			
Consumer Staples	5.98% (13.56)	0.659	0.095	-3.19% (15.09)	0.833	0.042	187
Energy	18.90% (20.17)	0.349	0.092	20.59% (15.88)	0.195	0.165	131
Financials	3.29% (4.07)	0.419	0.121	5.02% (3.73)	0.178	0.125	403
Health Care	-0.395601	0.098	0.021	-3.36% (4.78)	0.483	0.038	276
Industrials	-11.37%*** (4.40)	0.01	0.044	-12.19%*** (4.32)	0.005	0.043	658
Information Technology	-2.89% (4.56)	0.525	0.036	-6.27% (4.78)	0.189	0.035	432
Materials	4.91% (18.44)	0.79	0.119	-5.70% (55.37)	0.918	0.078	184
Real Estate	-6.48% (8.03)	0.42	0.262	-3.92% (8.75)	0.655	0.305	244
Utilities	9.27% (15.06)	0.538	0.268	0.80% (18.71)	0.966	0.223	68

Note: This table presents coefficients for the ‘Risk - Own Adoption’ factor from regressions of 12-month excess returns run separately for each GICS sector (22 total regressions: 11 sectors $\times$ 2 specifications). The ‘Controls Model’ includes book controls. The ‘FF6 Model’ includes the six Fama-French-Carhart factor exposures. All models include year-fixed effects. Standard errors clustered at the firm level in parentheses. The sample period is 2020–2024. *, **, *** represent statistical significance at the 10%, 5%, and 1% levels, respectively.

Appendix E: Additional Control Variable Summary Statistics

Table 30:

Descriptive Statistics for Fama-French-Carhart Factor Betas (3-, 6-, and 9-Month Horizons)

Variable	N	Mean	Std	Min	P25	Median	P75	Max
Factor Beta Controls (3-Month)								
MKTRF	3980	0.9785	0.4607	-0.405	0.741	0.979	1.231	2.146
SMB	3980	0.6751	0.7359	-1.228	0.209	0.654	1.092	2.705
HML	3980	0.1835	0.7712	-1.885	-0.247	0.15	0.569	2.539
RMW	3980	0.1147	0.7979	-2.498	-0.201	0.174	0.523	2.121

Variable	N	Mean	Std	Min	P25	Median	P75	Max
CMA	3980	0.0629	0.945	-2.885	-0.367	0.087	0.554	2.619
UMD	3980	-0.0445	0.6134	-1.659	-0.377	-0.052	0.269	1.801
Factor Beta Controls (6-Month)								
MKTRF	3970	0.9921	0.4548	-0.427	0.758	1.001	1.238	2.18
SMB	3970	0.6691	0.7461	-1.271	0.189	0.624	1.092	2.585
HML	3970	0.2028	0.6798	-1.578	-0.195	0.179	0.558	2.012
RMW	3970	0.1298	0.8282	-2.558	-0.213	0.179	0.567	2.271
CMA	3970	0.0103	1.0905	-3.463	-0.458	0.071	0.558	3.015
UMD	3970	-0.0305	0.5329	-1.458	-0.315	-0.029	0.267	1.404
Factor Beta Controls (9-Month)								
MKTRF	3843	0.9708	0.4706	-0.587	0.75	0.983	1.223	2.124
SMB	3843	0.6748	0.6995	-1.125	0.227	0.649	1.102	2.448
HML	3843	0.1546	0.7864	-2.027	-0.261	0.13	0.554	2.517
RMW	3843	0.1418	0.8669	-2.622	-0.236	0.197	0.599	2.445
CMA	3843	0.0365	1.1002	-3.212	-0.522	0.05	0.619	3.191
UMD	3843	-0.0296	0.615	-1.879	-0.322	-0.014	0.307	1.608

Note: This table presents summary statistics for the Fama-French-Carhart six-factor betas, estimated over rolling windows for the 3-, 6-, and 9-month horizons. All variables are winsorized. N is the number of firm-year observations.

Appendix F: Data Validation and Robustness

Table 31:
Pearson Correlation Matrix of Key Variables

	ER12m	LogMCap	P/B	ROA	β MKTRF	β SMB	β HML	β RMW	β CMA	β UMD
ER12m	1	-0.05	-0.07	0.05	-0.01	0.04	0.05	0.03	-0.02	0.02
LogMCap	-0.05	1	0.43	0.32	-0.03	-0.42	-0.10	-0.01	0.02	0.03
P/B	-0.07	0.43	1	0.26	0.03	-0.19	-0.20	-0.04	-0.06	0.06
ROA	0.05	0.32	0.26	1	0.00	-0.13	0.01	0.13	-0.03	0.19

	ER12m	LogMCap	P/B	ROA	β MKTRF	β SMB	β HML	β RMW	β CMA	β UMD
β MKTRF	-0.01	-0.03	0.03	0.00	1	0.14	0.07	0.00	0.03	0.01
β SMB	0.04	-0.42	-0.19	-0.13	0.14	1	-0.12	0.37	-0.04	0.05
β HML	0.05	-0.10	-0.20	0.01	0.07	-0.12	1	-0.22	-0.37	0.20
β RMW	0.03	-0.01	-0.04	0.13	0.00	0.37	-0.22	1	-0.06	-0.06
β CMA	-0.02	0.02	-0.06	-0.03	0.03	-0.04	-0.37	-0.06	1	-0.11
β UMD	0.02	0.03	0.06	0.19	0.01	0.05	0.20	-0.06	-0.11	1

Note: This table shows the Pearson correlation coefficients between the 12-month excess return (ER12m) and the main control variables for the 2020–2024 sample period.

Table 32:
Descriptive Statistics of Key Variables Before Winsorization

Variable	N	Mean	Std	Min	Med.	Max	Skew	Kurt
3m ExRet	3325	0.29%	21.06%	-91.98%	-0.84%	255.29%	1.96	17.95
6m ExRet	3270	2.76%	28.99%	-92.70%	0.29%	398.43%	2.65	25.13
9m ExRet	3220	5.18%	38.88%	-93.48%	1.60%	452.26%	2.67	19.63
12m ExRet	3174	4.89%	42.97%	-94.35%	0.13%	490.26%	2.87	19.59
Log M Cap	3989	22.14	1.68	16.22	22.02	28.85	0.37	0.11
P/B	3985	3.29	83.50	-4,471.06	2.35	1,351.32	-40.92	2,168.53
ROA	3984	1.22%	3.29%	-40.59%	1.06%	82.24%	2.57	113.71
MKTRF Beta	3980	0.98	0.72	-7.54	0.98	13.44	0.83	49.49
SMB Beta	3980	0.67	1.13	-10.25	0.65	15.26	0.07	29.64
HML Beta	3980	0.18	1.23	-15.70	0.15	17.10	0.73	45.26
RMW Beta	3980	0.10	1.21	-13.50	0.17	9.36	-1.63	27.05
CMA Beta	3980	0.05	1.67	-19.13	0.09	19.94	-0.31	47.40
UMD Beta	3980	-0.05	0.88	-10.97	-0.05	10.12	-0.37	33.22


Note: This table presents summary statistics for the main dependent and control variables before winsorization. The sample period is 2020–2024.

Table 33:
Descriptive Statistics of Key Variables After Winsorization

Variable	N	Mean	Std	Min	Med.	Max	Skew	Kurt
3m ExRet	3325	0.03%	18.50%	-46.47%	-0.84%	62.45%	0.43	1.08
6m ExRet	3270	2.35%	25.41%	-57.54%	0.29%	97.54%	0.74	1.94
9m ExRet	3220	4.51%	33.99%	-66.74%	1.60%	138.90%	0.98	2.32
12m ExRet	3174	4.12%	37.60%	-72.99%	0.13%	167.17%	1.26	3.57
Log M Cap	3989	22.14	1.64	18.91	22.02	26.41	0.31	-0.40
P/B	3985	4.30	9.68	-27.94	2.35	63.95	3.19	19.28
ROA	3984	1.22%	2.36%	-8.00%	1.06%	8.66%	-0.36	3.73
MKTRF Beta	3980	0.97	0.53	-1.39	0.98	2.66	-0.80	4.88
SMB Beta	3980	0.66	0.86	-2.97	0.65	3.61	-0.29	4.24
HML Beta	3980	0.18	0.87	-3.15	0.15	3.54	0.16	4.15
RMW Beta	3980	0.11	0.95	-4.09	0.17	3.54	-0.89	6.09
CMA Beta	3980	0.06	1.11	-4.83	0.09	4.05	-0.63	5.64
UMD Beta	3980	-0.04	0.69	-2.37	-0.05	2.79	0.46	4.04

Note: This table presents summary statistics for the main dependent and control variables after winsorization. These are the final variables used in the analyses. The sample period is 2020–2024.


Figure 8:
Distribution of AI Factor Grades Assigned for Talent & Investment and AI Washing



Note: This figure shows the percentage of firm-years (N=3,995) assigned each grade (A–E) by the three LLMs for the ‘Talent & Investment’ and ‘AI Washing’ factors.

Figure 9:


Distribution of AI Factor Grades Assigned for Risk - Own Adoption and Risk - External Threats



Note: This figure shows the percentage of firm-years (N=3,995) assigned each grade (A–C) by the three LLMs for the ‘Risk - Own Adoption’ and ‘Risk - External Threats’ factors.


Figure 10:

Distribution of AI Factor Grades Assigned for Risk - Non-Adoption and Forward-Looking




Note: This figure shows the percentage of firm-years (N=3,995) assigned each grade (A–E) by the three LLMs for the ‘Risk - Non-Adoption’ and ‘Forward-Looking’ factors.

Figure 11:
Inter-Model Rank Correlation (Spearman's Rho) of AI Factor Scores



Note: This heatmap shows the Spearman's rho (ρ) rank correlation for AI factor scores between pairs of LLMs. Darker tiles indicate stronger agreement. The sample covers 3,995 firm-years from 2020–2024.

Figure 12:
Inter-Model Linear Correlation (Pearson's r) of AI Factor Scores



Note: This heatmap shows the Pearson's r linear correlation for AI factor scores between pairs of LLMs. Darker tiles indicate stronger correlation. The sample covers 3,995 firm-years from 2020–2024.

Table 34:
Inter-Model Reliability Metrics for AI Factor Scores (GPT-4o-Mini vs. Gemini Flash 1.5)

Factor	Kendall's tau	Spearman's rho	Cohen's weighted kappa	Exact Match	Within 1	N
Strategic Depth	0.59	0.64	0.58	66.60	87.50	3995
Disclosure Sentiment	0.31	0.34	0.31	55.20	69.20	3995
Risk - Own Adoption	0.60	0.61	0.62	85.50	99.50	3995

Factor	Kendall's tau	Spearman's rho	Cohen's weighted kappa	Exact Match	Within 1	N
Risk - External Threats	0.34	0.34	0.31	64.10	99.70	3995
Risk - Non-Adoption	0.45	0.47	0.42	61.60	98.30	3995
Forward-Looking	0.46	0.48	0.39	79.40	88.70	3995
AI Washing Index	0.10	0.10	0.03	66.90	67.30	3995
Talent & Investment	0.44	0.47	0.51	68.50	91.40	3995

Note: This table reports inter-model reliability metrics between GPT-4o-Mini and Gemini Flash 1.5. Kendall's tau and Spearman's rho are rank correlation coefficients. Cohen's weighted kappa measures chance-corrected agreement for ordinal data. 'Exact Match' is the percentage of identical grades; 'Within 1' is the percentage of grades within one level. N=3,995 firm-year assessments for 2020–2024.

Table 35:
Inter-Model Reliability Metrics for AI Factor Scores (Gemini Flash 1.5 vs. Gemini Flash 2.5)

Factor	Kendall's tau	Spearman's rho	Cohen's weighted kappa	Exact Match	Within 1	N
Strategic Depth	0.66	0.70	0.66	69.60	89.70	3995
Disclosure Sentiment	0.34	0.36	0.33	45.60	56.50	3995
Risk - Own Adoption	0.71	0.72	0.68	88.70	98.00	3995
Risk - External Threats	0.37	0.38	0.40	61.80	99.50	3995
Risk - Non-Adoption	0.54	0.57	0.57	61.30	98.50	3995
Forward-Looking	0.54	0.56	0.46	82.20	91.50	3995
AI Washing Index	0.13	0.14	0.03	86.70	87.00	3995

Factor	Kendall's tau	Spearman's rho	Cohen's weighted kappa	Exact Match	Within 1	N
Talent & Investment	0.31	0.33	0.29	68.30	88.50	3995

Note: This table reports inter-model reliability metrics between Gemini Flash 1.5 and Gemini Flash 2.5. Kendall's tau and Spearman's rho are rank correlation coefficients. Cohen's weighted kappa measures chance-corrected agreement for ordinal data. 'Exact Match' is the percentage of identical grades; 'Within 1' is the percentage of grades within one level. N=3,995 firm-year assessments for 2020–2024.

Table 36:
Diagnostic Tests for Regression Models

Test	P value	Result
Breusch-Pagan	6.10E-15	Heteroskedasticity
White Test	6.17E-19	Heteroskedasticity
Durbin-Watson	1.91474604	No autocorrelation

Note: This table summarises results from diagnostic tests on the main regression models. The Breusch-Pagan and White tests check for heteroskedasticity. The Durbin-Watson statistic tests for first-order serial correlation.

Table 37:
Variance Inflation Factors (VIF) for Regression Models

AI Factor	Factor Clean	Factor Type	VIF
AI_Washing	AI Washing	AI Factor	3.30041114
AI_Washing	Beta Mktrf	FF Factor	4.09168705
AI_Washing	Beta Smb	FF Factor	2.09960645
AI_Washing	Beta Hml	FF Factor	1.30145539
AI_Washing	Beta Rmw	FF Factor	1.24385144
AI_Washing	Beta Cma	FF Factor	1.25733613
AI_Washing	Beta Umd	FF Factor	1.06685207
Disclosure_Sentiment	Disclosure Sentiment	AI Factor	4.19660295
Disclosure_Sentiment	Beta Mktrf	FF Factor	4.84900674
Disclosure_Sentiment	Beta Smb	FF Factor	2.10739769
Disclosure_Sentiment	Beta Hml	FF Factor	1.28850189
Disclosure_Sentiment	Beta Rmw	FF Factor	1.24730498

AI Factor	Factor Clean	Factor Type	VIF
Disclosure_Sentiment	Beta Cma	FF Factor	1.25032622
Disclosure_Sentiment	Beta Umd	FF Factor	1.06688475
Forward_Looking	Forward-Looking	AI Factor	4.90696466
Forward_Looking	Beta Mktrf	FF Factor	5.23418426
Forward_Looking	Beta Smb	FF Factor	2.14951961
Forward_Looking	Beta Hml	FF Factor	1.28442932
Forward_Looking	Beta Rmw	FF Factor	1.24723702
Forward_Looking	Beta Cma	FF Factor	1.25058684
Forward_Looking	Beta Umd	FF Factor	1.06847623
Strategic_Depth	Strategic Depth	AI Factor	2.85300499
Strategic_Depth	Beta Mktrf	FF Factor	3.80262312
Strategic_Depth	Beta Smb	FF Factor	2.08600494
Strategic_Depth	Beta Hml	FF Factor	1.30757816
Strategic_Depth	Beta Rmw	FF Factor	1.24364697
Strategic_Depth	Beta Cma	FF Factor	1.25765947
Strategic_Depth	Beta Umd	FF Factor	1.06685345
Talent_Investment	Talent & Investment	AI Factor	4.95242672
Talent_Investment	Beta Mktrf	FF Factor	5.3857436
Talent_Investment	Beta Smb	FF Factor	2.13364521
Talent_Investment	Beta Hml	FF Factor	1.29163016
Talent_Investment	Beta Rmw	FF Factor	1.24800857
Talent_Investment	Beta Cma	FF Factor	1.25203181
Talent_Investment	Beta Umd	FF Factor	1.06702905
Risk_External	Risk External	AI Factor	5.42273201
Risk_External	Beta Mktrf	FF Factor	5.48647869
Risk_External	Beta Smb	FF Factor	2.17950028

AI Factor	Factor Clean	Factor Type	VIF
Risk_External	Beta Hml	FF Factor	1.27246113
Risk_External	Beta Rmw	FF Factor	1.25136427
Risk_External	Beta Cma	FF Factor	1.24729342
Risk_External	Beta Umd	FF Factor	1.06891629
Risk_Non_Adoption	Risk Non - Adoption	AI Factor	5.44791123
Risk_Non_Adoption	Beta Mktrf	FF Factor	5.57401007
Risk_Non_Adoption	Beta Smb	FF Factor	2.17194551
Risk_Non_Adoption	Beta Hml	FF Factor	1.27597782
Risk_Non_Adoption	Beta Rmw	FF Factor	1.24976934
Risk_Non_Adoption	Beta Cma	FF Factor	1.24804522
Risk_Non_Adoption	Beta Umd	FF Factor	1.06887092
Risk_Own_Adoption	Risk Own Adoption	AI Factor	5.5209148
Risk_Own_Adoption	Beta Mktrf	FF Factor	5.60099503
Risk_Own_Adoption	Beta Smb	FF Factor	2.18019879
Risk_Own_Adoption	Beta Hml	FF Factor	1.27607625
Risk_Own_Adoption	Beta Rmw	FF Factor	1.24988556
Risk_Own_Adoption	Beta Cma	FF Factor	1.24759023
Risk_Own_Adoption	Beta Umd	FF Factor	1.06907371

Note: This table reports the Variance Inflation Factors (VIFs) for the independent variables in the main regression models. VIF measures multicollinearity. A VIF value greater than 10 is typically considered problematic.